

YSC4231: Parallel, Concurrent and Distributed Programming

Theory Assignment 3

In this theory assignment, some of the problems refer to an accompanying repository.¹ Feel free to clone it and play with the code, modifying it in any way you like. For your solution, though, you only have to submit a PDF with explanations and proofs — no code is required.

Problem 1. The `AtomicInteger` Java class (in the `java.util.concurrent.atomic` package) is a container for an integer value. One of its methods is

```
def compareAndSet(expect: Int, update: Int): Boolean.
```

This method compares the object's current value to `expect`. If the values are equal, then it atomically replaces the object's value with `update` and returns `true`. Otherwise, it leaves the object's value unchanged, and returns `false`. This class also provides `get() : Int` which returns the object's actual value.

Consider the FIFO queue implementation given by the class `IQueue`. It stores its items in an array `items`. It has two `AtomicInteger` fields: `tail` is the index of the next slot from which to remove an item, and `head` is the index of the next slot in which to place an item. Give an example showing that this implementation is *not* linearizable. Use the tests to help your reasoning.

Problem 2. Consider the class `ReaderWriter`. According to what you have been told about the Java memory model, will the reader method *divide by zero* if

- (a) `writer()` and `reader()` are invoked by two concurred threads (a writer thread and a reader thread, correspondingly) just once each?
- (b) `write()` and `reader()` are repeatedly invoked by two concurred threads (writer/reader) multiple times?

Please, explain your reasoning and supply examples of executions, if necessary. Feel free to play with the class `RWTest`, changing it to emulate both these scenarios (hint: you can create and start new threads within loops).

Problem 3. For each of the histories shown in Figures 1 and 2, are they sequentially consistent? Linearizable? Justify your answer.




Figure 1: First history for Problem 3.

Problem 4. This problem examines a stack implementation `AGMStack` whose `push()` method does not have a single fixed linearization point *in the code*.

The stack stores its items in an `items` array. The `tail` field is an `AtomicInteger`, initially zero. The `push()` method reserves a slot by incrementing `tail`, and then stores the item at that location. Note that these two steps are not atomic: there is an interval after `tail` has been incremented but before the item has been stored in the array.

¹<https://github.com/ysc4231/hw05-linearizability-examples>




Figure 2: Second history for Problem 3.

The `pop()` method reads the value of `tail` and then traverses the array in descending order from the `tail` to slot zero. For each slot, it swaps `None` with the current contents, returning the first non-`None` item it finds. If all slots are `None`, the method returns `None`, indicating an empty stack.

- Give an example execution showing that the linearization point for `push()` cannot occur at the line

```
val j = tail.getAndIncrement()
```

Hint: give an execution where two `push()` calls are *not linearized* (i.e., take effect) in the order they execute this line.

- Give another example execution showing that the linearization point for `push()` cannot occur at the line `items.set(j, Some(x))`
- Since these are the only two memory accesses in `push()`, we must conclude that `push()` has no single fixed linearization point. Does this mean `push()` is not linearizable? Can you then show an execution history of `AGMStack` that is not linearizable or argue that *all* histories produced by `AGMStack` are linearizable somehow (for the latter option you might want to split the set of all histories, which is of course infinite, into a finite set of “classes” of histories, and then argue for linearizability of each such class)?