

# YSC3248: Parallel, Concurrent and Distributed Programming



Concurrent Linked Lists  
Part II

# Last Lecture:


## Fine-Grained Synchronization

- Instead of using a single lock ...
- Split object into
  - Independently-synchronized components
- Methods conflict when they access
  - The same component ...
  - At the same time


# Hand-Over-Hand Again




# Hand-Over-Hand Again




# Hand-Over-Hand Again




# Hand-Over-Hand Again



# Hand-Over-Hand Again




# Hand-Over-Hand Again







# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node



# Removing a Node




# Removing a Node







# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node



# Removing a Node



# Removing a Node





# Remove method

```
def remove(item: T): Boolean = {  
    var pred, curr: Node = null  
    val key = item.hashCode  
  
    try { ... } finally {  
        curr.unlock()  
        pred.unlock()  
    }  
}
```

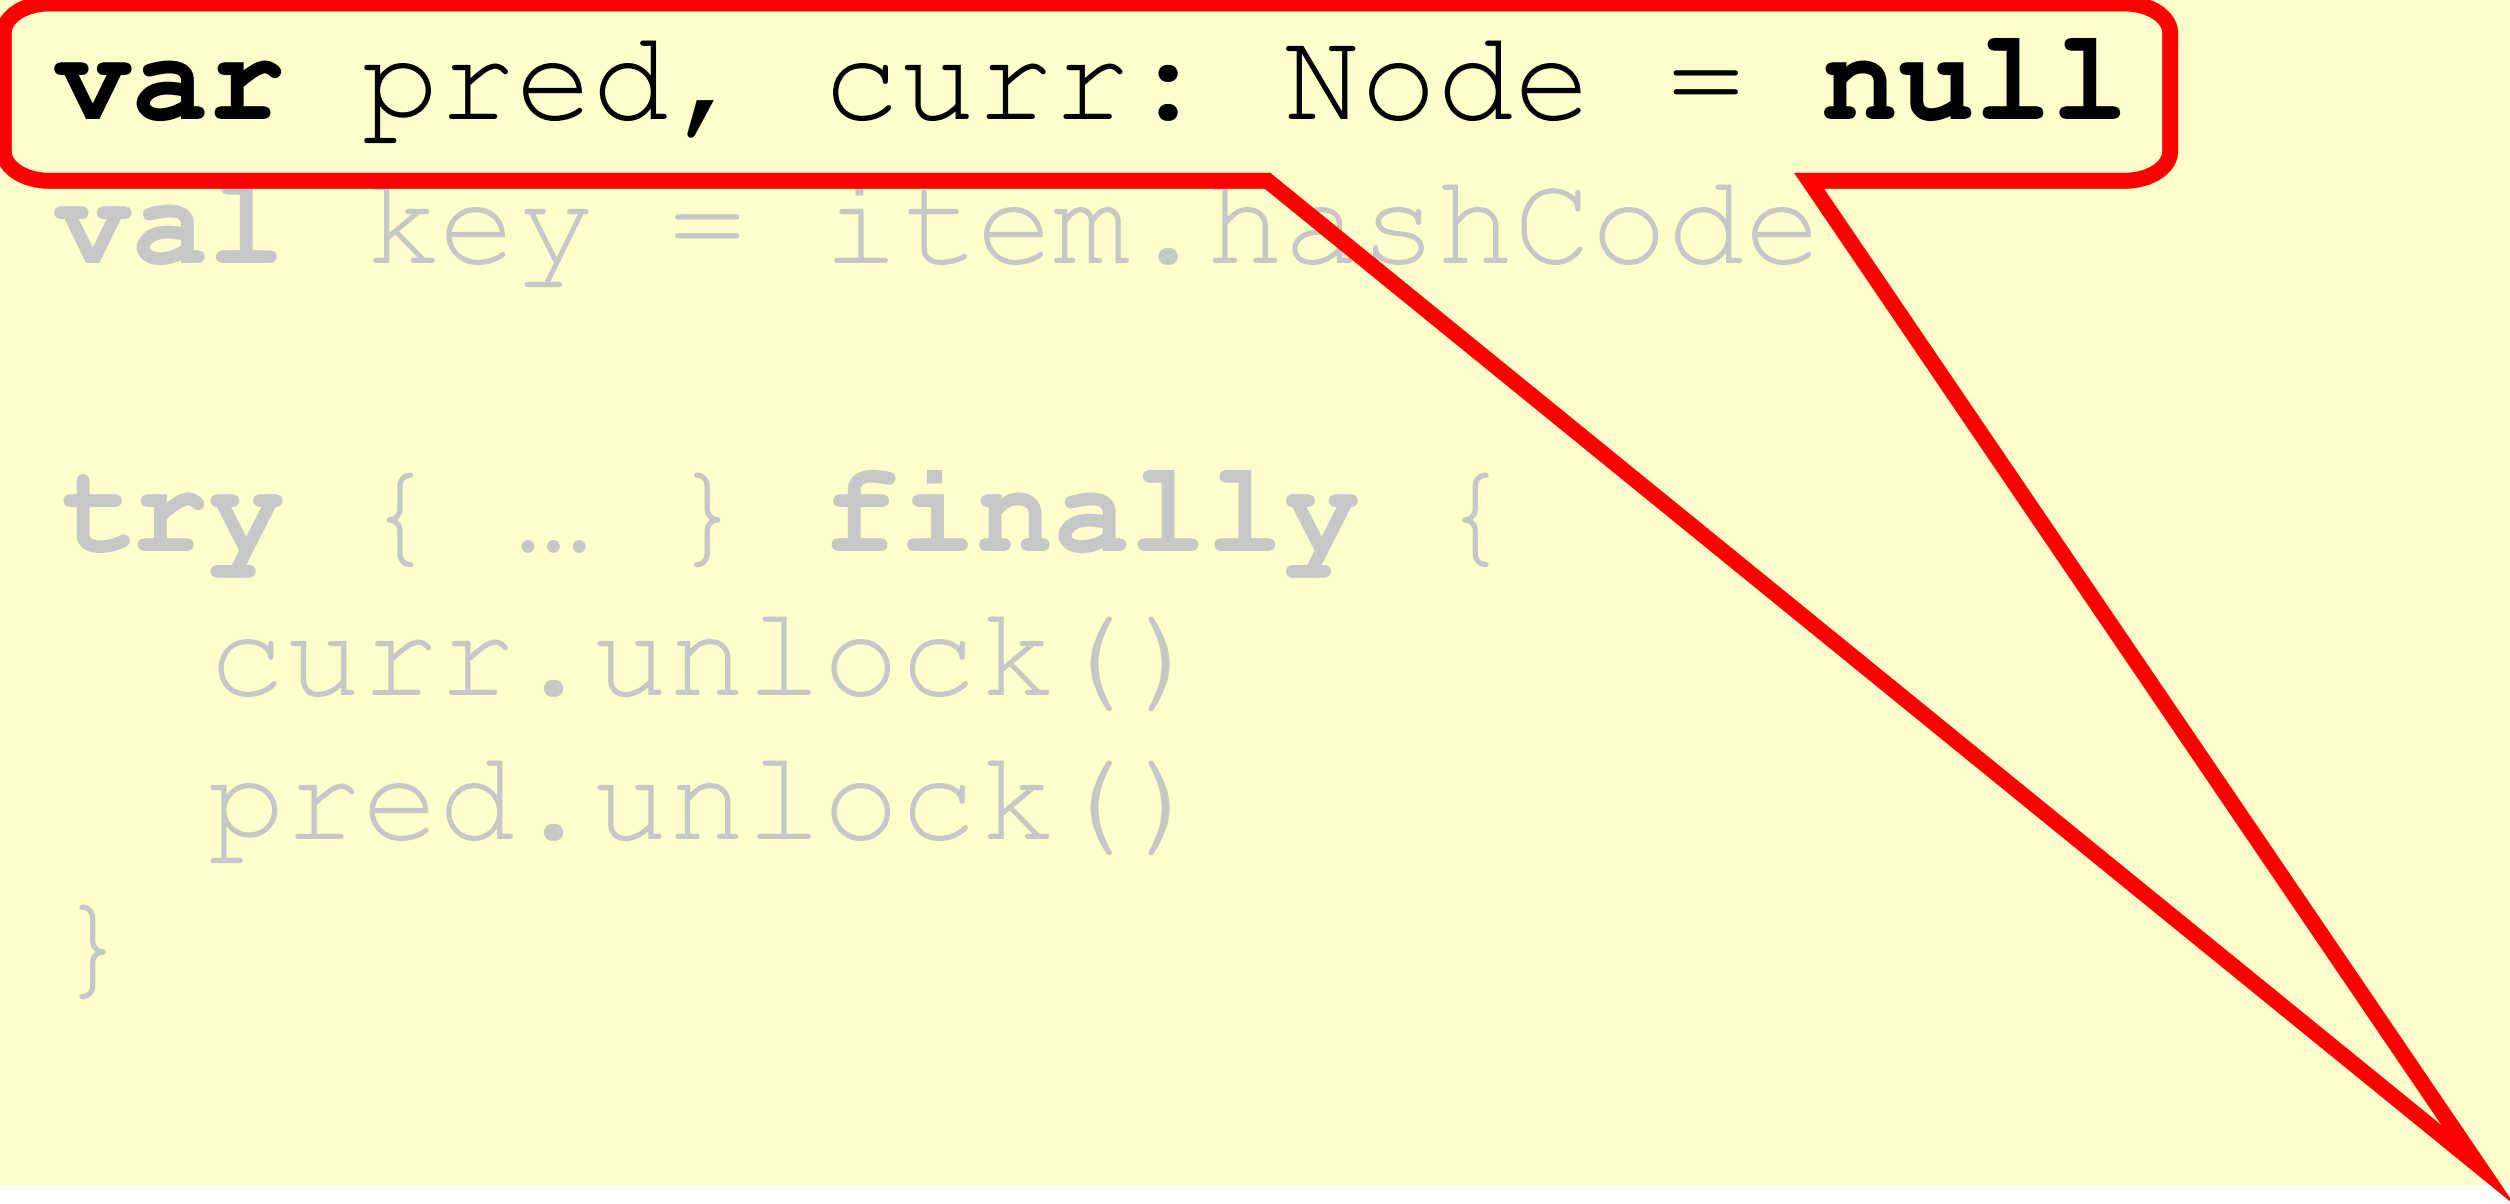
# Remove method

```
def remove(item: T): Boolean = {  
    var pred, curr: Node = null  
    val key = item.hashCode  
  
    try { ... } finally {  
        curr.unlock()  
        pred.unlock()  
    }  
}
```

**Key used to order node**

# Remove method

```
def remove(item: T): Boolean = {  
    var pred, curr: Node = null  
    val key = item.hashCode  
  
    try { ... } finally {  
        curr.unlock()  
        pred.unlock()  
    }  
}
```



**Predecessor and current nodes**

# Remove method

```
def remove(item: T): Boolean = {  
  var pred, curr: Node = null  
  val key = item.hashCode  
  
  try { ... } finally {  
    curr.unlock()  
    pred.unlock()  
  }  
}
```

**Make sure  
locks released**

# Remove method

```
def remove(item: T): Boolean = {  
  var pred, curr: Node = null  
  val key = item.hashCode  
  
  try { ... } finally {  
    curr.unlock()  
    pred.unlock()  
  }  
}
```

**Everything else**


# Remove method

```
try {  
    pred = head  
    pred.lock()  
    curr = pred.next  
    curr.lock()  
    ...  
} finally { ... }
```

# Remove method

**lock pred == head**


```
try {  
    pred = head  
    pred.lock()  
    curr = pred.next  
    curr.lock()  
    ...  
} finally { ... }
```



# Remove method

```
try {  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```

**Lock current**






# Remove method


```
try {  
    pred = head  
    pred.lock()  
    curr = pred.next  
    curr.lock()  
    ...  
} finally { ... }
```

**Traversing list**



# Remove: searching


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```



# Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```


**Search key range**



# Remove: searching

```
while (curr.key <= key) :  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```


**At start of each loop:  
curr and pred locked**



# Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

**If item found, remove node**




# Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

**pred.unlock()**


**Unlock predecessor**



# Remove: searching

Only one node locked!


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```



# Remove: searching

**demote current**

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```






# Remove: searching

## Find and lock new current


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = currNode  
    curr = curr.next  
    curr.lock()  
}  
return false
```



# Remove: searching

**Loop invariant restored**

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = currNode  
    curr = curr.next  
    curr.lock()  
}  
return false
```



# Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}
```

**Otherwise, not present**

**return false**

# Why does this work?

- To remove node  $e$ 
  - Must lock  $e$
  - Must lock  $e$ 's predecessor
- Therefore, if you lock a node
  - It can't be removed
  - And neither can its successor

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

- pred reachable from head
- curr is pred.next
- So curr.item is in the set

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

**Linearization point if  
item is present**

# Why remove() is linearizable

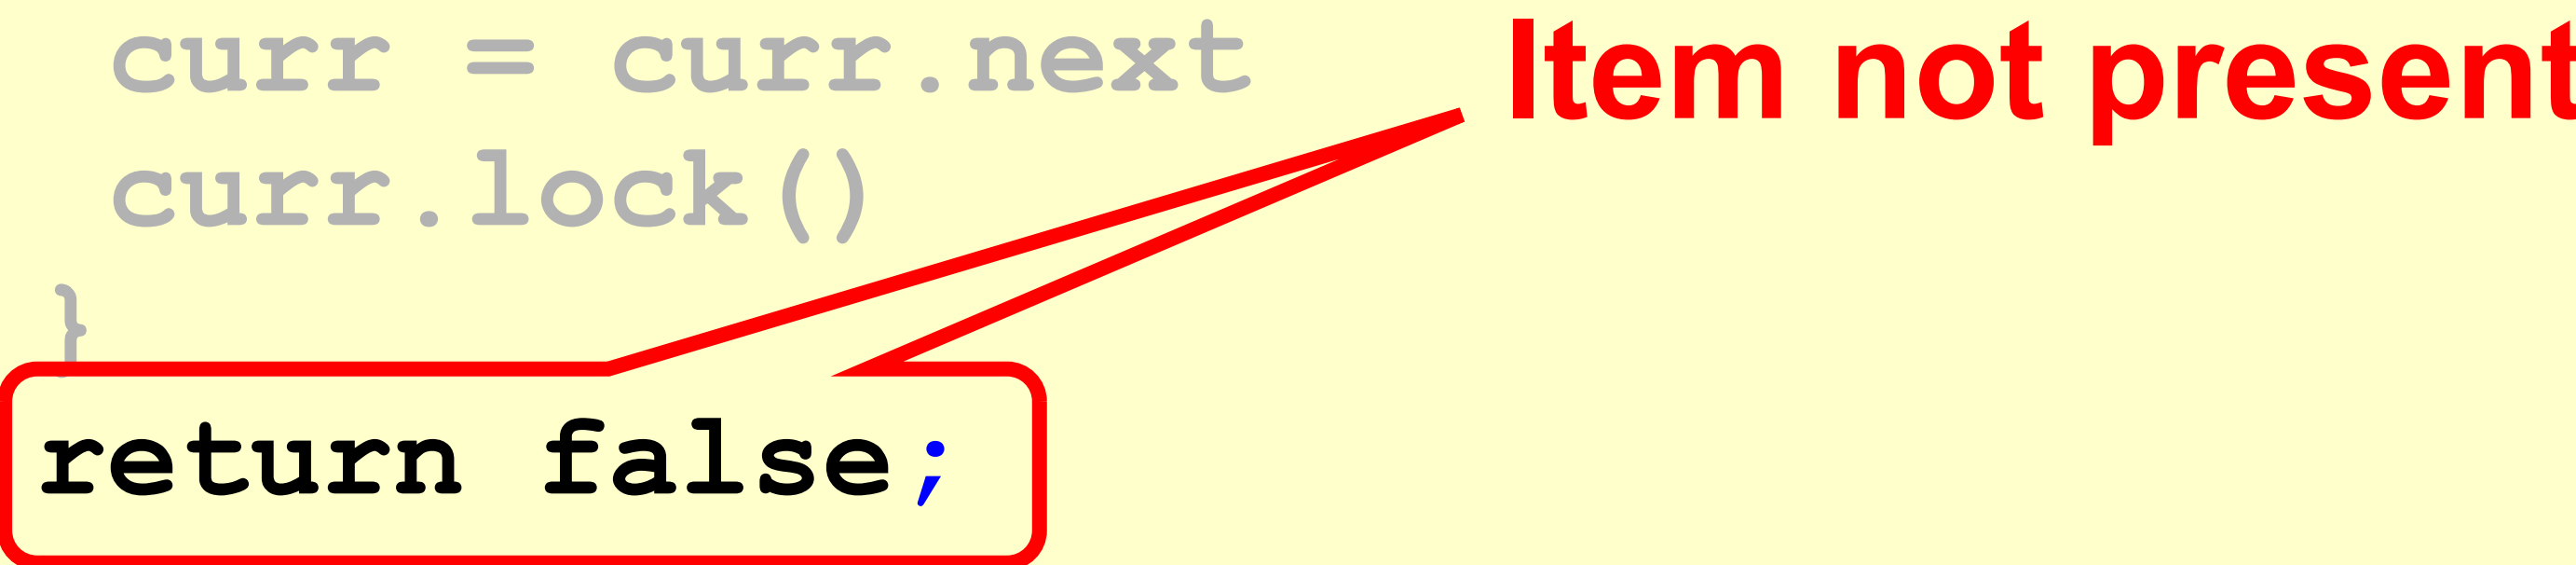
```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

**Node locked, so no other  
thread can remove it ....**

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false;
```

**Item not present**





# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}
```

**return false**

- pred **reachable** from head
- curr **is** pred.next
- pred.key < **key**
- **key** < curr.key

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Linearization point**



# Adding Nodes

- To add node  $e$ 
  - Must lock predecessor
  - Must lock successor
- Neither can be deleted
  - (Is successor lock actually required?)

# Same Abstraction Map

- $S(\text{head}) =$   
 $\{ x \mid \text{there exists } a \text{ such that}$ 
  - $a \text{ reachable from head and}$
  - $a.\text{item} = x$ $\}$

# Rep Invariant

- Easy to check that
  - tail always reachable from head
  - Nodes sorted, no duplicates


# Drawbacks

- Better than coarse-grained lock
  - Threads can traverse in parallel
- Still not ideal
  - Long chain of acquire/release
  - Inefficient

# Optimistic Synchronization


- Find nodes without locking
- Lock nodes
- Check that everything is OK

# Optimistic: Traverse without Locking







# Optimistic: Lock and Load




# Optimistic: Lock and Load




# What could go wrong?




# What could go wrong?




# What could go wrong?




# What could go wrong?



# What could go wrong?

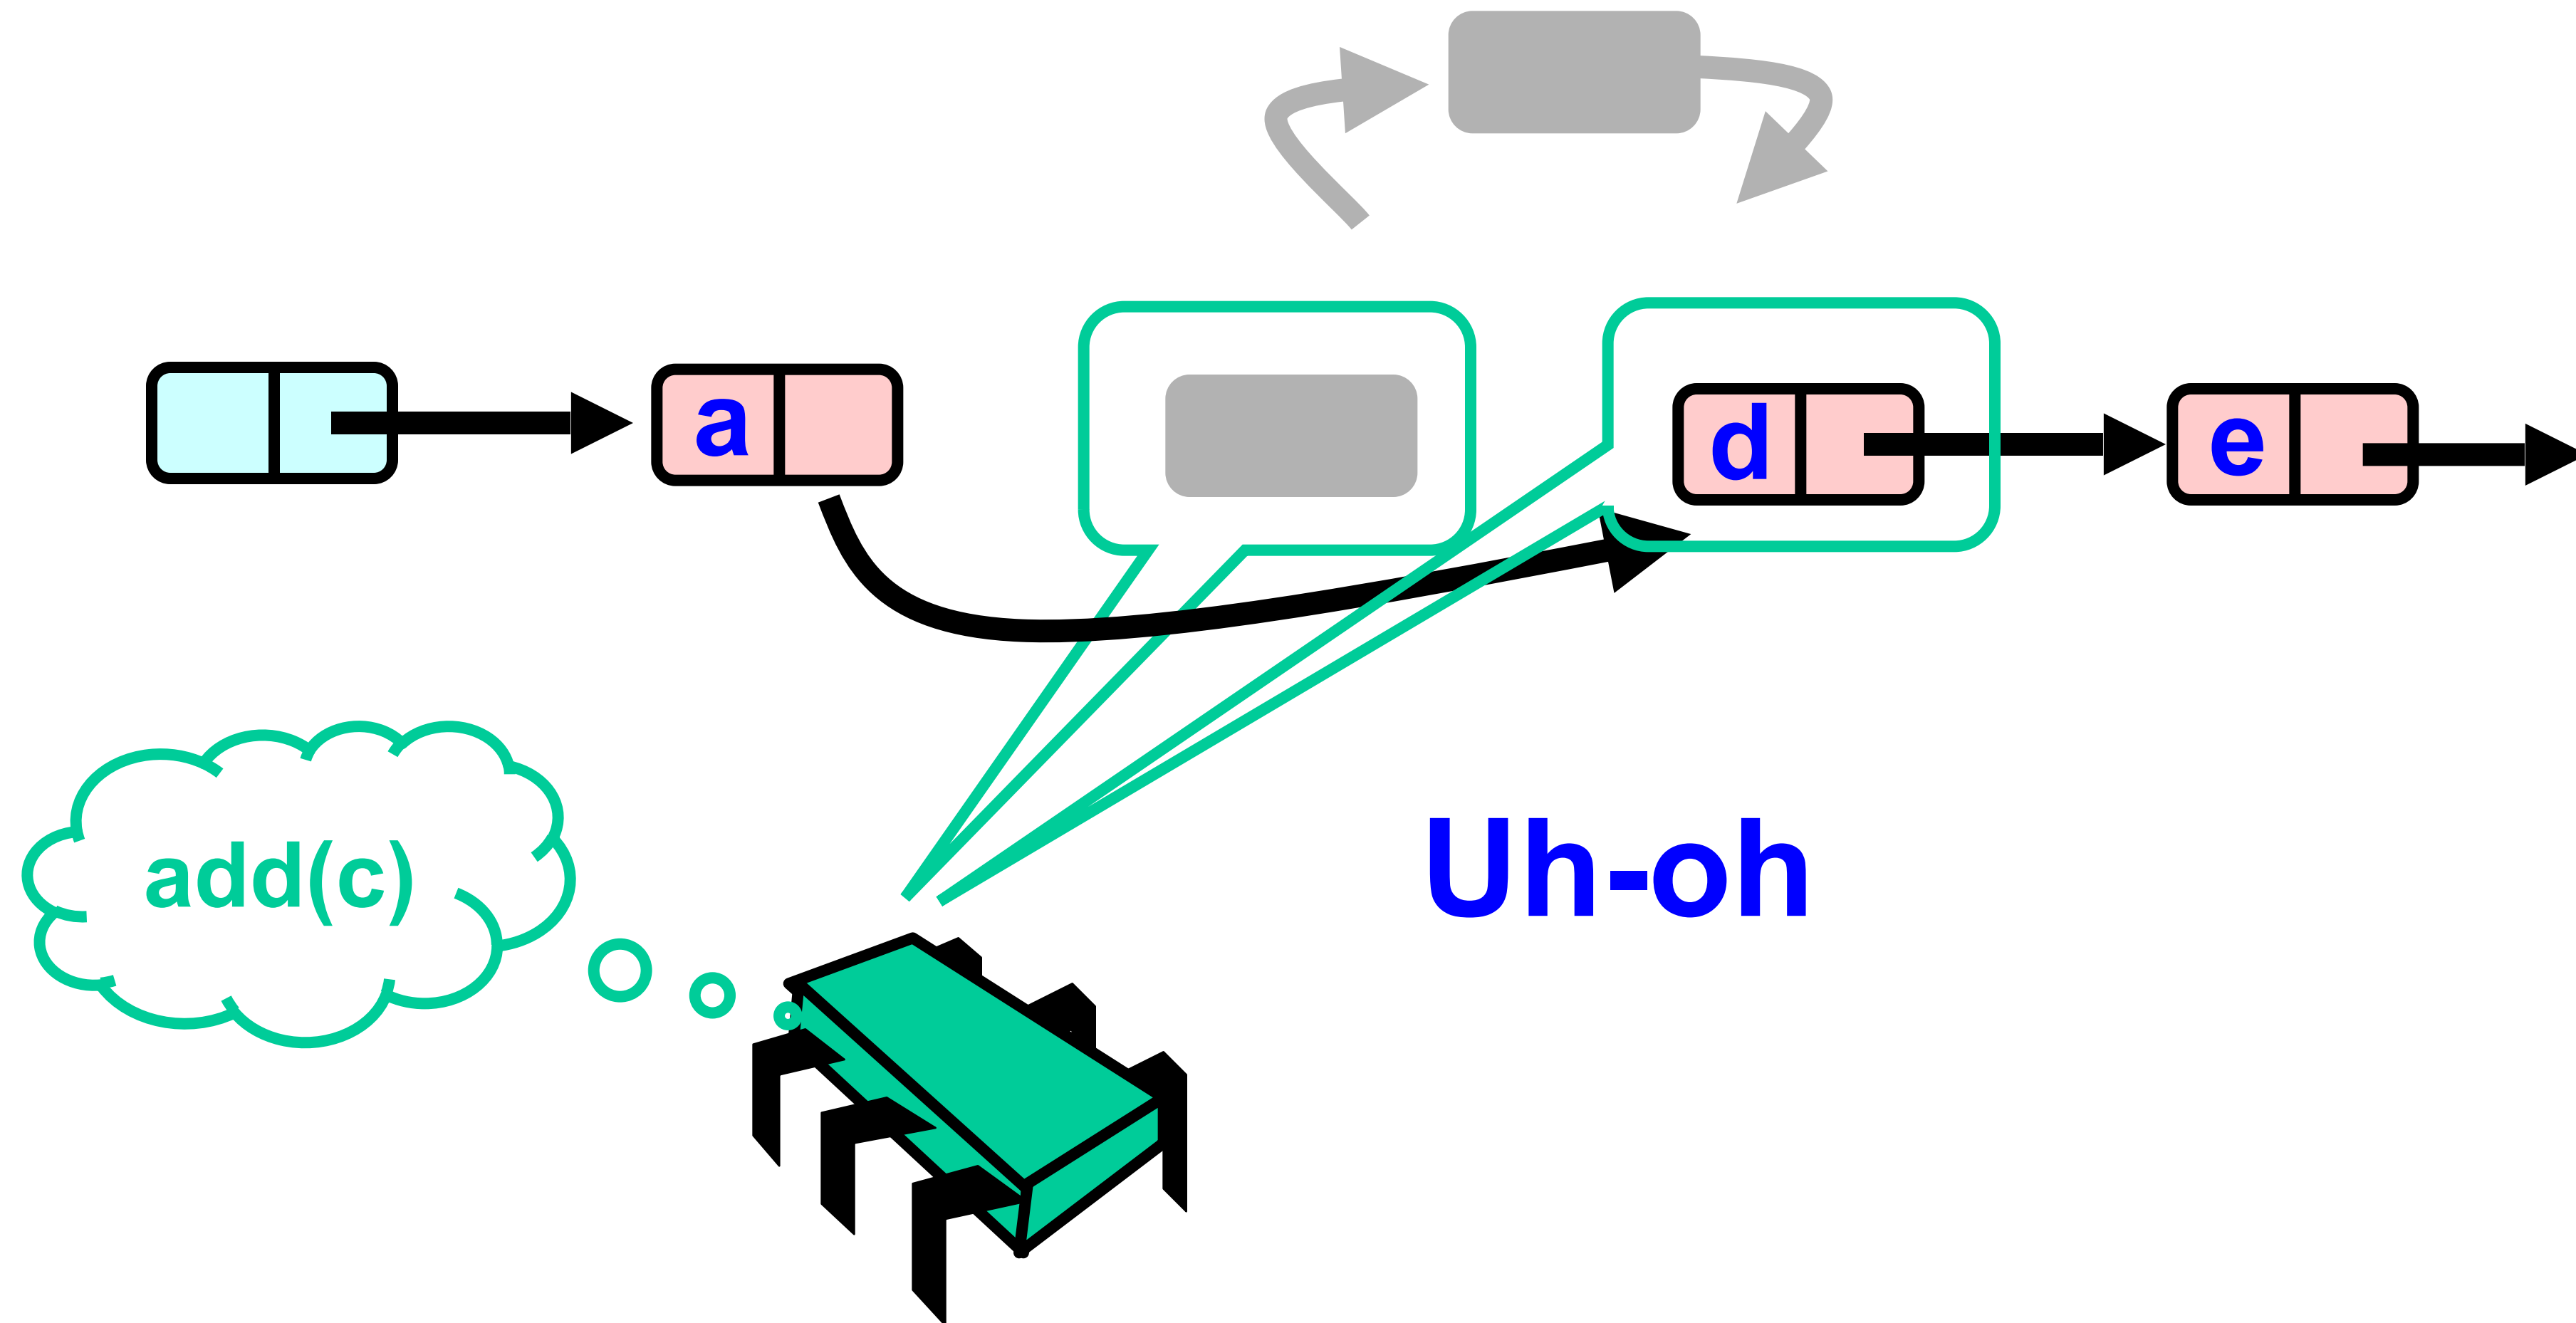


# What could go wrong?

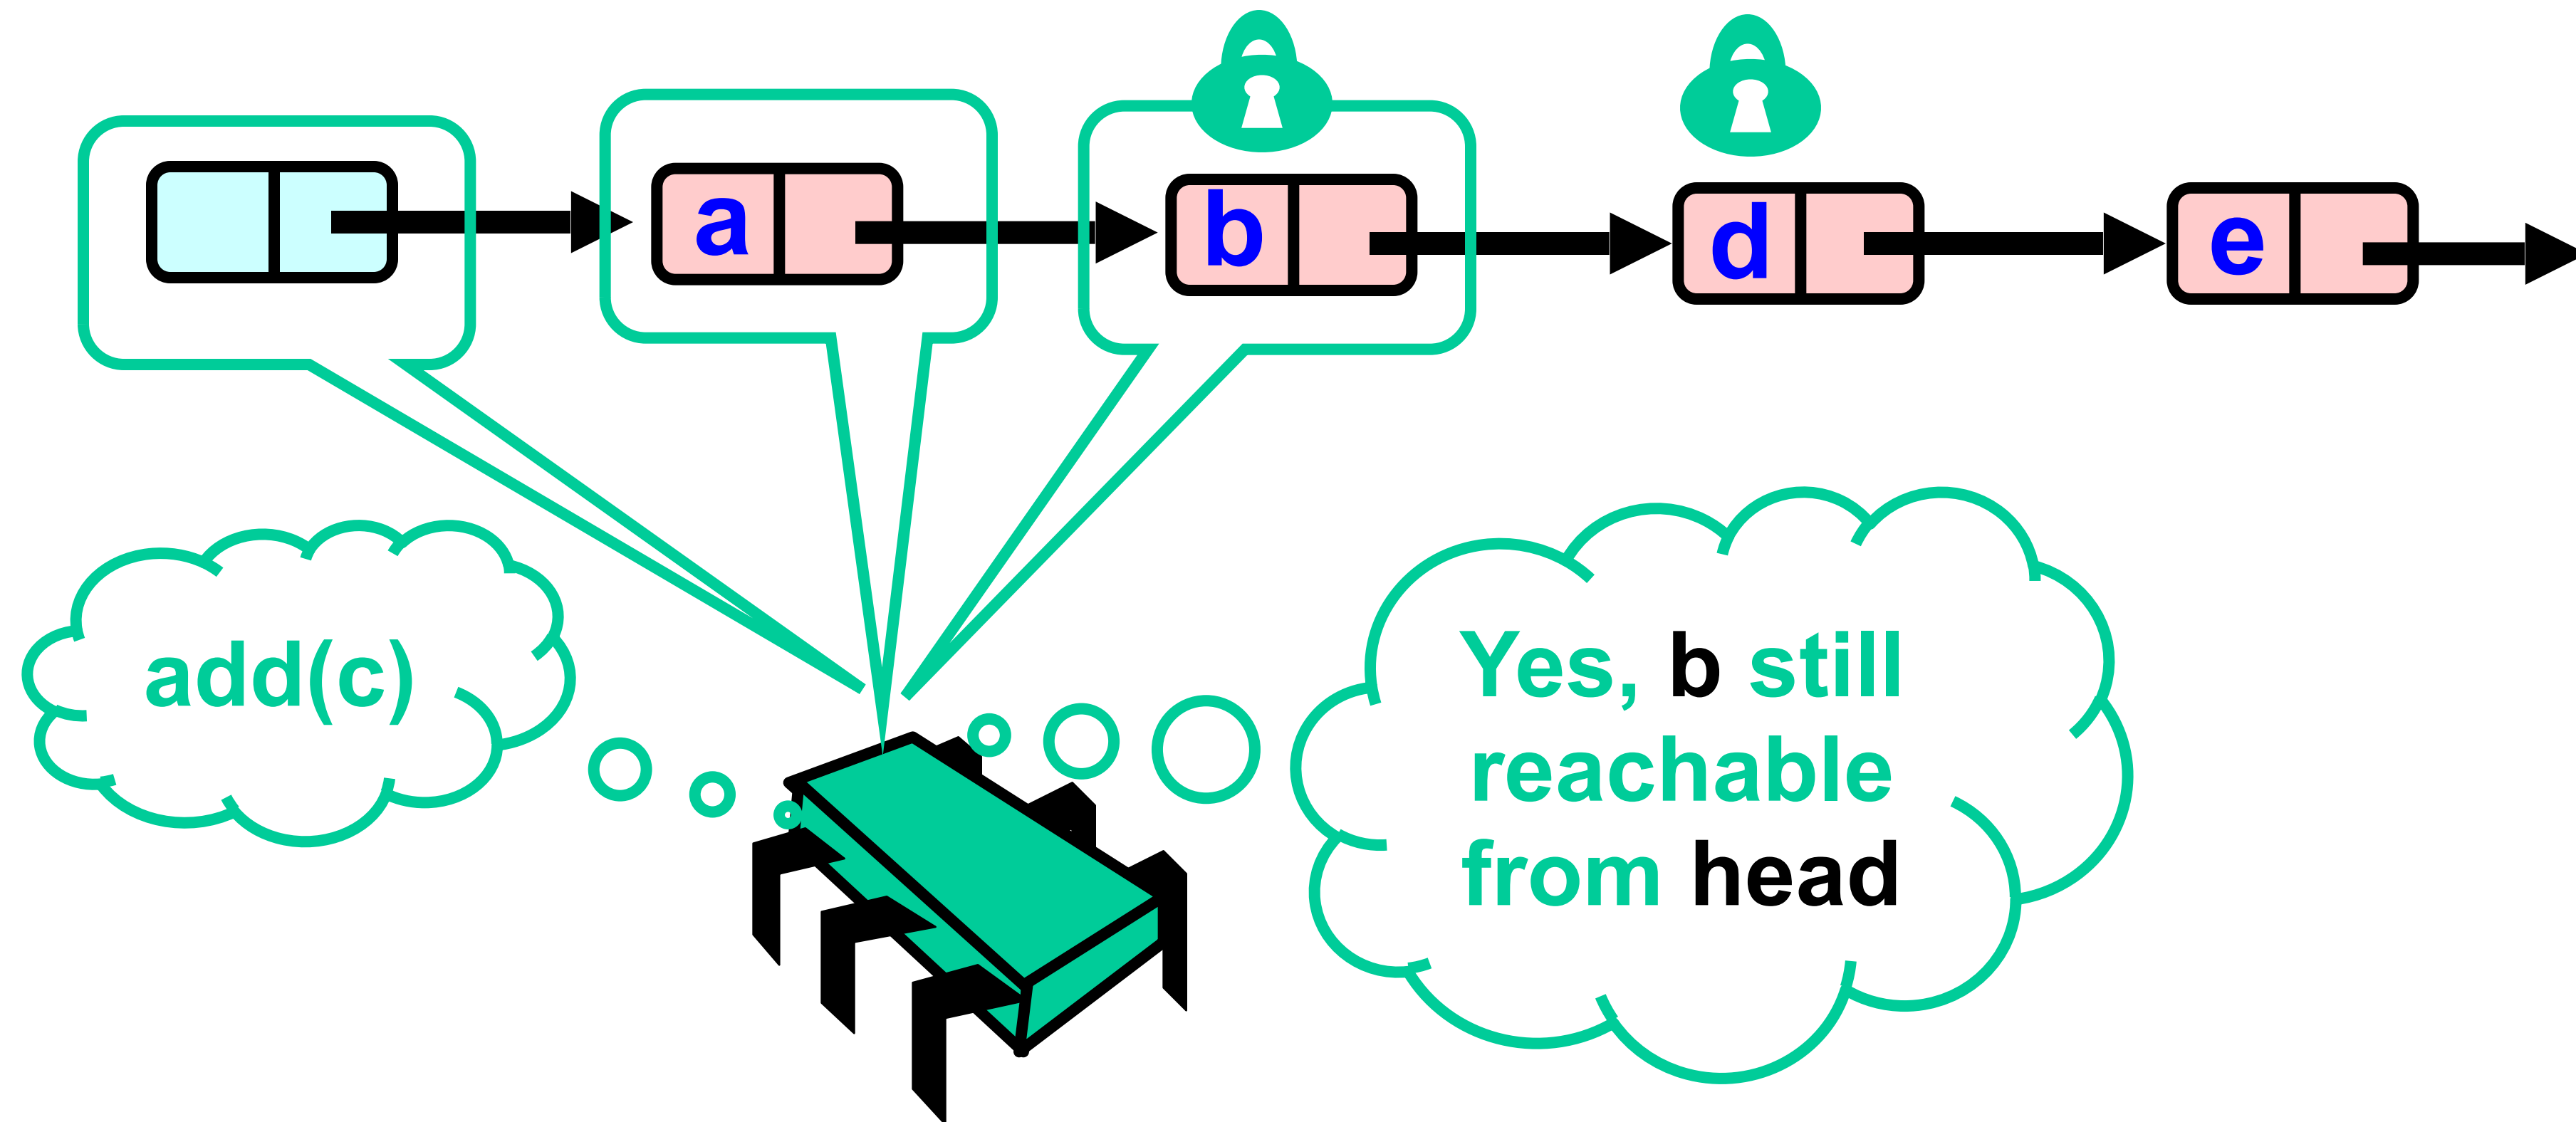





# What could go wrong?




# Validate – Part 1




# What Else Could Go Wrong?




# What Else Could Go Wrong?



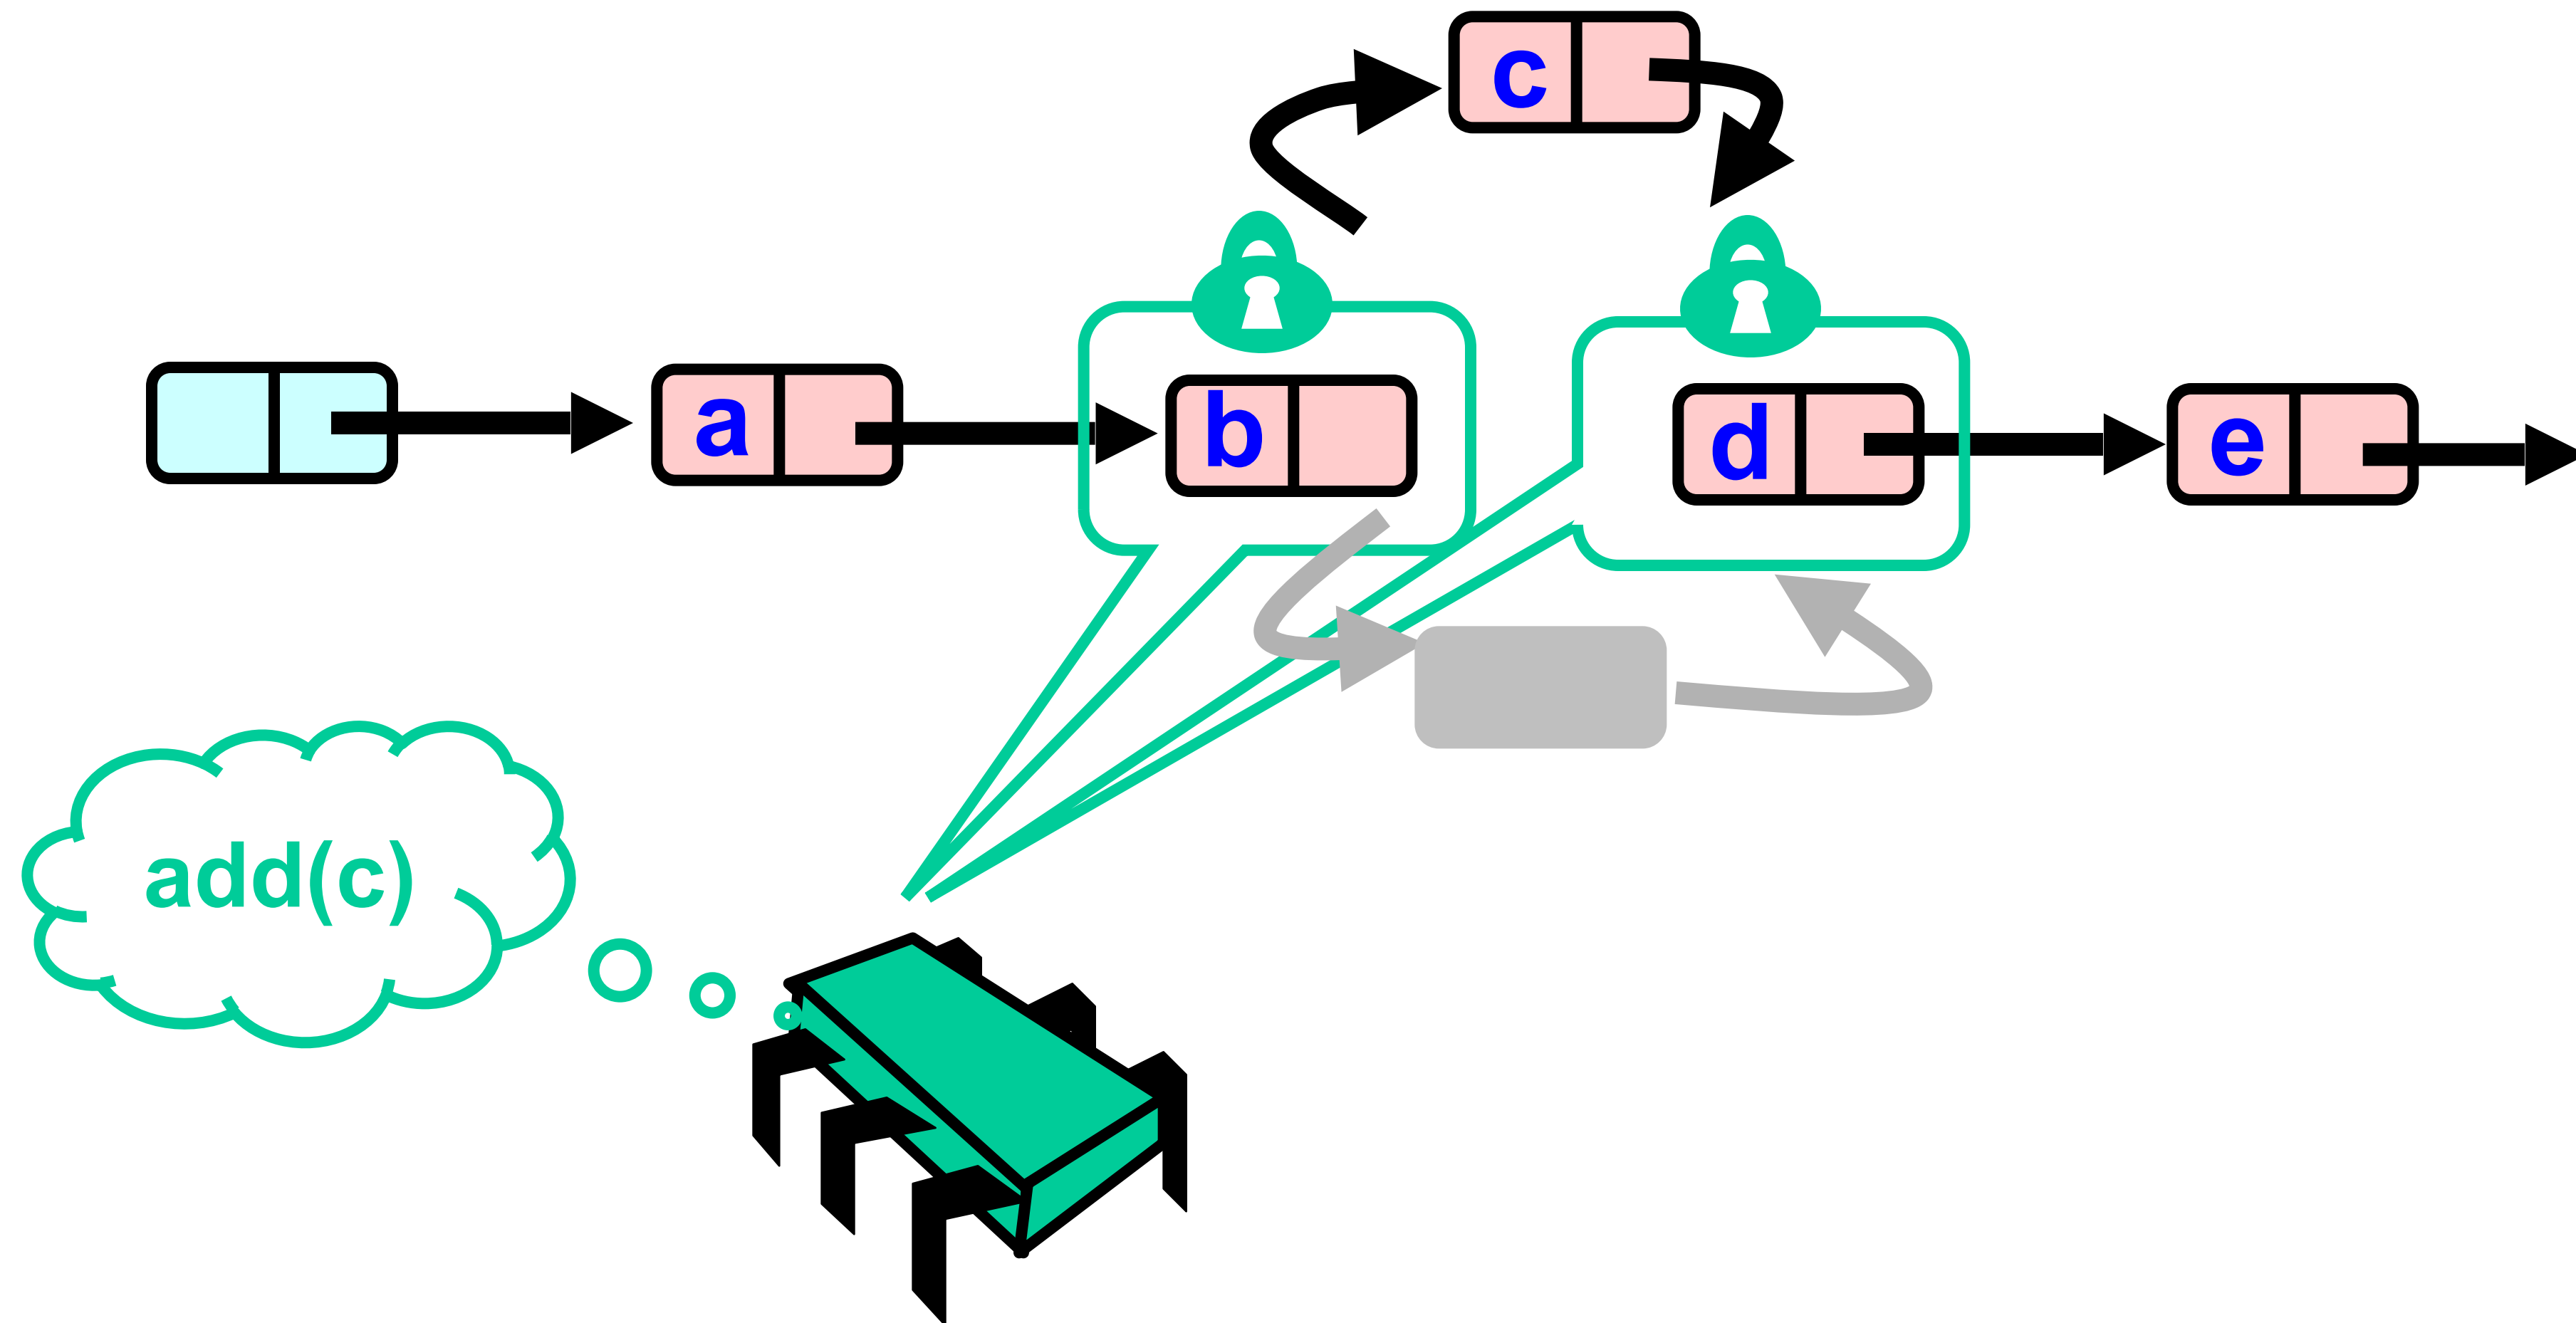
# What Else Could Go Wrong?



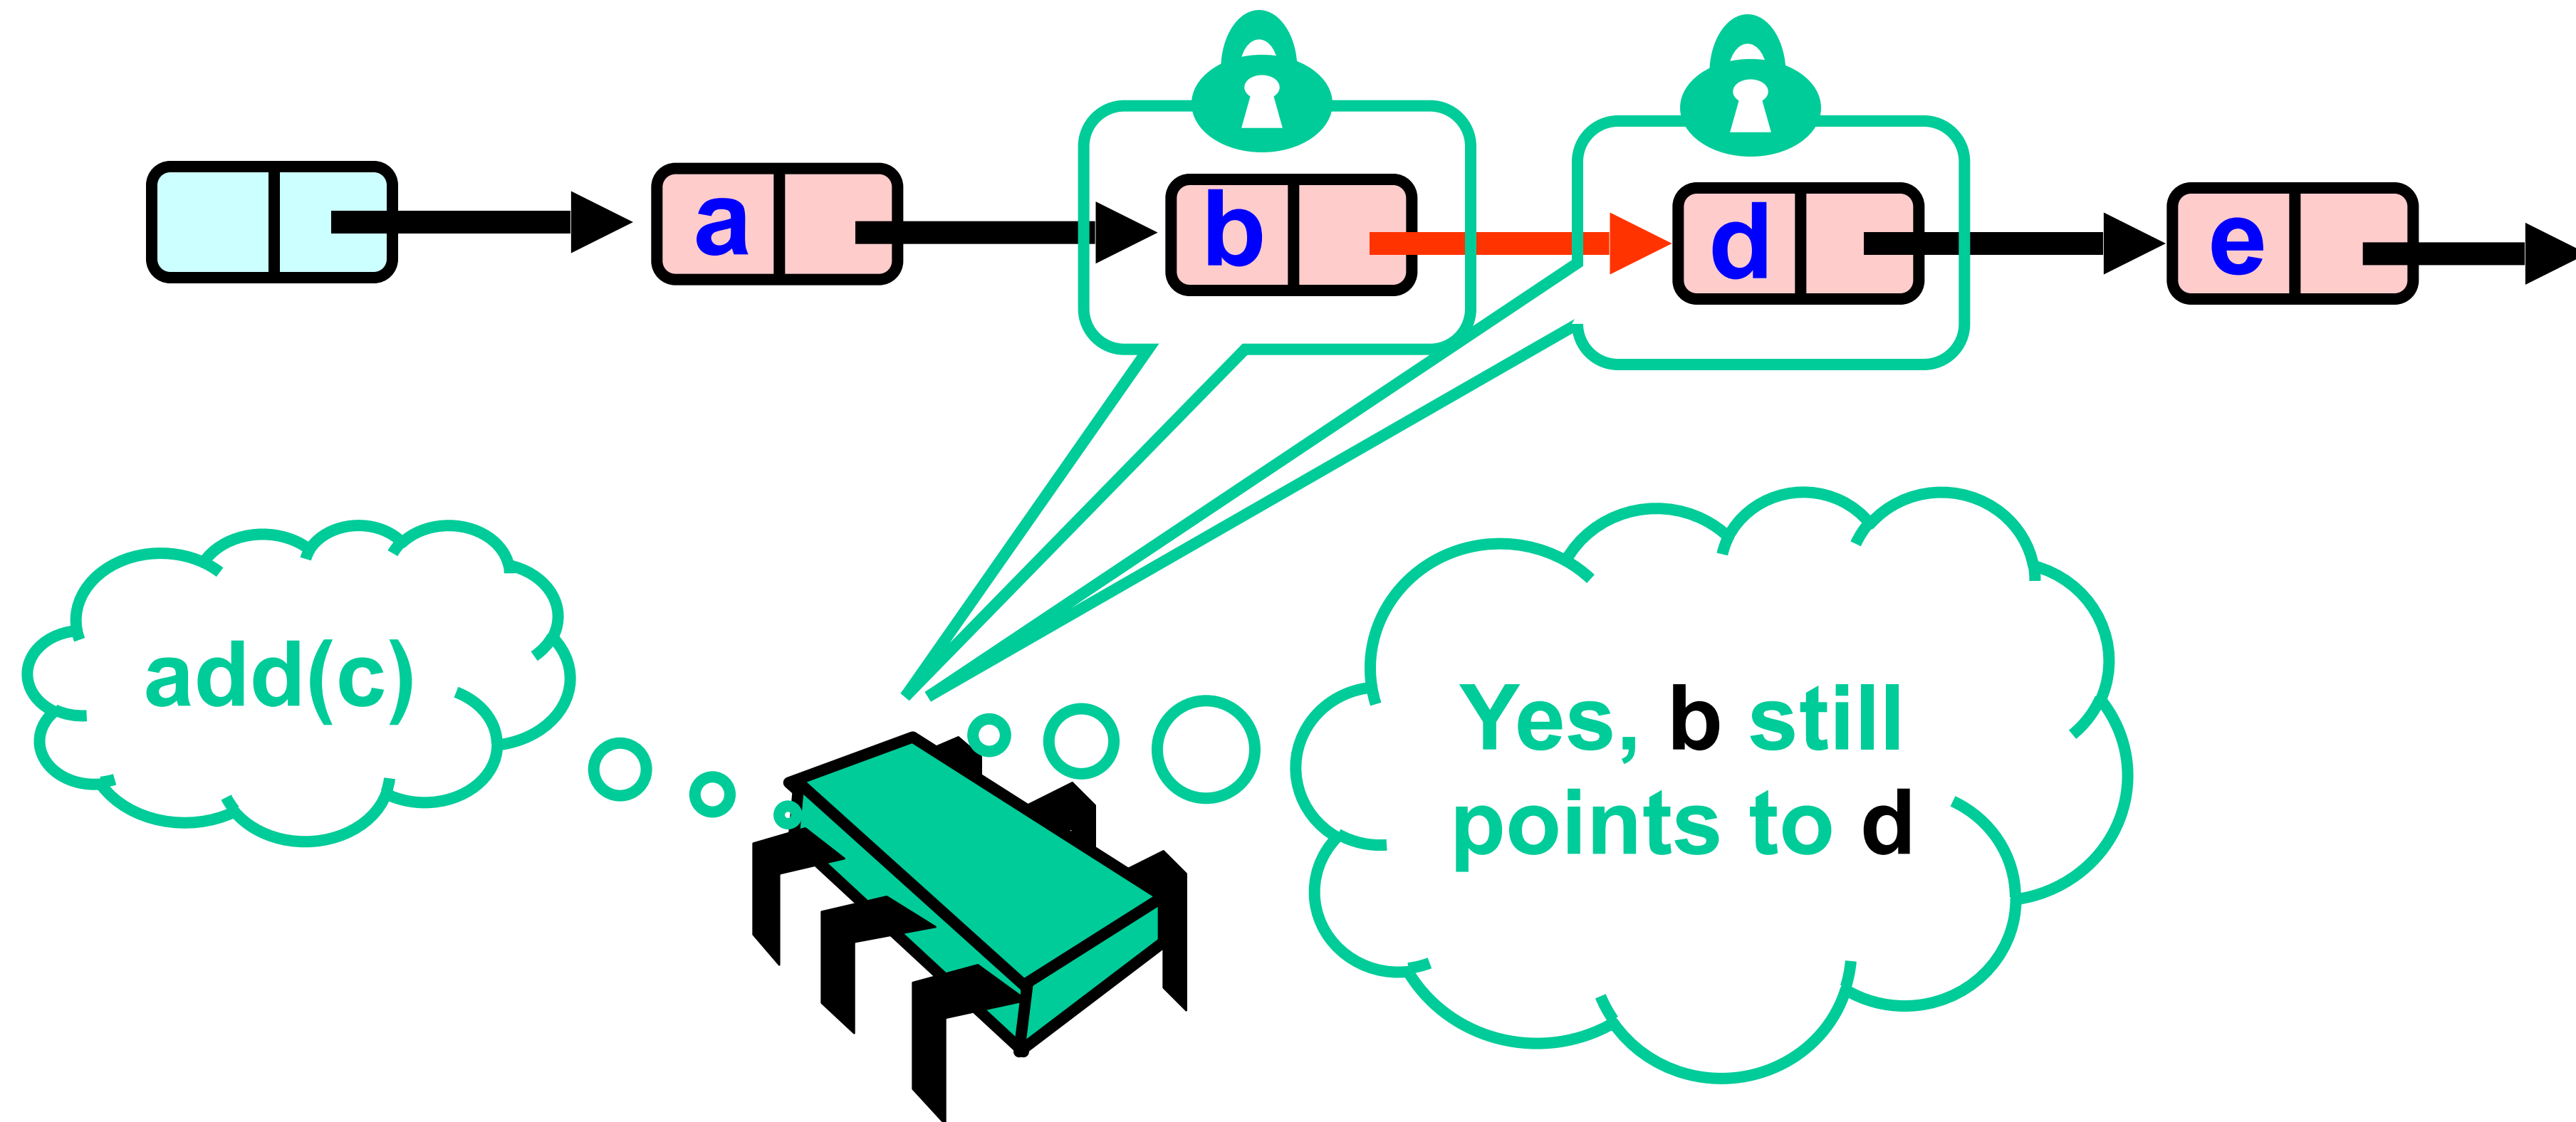
# What Else Could Go Wrong?



# What Else Could Go Wrong?




# Validate Part 2 (while holding locks)





# Optimistic: Linearization Point



# Same Abstraction Map

- $S(\text{head}) =$   
 $\{ x \mid \text{there exists } a \text{ such that}$ 
  - $a \text{ reachable from head and}$
  - $a.\text{item} = x$ $\}$


# Invariants

- Careful: we may traverse deleted nodes
- But we establish properties by
  - Validation
  - After we lock target nodes

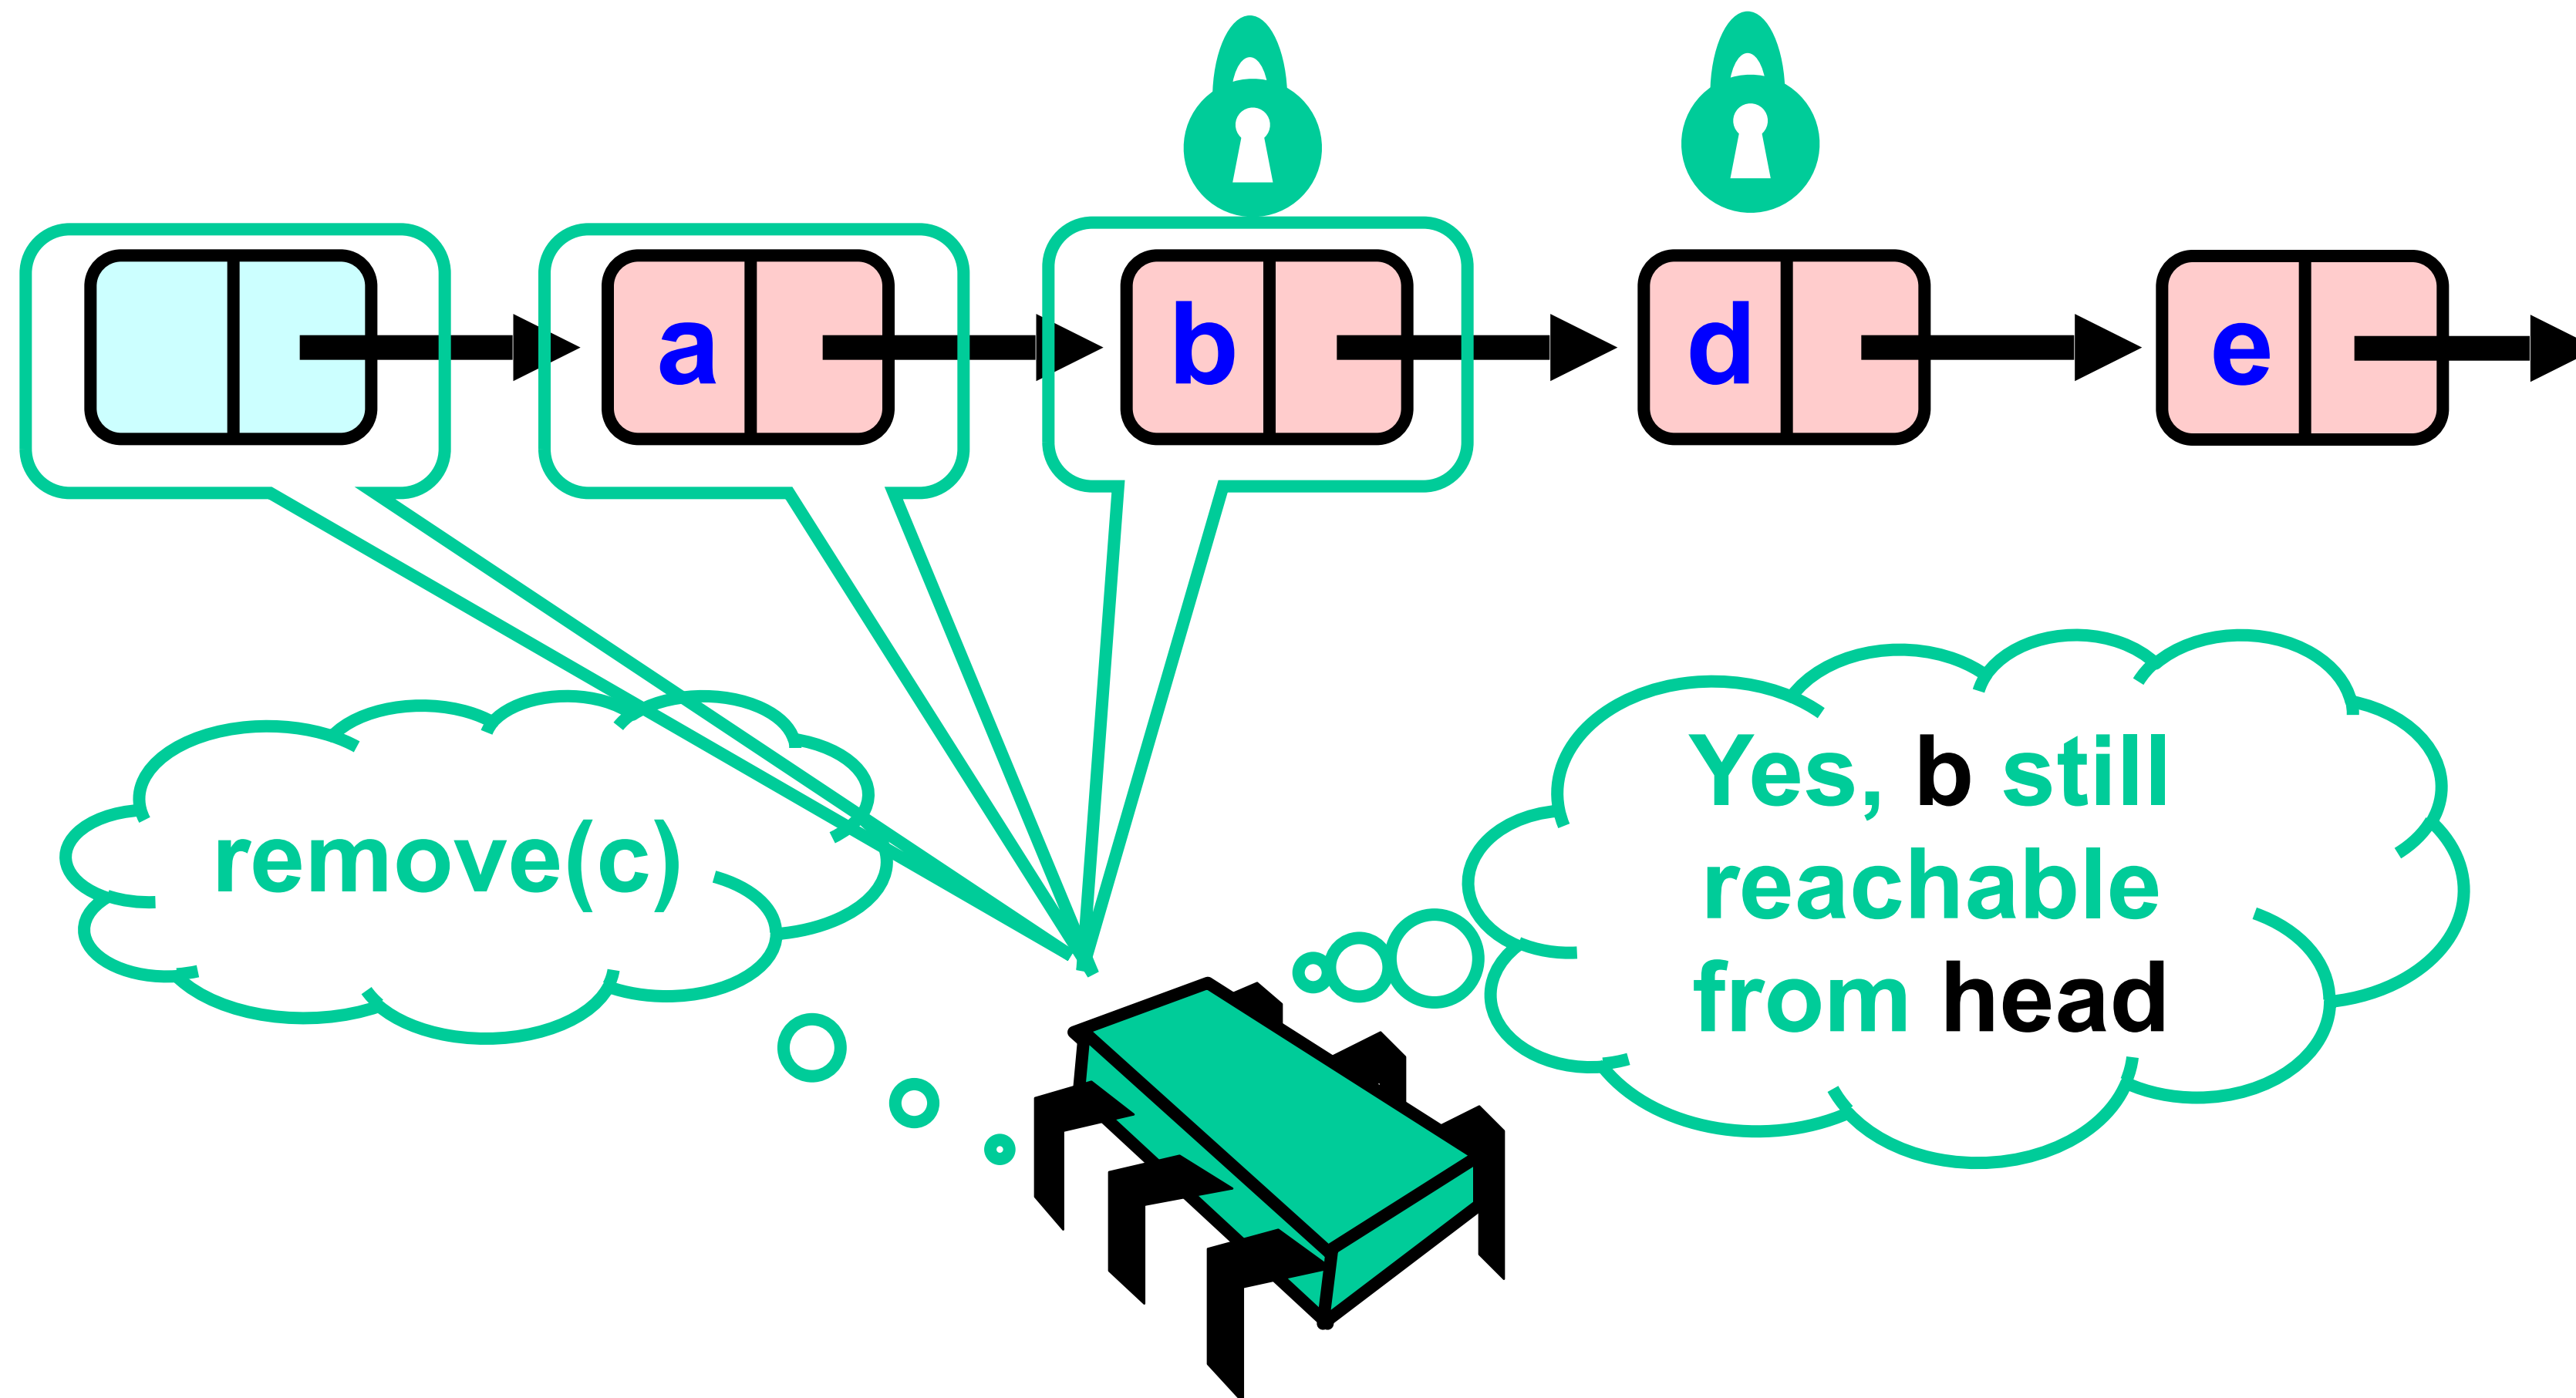
# Removal

- If
  - Nodes **b** and **c** both locked
  - Node **b** still accessible
  - Node **c** still successor to **b**
- Then
  - Neither will be deleted
  - OK to delete and return true


# Unsuccessful Remove




# Validate (1)



# Validate (2)



# OK Computer





# Correctness

- If
  - Nodes **b** and **d** both locked
  - Node **b** still accessible
  - Node **d** still successor to **b**
- Then
  - Neither will be deleted
  - No thread can add **c** after **b**
  - OK to return false


# Validation

```
def validate(pred: Node, curr: Node): Boolean = {  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

# Validation

```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```


**Predecessor &  
current nodes**




# Validation

```
def validate(pred: Node, curr: Node):  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

**Start at the  
beginning**




# Validation



```
def validate(pred: Node, curr: Node):  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

**Search range of keys**


# Validation



```
def validate(pred: Node, curr: Node):  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

**Predecessor reachable**

# Validation




```
def validate(pred: Node, curr: Node):  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

Is current node next?

# Validation

Otherwise move on

```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```






# Validation


## Predecessor not reachable

```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```




# Remove: searching

```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```



# Remove: searching


```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```



**Search key**

# Remove: searching


```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```



**Loop until no synchronization conflict  
(see the code further)**

# Remove: searching

```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```




**Examine predecessor and current nodes**

# Remove: searching

```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```

**Search by key**



# On Exit from While-True-Loop

- If item is present
  - curr holds item
  - pred just before curr
- If item is absent
  - curr has first higher key
  - pred just before curr
- Assuming no synchronization problems

# Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) { // present in list
            pred.next = curr.next
            return true
        } else { // not present in list
            return false
        }
    }
} finally { // always unlock
    pred.unlock(); curr.unlock()
}
```



# Remove Method

```
pred.lock(); curr.lock()
try {
    if (!validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally { // always unlock
    pred.unlock(); curr.unlock()
}
```

**Always unlock**

# Remove Method

```
pred.lock(); curr.lock()
```

```
try {  
    if (validate(pred, curr)) {  
        if (curr.key == key) {  
            pred.next = curr.next  
            return true  
        } else {  
            return false  
        }  
    }  
} finally {  
    pred.unlock(); curr.unlock()  
}
```

**Lock both nodes**

# Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally {
    pred.unlock(); curr.unlock()
}
```

**Check for synchronization conflicts**

# Remove Method


```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally {
    pred.unlock(); curr.unlock()
}
```

**target found, remove  
node**

# Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally {
    pred.unlock(); curr.unlock()
}
```

**target not found**



# Optimistic List

- Limited hot-spots
  - Targets of `add()`, `remove()`, `contains()`
  - No contention on traversals
- Moreover
  - Traversals are wait-free
  - Food for thought ...

# So Far, So Good

- Much less lock acquisition/release
  - Performance
  - Concurrency
- Problems
  - Need to traverse list twice
  - `contains()` method acquires locks

# Evaluation

- Optimistic is effective if
  - cost of scanning twice without locks is less than
  - cost of scanning once with locks
- Drawback
  - `contains()` acquires locks
  - 90% of calls in many apps




# Lazy List

- Like optimistic, except
  - Scan once
  - `contains(x)` never locks ...
- Key insight
  - Removing nodes causes trouble
  - Do it “lazily”


# Lazy List

- **remove ()**
  - Scans list (as before)
  - Locks predecessor & current (as before)
- Logical delete
  - Marks current node as removed (new!)
- Physical delete
  - Redirects predecessor's next (as before)


# Lazy Removal




# Lazy Removal




# Lazy Removal



# Lazy Removal



# Lazy Removal



# Lazy List


- All Methods
  - Scan through locked and marked nodes
  - Removing a node doesn't slow down other method calls ...
- Must still lock pred and curr nodes.




# Validation

- No need to rescan list!
- Check that `pred` is not marked
- Check that `curr` is not marked
- Check that `pred` points to `curr`


# Business as Usual




# Business as Usual




# Business as Usual




# Business as Usual




# Business as Usual



# Business as Usual




# Business as Usual






# Business as Usual



# Business as Usual



# New Abstraction Map

- $S(\text{head}) =$   
 $\{ x \mid \text{there exists node } a \text{ such that}$ 
  - $a$  reachable from head and
  - $a.\text{item} = x$  and
  - $a$  is unmarked $\}$

# Invariant

- If not marked then item in the set
- and is reachable from head
- and if not yet traversed it is reachable from pred

# Validation

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```

# List Validate Method

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```

**Predecessor not  
Logically removed**

# List Validate Method

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```



**Current not  
Logically removed**

# List Validate Method

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```

**Predecessor still  
Points to current**



# Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true  
            pred.next = curr.next  
            return true;  
        } else {  
            return false  
        }  
    } finally {  
        pred.unlock()  
        curr.unlock()  
    }  
}
```

# Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true  
            pred.next = curr.next  
            return true  
        } else {  
            return false  
        }  
    }  
    finally {  
        pred.unlock()  
        curr.unlock()  
    }  
}
```

**Validate as before**

# Remove

```
try {  
    pred.lock(); curr.lock();  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true;  
            pred.next = curr.next;  
            return true;  
        } else {  
            return false;  
        }  
    }  
} finally {  
    pred.unlock();  
    curr.unlock();  
}
```

**Key found**

# Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true;  
            pred.next = curr.next  
            return true  
        } else {  
            return false  
        }  
    } finally {  
        pred.unlock()  
        curr.unlock()  
    }  
}
```

**Logical remove**

# Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true  
            pred.next = curr.next;  
            return true  
        } else {  
            return false  
        }  
    }  
} finally {  
    pred.unlock()  
    curr.unlock()  
}
```

**physical remove**

# Contains

```
def contains(item: T) = {  
    val key = item.hashCode  
    var curr = this.head  
    while (curr.key < key) curr = curr.next  
    curr.key == key && !curr.marked  
}
```

# Contains

```
def contains(item: T) = {  
    val key = item.hashCode  
    var curr = this.head  
    while (curr.key < key) curr = curr.next  
    curr.key == key && !curr.marked  
}
```

**Start at the head**

# Contains

```
def contains(item: T) = {  
    val key = item.hashCode  
    var curr = this.head  
    while (curr.key < key) curr = curr.next  
    curr.key == key && !curr.marked  
}
```

**Search key range**



# Contains

```
def contains(item: T) = {  
  val key = item.hashCode  
  var curr = this.head  
  while (curr.key < key) curr = curr.next  
  curr.key == key && !curr.marked  
}
```


**Traverse *without locking***  
**(nodes may have been removed)**

# Contains

```
def contains(item: T) = {  
  val key = item.hashCode  
  var curr = this.head  
  while (curr.key < key) curr = curr.next  
  curr.key == key && !curr.marked  
}
```

**Present and undeleted?**


# Summary: Wait-free Contains



Use Mark bit + list ordering

1. Not marked  $\rightarrow$  in the set
2. Marked or missing  $\rightarrow$  not in the set

# Lazy List



Lazy add () and remove () + Wait-free contains ()

# Evaluation

- Good:
  - **contains()** doesn't lock
  - In fact, its wait-free!
  - Good because typically high % contains()
  - Uncontended calls don't re-traverse
- Bad
  - Contended **add()** and **remove()** calls must re-traverse
  - Traffic jam if one thread delays

# Traffic Jam

- Any concurrent data structure based on mutual exclusion has a weakness
- If one thread
  - Enters critical section
  - And “eats the big muffin”
    - Cache miss, page fault, descheduled ...
  - Everyone else using that lock is stuck!
  - Need to trust the scheduler....

# Reminder: Lock-Free Data Structures



- No matter what ...
  - Guarantees minimal progress in any execution
  - i.e. Some thread will always complete a method call
  - Even if others halt at malicious times
  - Implies that implementation can't use locks


# Lock-free Lists

- Next logical step
  - Wait-free `contains()`
  - lock-free `add()` and `remove()`
- Use only `compareAndSet()`
  - What could go wrong?



# Lock-free Lists

Logical Removal




Physical Removal

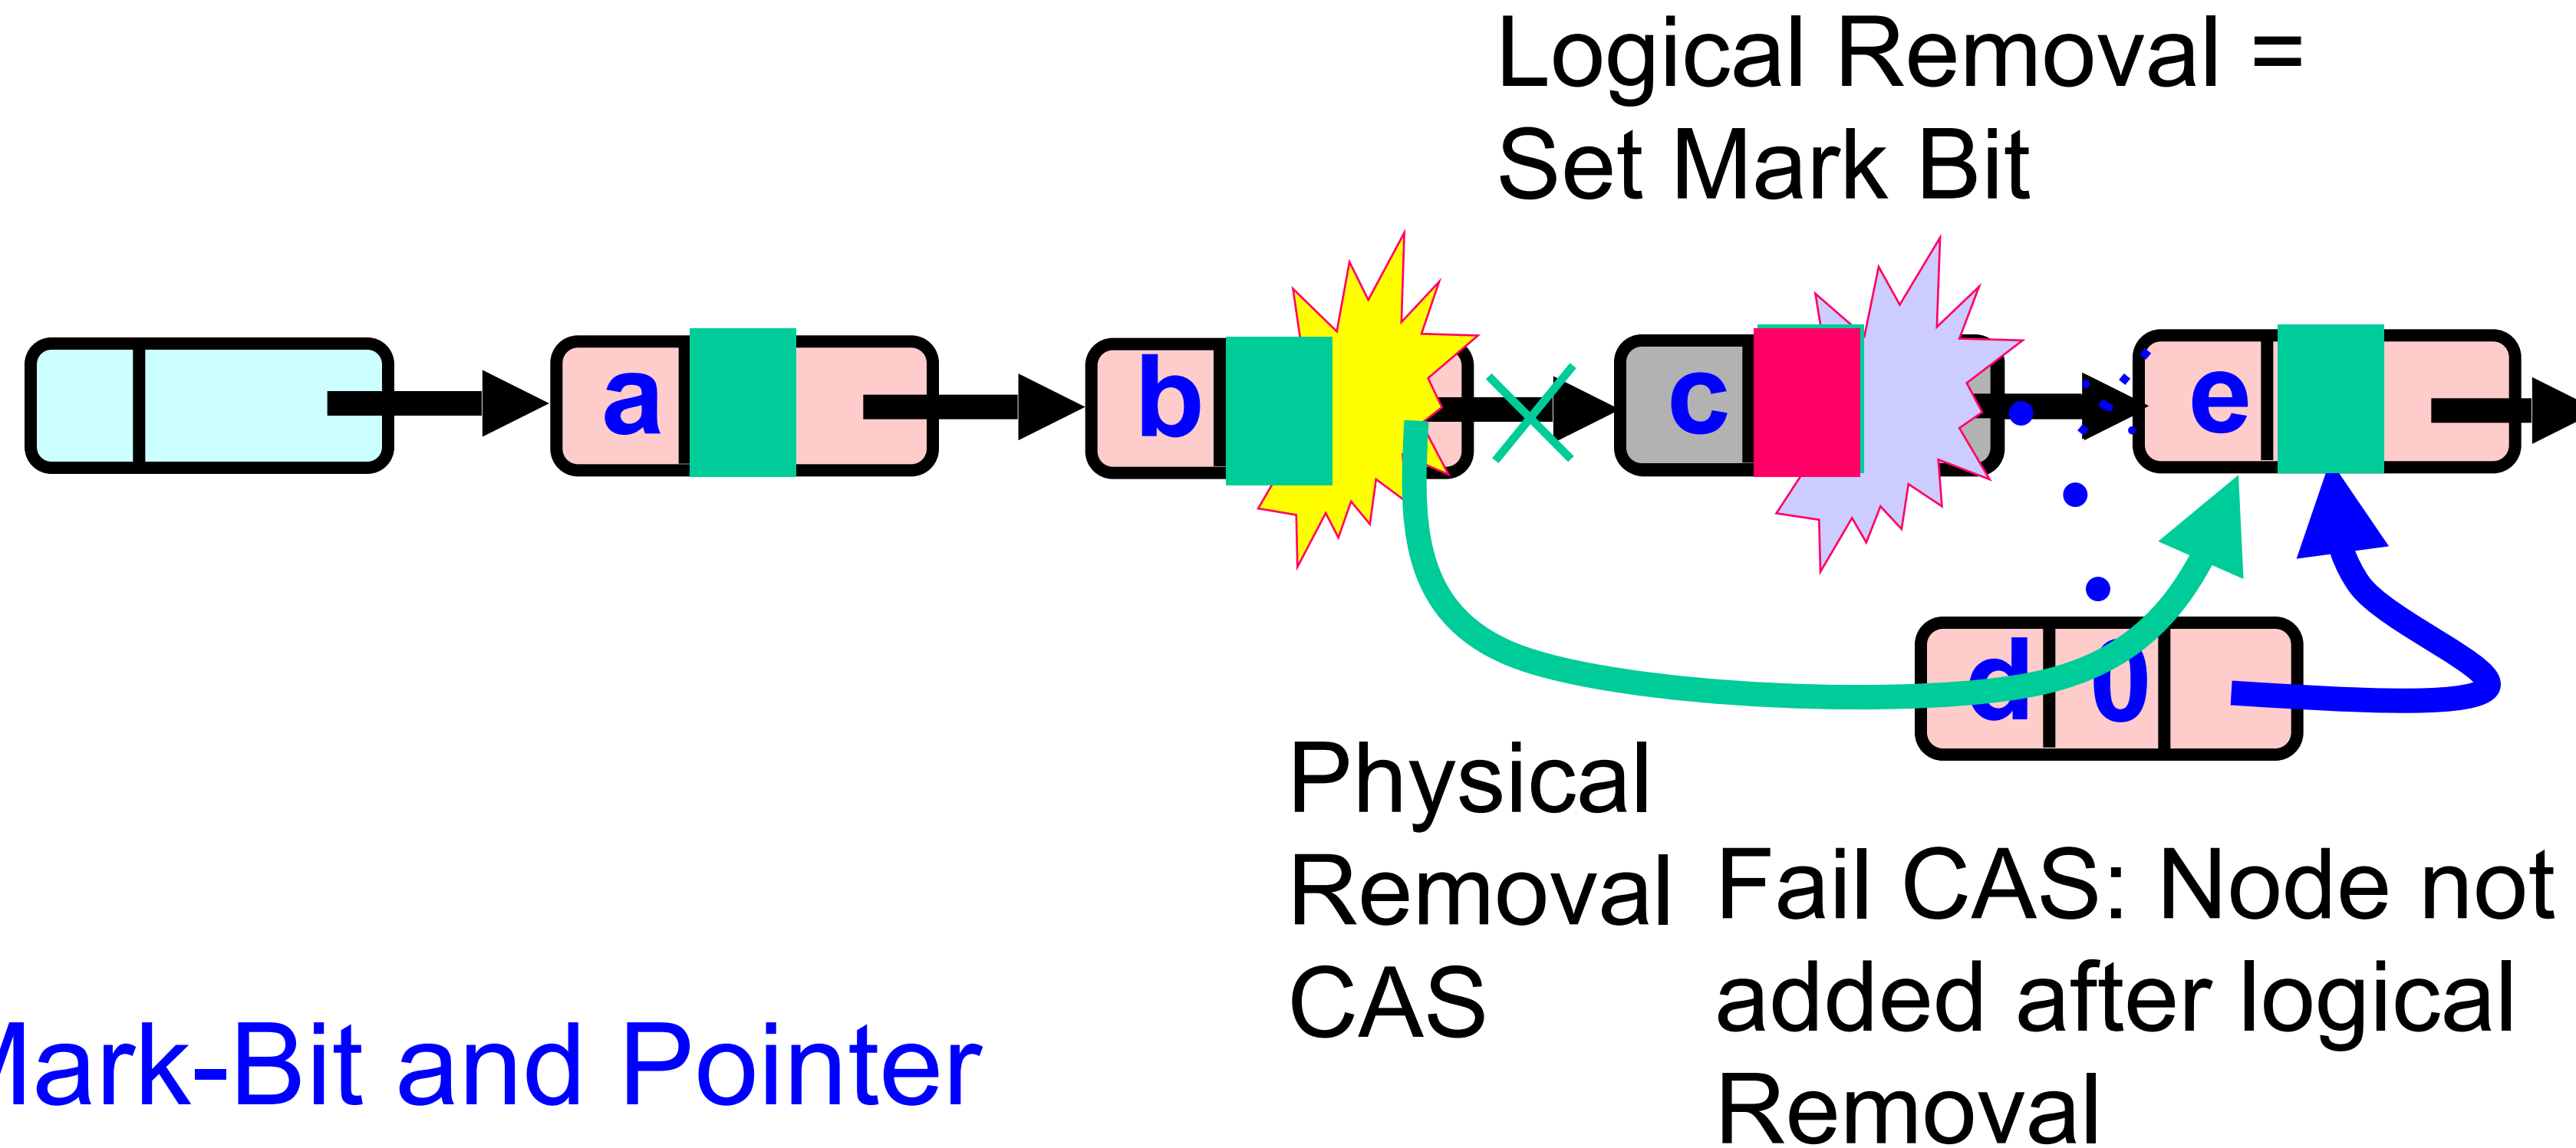
Use CAS to verify pointer  
is correct

Not enough!

# Problem...



# The Solution: Combine Bit and Pointer




Mark-Bit and Pointer  
are CASed together  
(**AtomicMarkableReference**)

# Solution

- Use AtomicMarkableReference
- Atomically
  - Swing reference and
  - Update flag
- Remove in two steps
  - Set mark bit in next field
  - Redirect predecessor's pointer

# Marking a Node


- **AtomicMarkableReference** class
  - Java.util.concurrent.atomic package




# Performance

- Different list-based set implementations
- 16-node machine
- Vary percentage of **contains()** calls

# High Contains Ratio

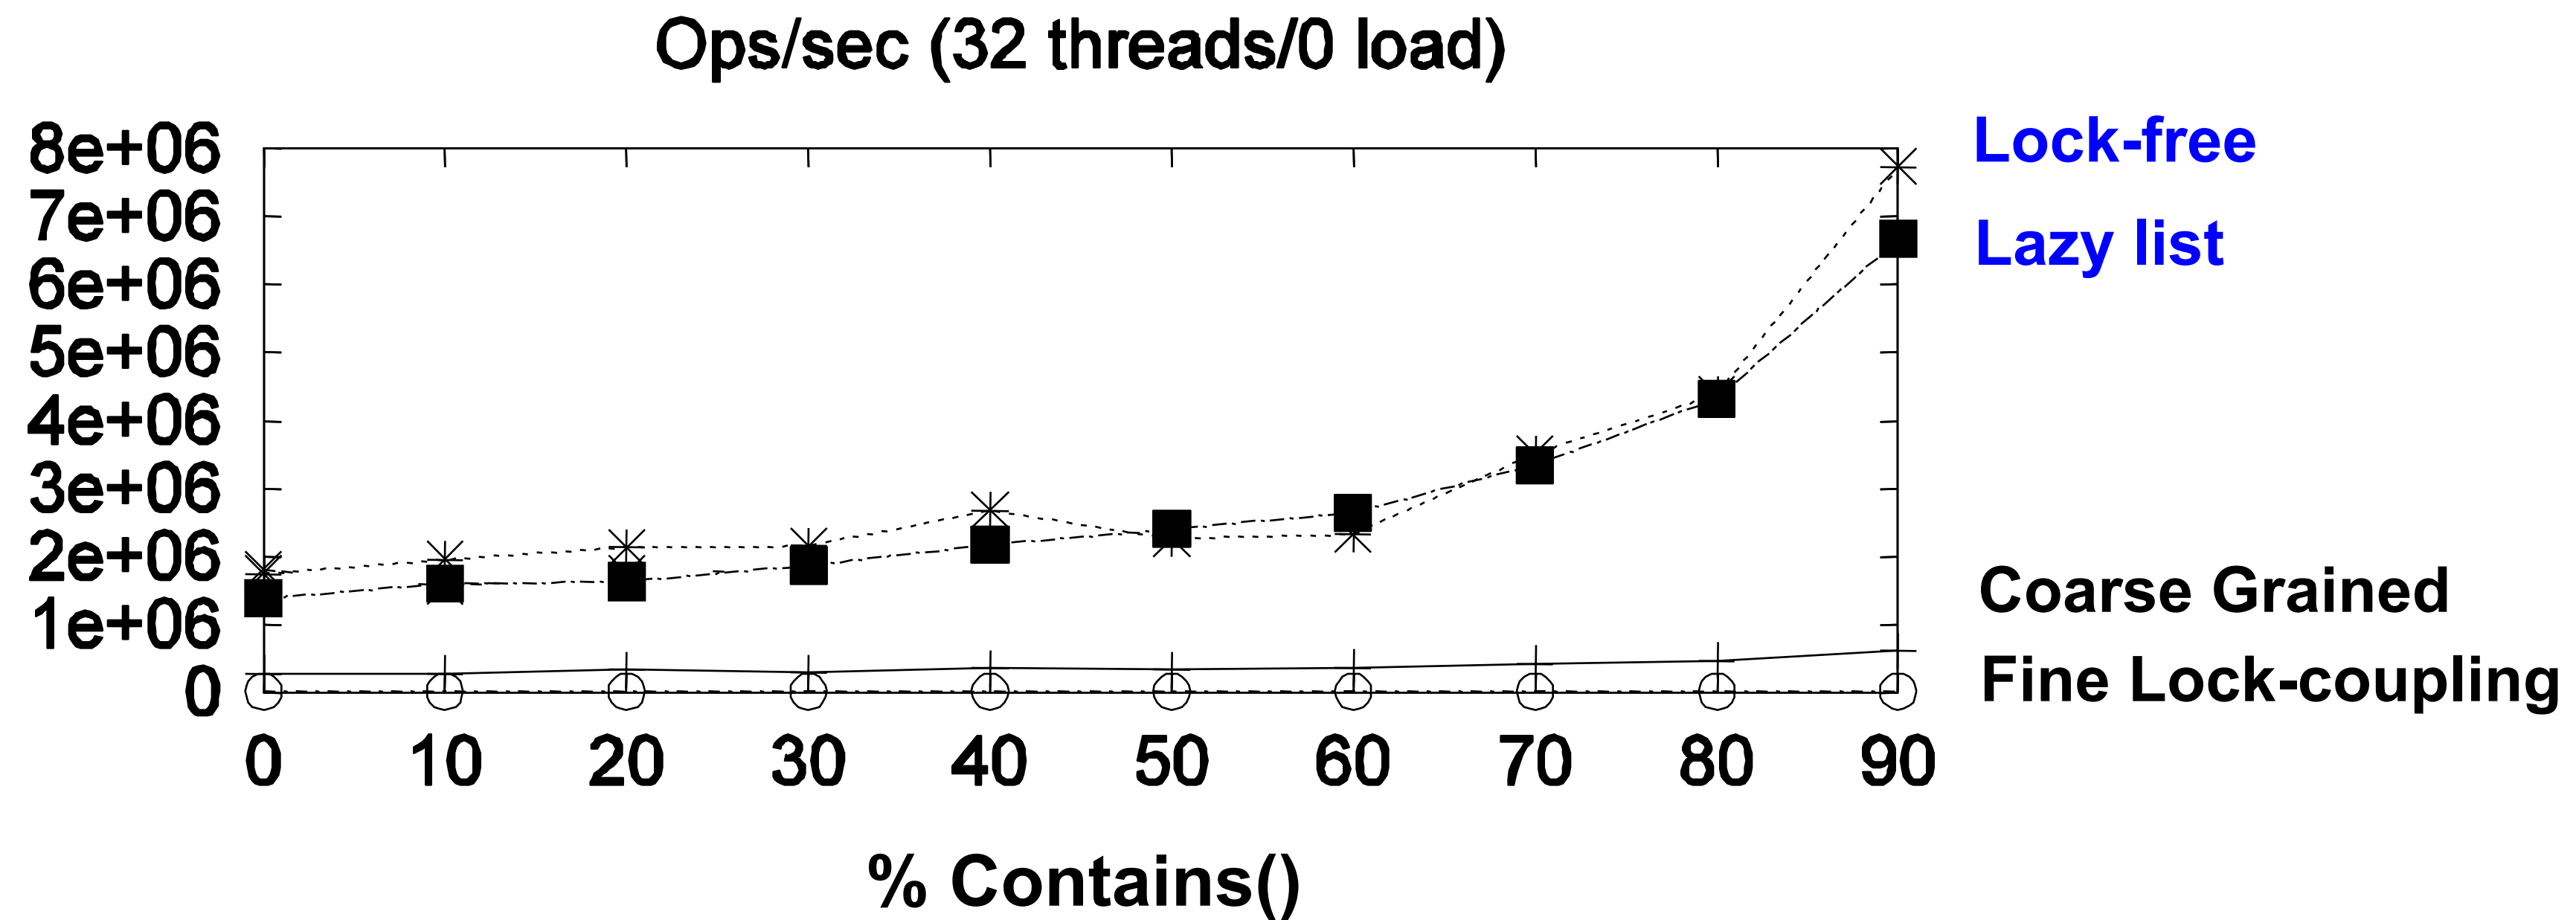


# Low Contains Ratio





# As Contains Ratio Increases



# Summary

- Coarse-grained locking
- Fine-grained locking (“hand-over-hand”)
- Optimistic synchronization
- Lazy synchronization
- Lock-free synchronization

# “To Lock or Not to Lock”

- Locking vs. Non-blocking:
  - Extremist views on both sides
  - Locking: longs waits
  - Non-blocking: long “clean-ups”
- The answer: nobler to compromise
  - Example: Lazy list combines blocking **add()** and **remove()** and a wait-free **contains()**
  - Remember: Blocking/non-blocking is a property of a method

This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](https://creativecommons.org/licenses/by-sa/2.5/).

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.