

YSC3248: Parallel, Concurrent and Distributed Programming



Concurrent Linked Lists
Part II

Last Lecture:


Fine-Grained Synchronization

- Instead of using a single lock ...
- Split object into
 - Independently-synchronized components
- Methods conflict when they access
 - The same component ...
 - At the same time


Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node



Removing a Node



Removing a Node



Remove method

```
def remove(item: T): Boolean = {  
    var pred, curr: Node = null  
    val key = item.hashCode  
  
    try { ... } finally {  
        curr.unlock()  
        pred.unlock()  
    }  
}
```

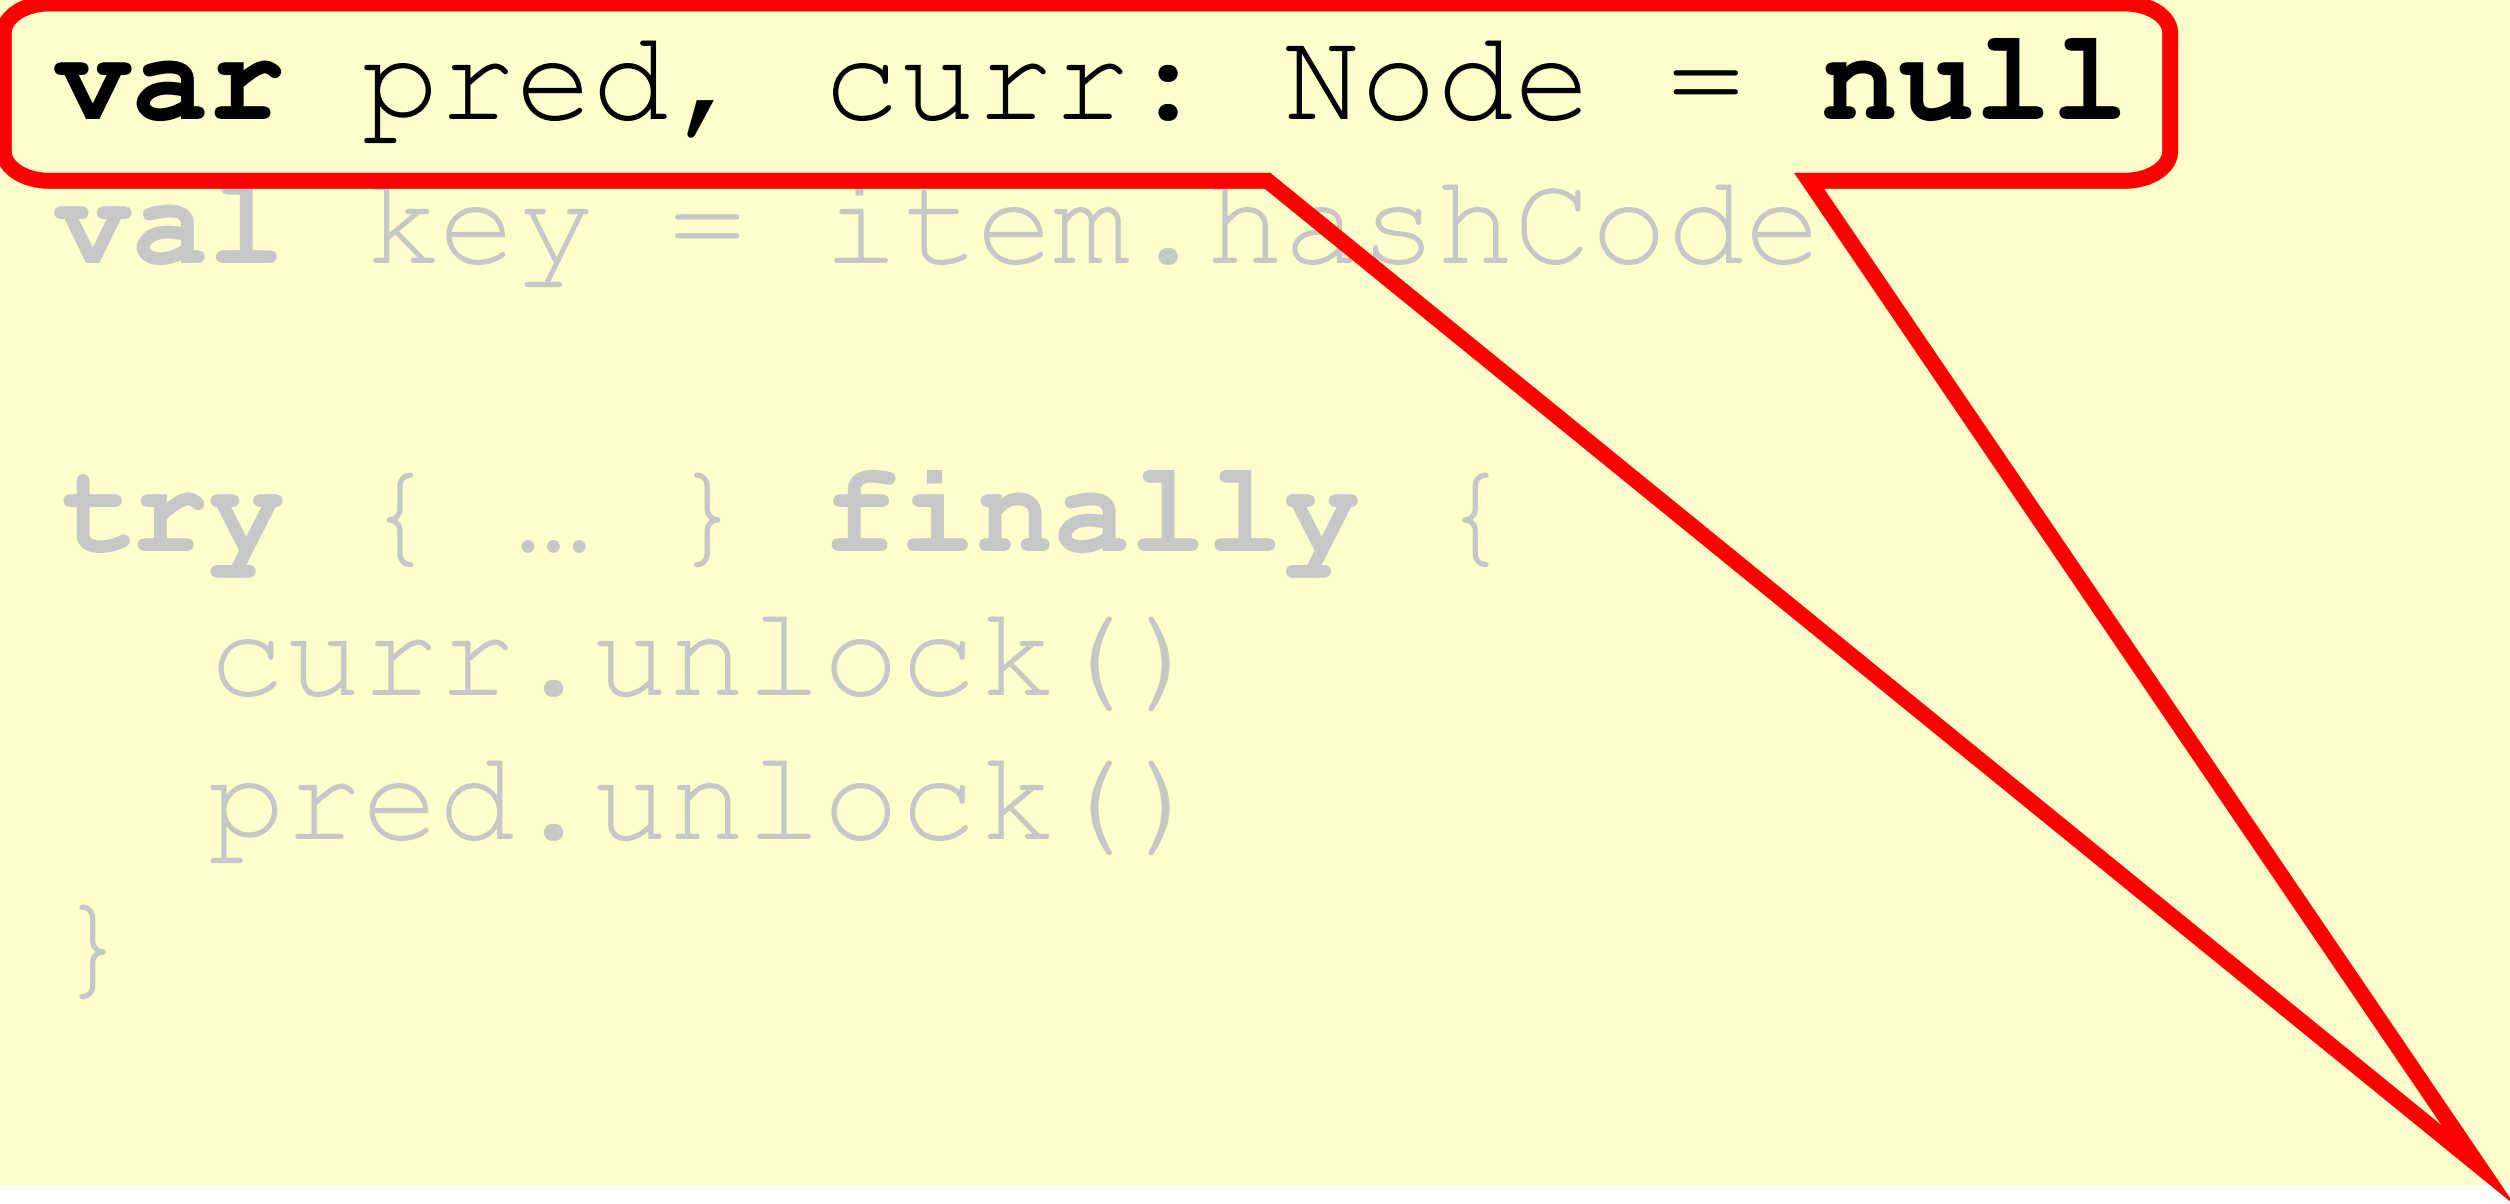
Remove method

```
def remove(item: T): Boolean = {  
    var pred, curr: Node = null  
    val key = item.hashCode  
  
    try { ... } finally {  
        curr.unlock()  
        pred.unlock()  
    }  
}
```

Key used to order node

Remove method

```
def remove(item: T): Boolean = {  
  var pred, curr: Node = null  
  val key = item.hashCode  
  
  try { ... } finally {  
    curr.unlock()  
    pred.unlock()  
  }  
}
```



Predecessor and current nodes

Remove method

```
def remove(item: T): Boolean = {  
  var pred, curr: Node = null  
  val key = item.hashCode  
  
  try { ... } finally {  
    curr.unlock()  
    pred.unlock()  
  }  
}
```

**Make sure
locks released**

Remove method

```
def remove(item: T): Boolean = {  
  var pred, curr: Node = null  
  val key = item.hashCode  
  
  try { ... } finally {  
    curr.unlock()  
    pred.unlock()  
  }  
}
```

Everything else


Remove method

```
try {  
    pred = head  
    pred.lock()  
    curr = pred.next  
    curr.lock()  
    ...  
} finally { ... }
```

Remove method

lock pred == head


```
try {  
  pred = head  
  pred.lock()  
  curr = pred.next  
  curr.lock()  
  ...  
} finally { ... }
```



Remove method

```
try {  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```


Lock current



Remove method


```
try {  
    pred = head  
    pred.lock()  
    curr = pred.next  
    curr.lock()  
    ...  
} finally { ... }
```

Traversing list



Remove: searching


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```


Search key range



Remove: searching

```
while (curr.key <= key) :  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```


**At start of each loop:
curr and pred locked**



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

If item found, remove node




Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

pred.unlock()


Unlock predecessor



Remove: searching

Only one node locked!


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

demote current


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

Find and lock new current


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = currNode  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

Loop invariant restored

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = currNode  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}
```

Otherwise, not present

return false

Why does this work?

- To remove node e
 - Must lock e
 - Must lock e 's predecessor
- Therefore, if you lock a node
 - It can't be removed
 - And neither can its successor

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

- **pred** reachable from head
- **curr** is pred.next
- So **curr.item** is in the set

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

**Linearization point if
item is present**

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

**Node locked, so no other
thread can remove it**

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}
```

return false

Item not present

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}
```

return false

- **pred** reachable from head
- **curr** is pred.next
- **pred.key** < key
- **key** < **curr.key**

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

Linearization point



Adding Nodes

- To add node e
 - Must lock predecessor
 - Must lock successor
- Neither can be deleted
 - (Is successor lock actually required?)

Same Abstraction Map

- $S(\text{head}) =$
 $\{ x \mid \text{there exists } a \text{ such that}$
 - $a \text{ reachable from head and}$
 - $a.\text{item} = x$ $\}$

Rep Invariant

- Easy to check that
 - tail always reachable from head
 - Nodes sorted, no duplicates


Drawbacks

- Better than coarse-grained lock
 - Threads can traverse in parallel
- Still not ideal
 - Long chain of acquire/release
 - Inefficient


Optimistic Synchronization

- Find nodes without locking
- Lock nodes
- Check that everything is OK


Optimistic: Traverse without Locking




Optimistic: Lock and Load




Optimistic: Lock and Load




What could go wrong?




What could go wrong?




What could go wrong?



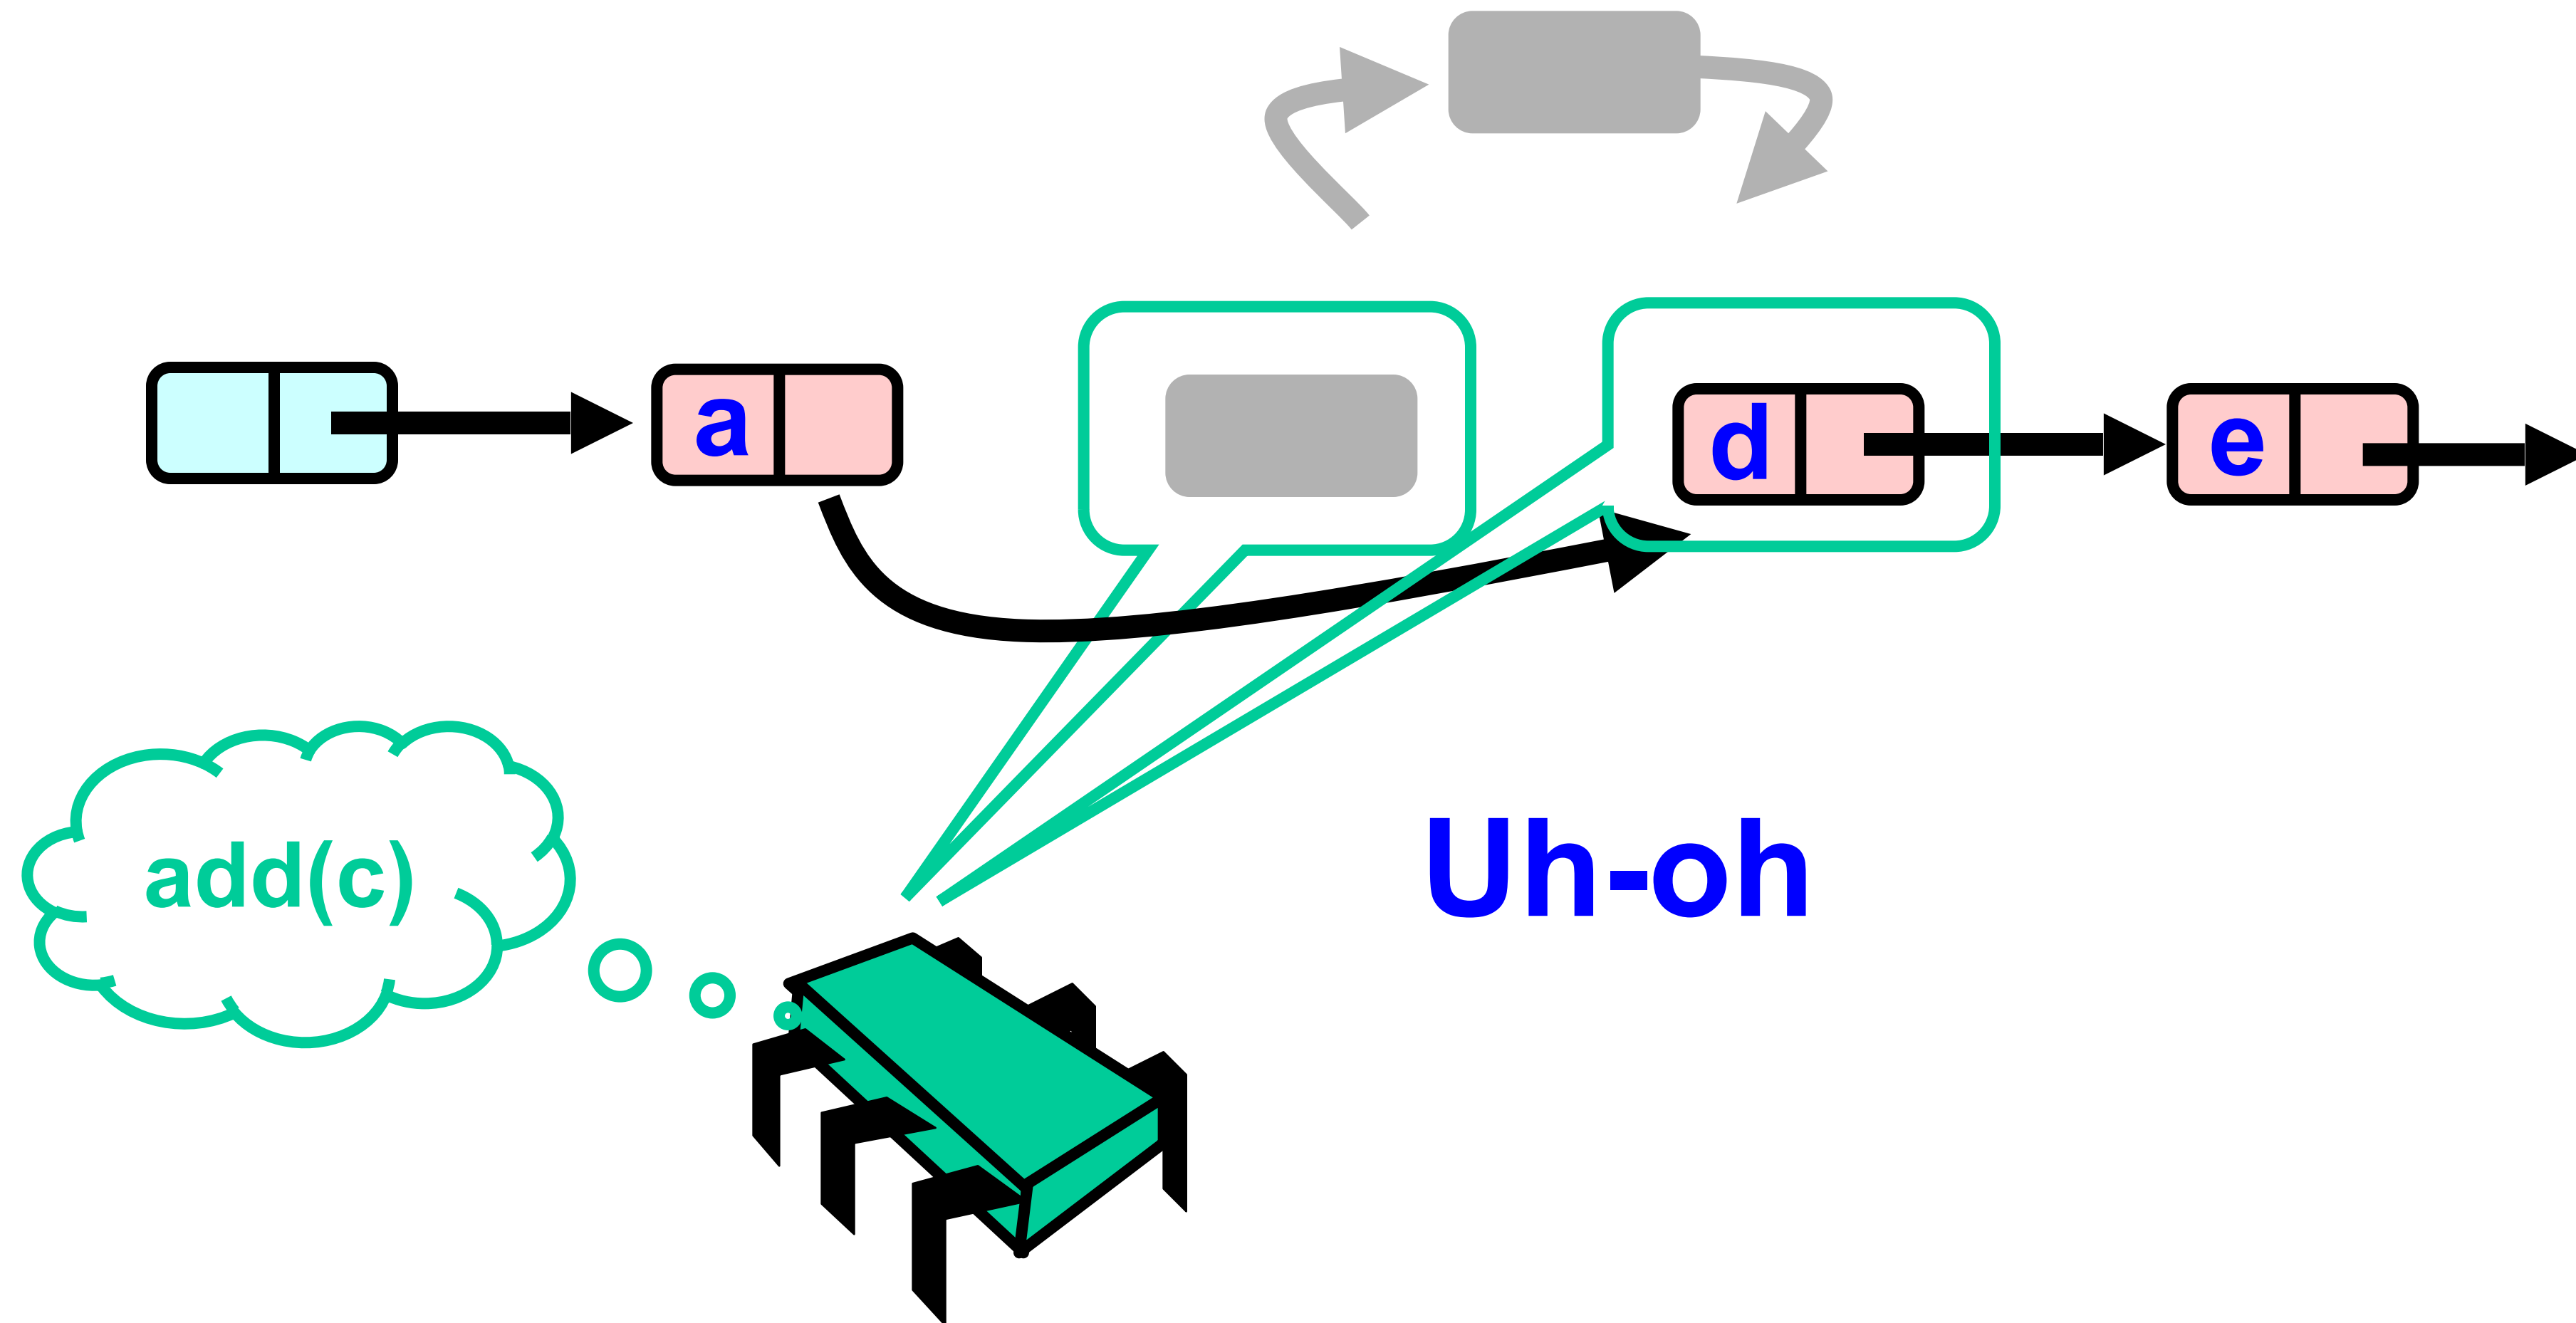
What could go wrong?



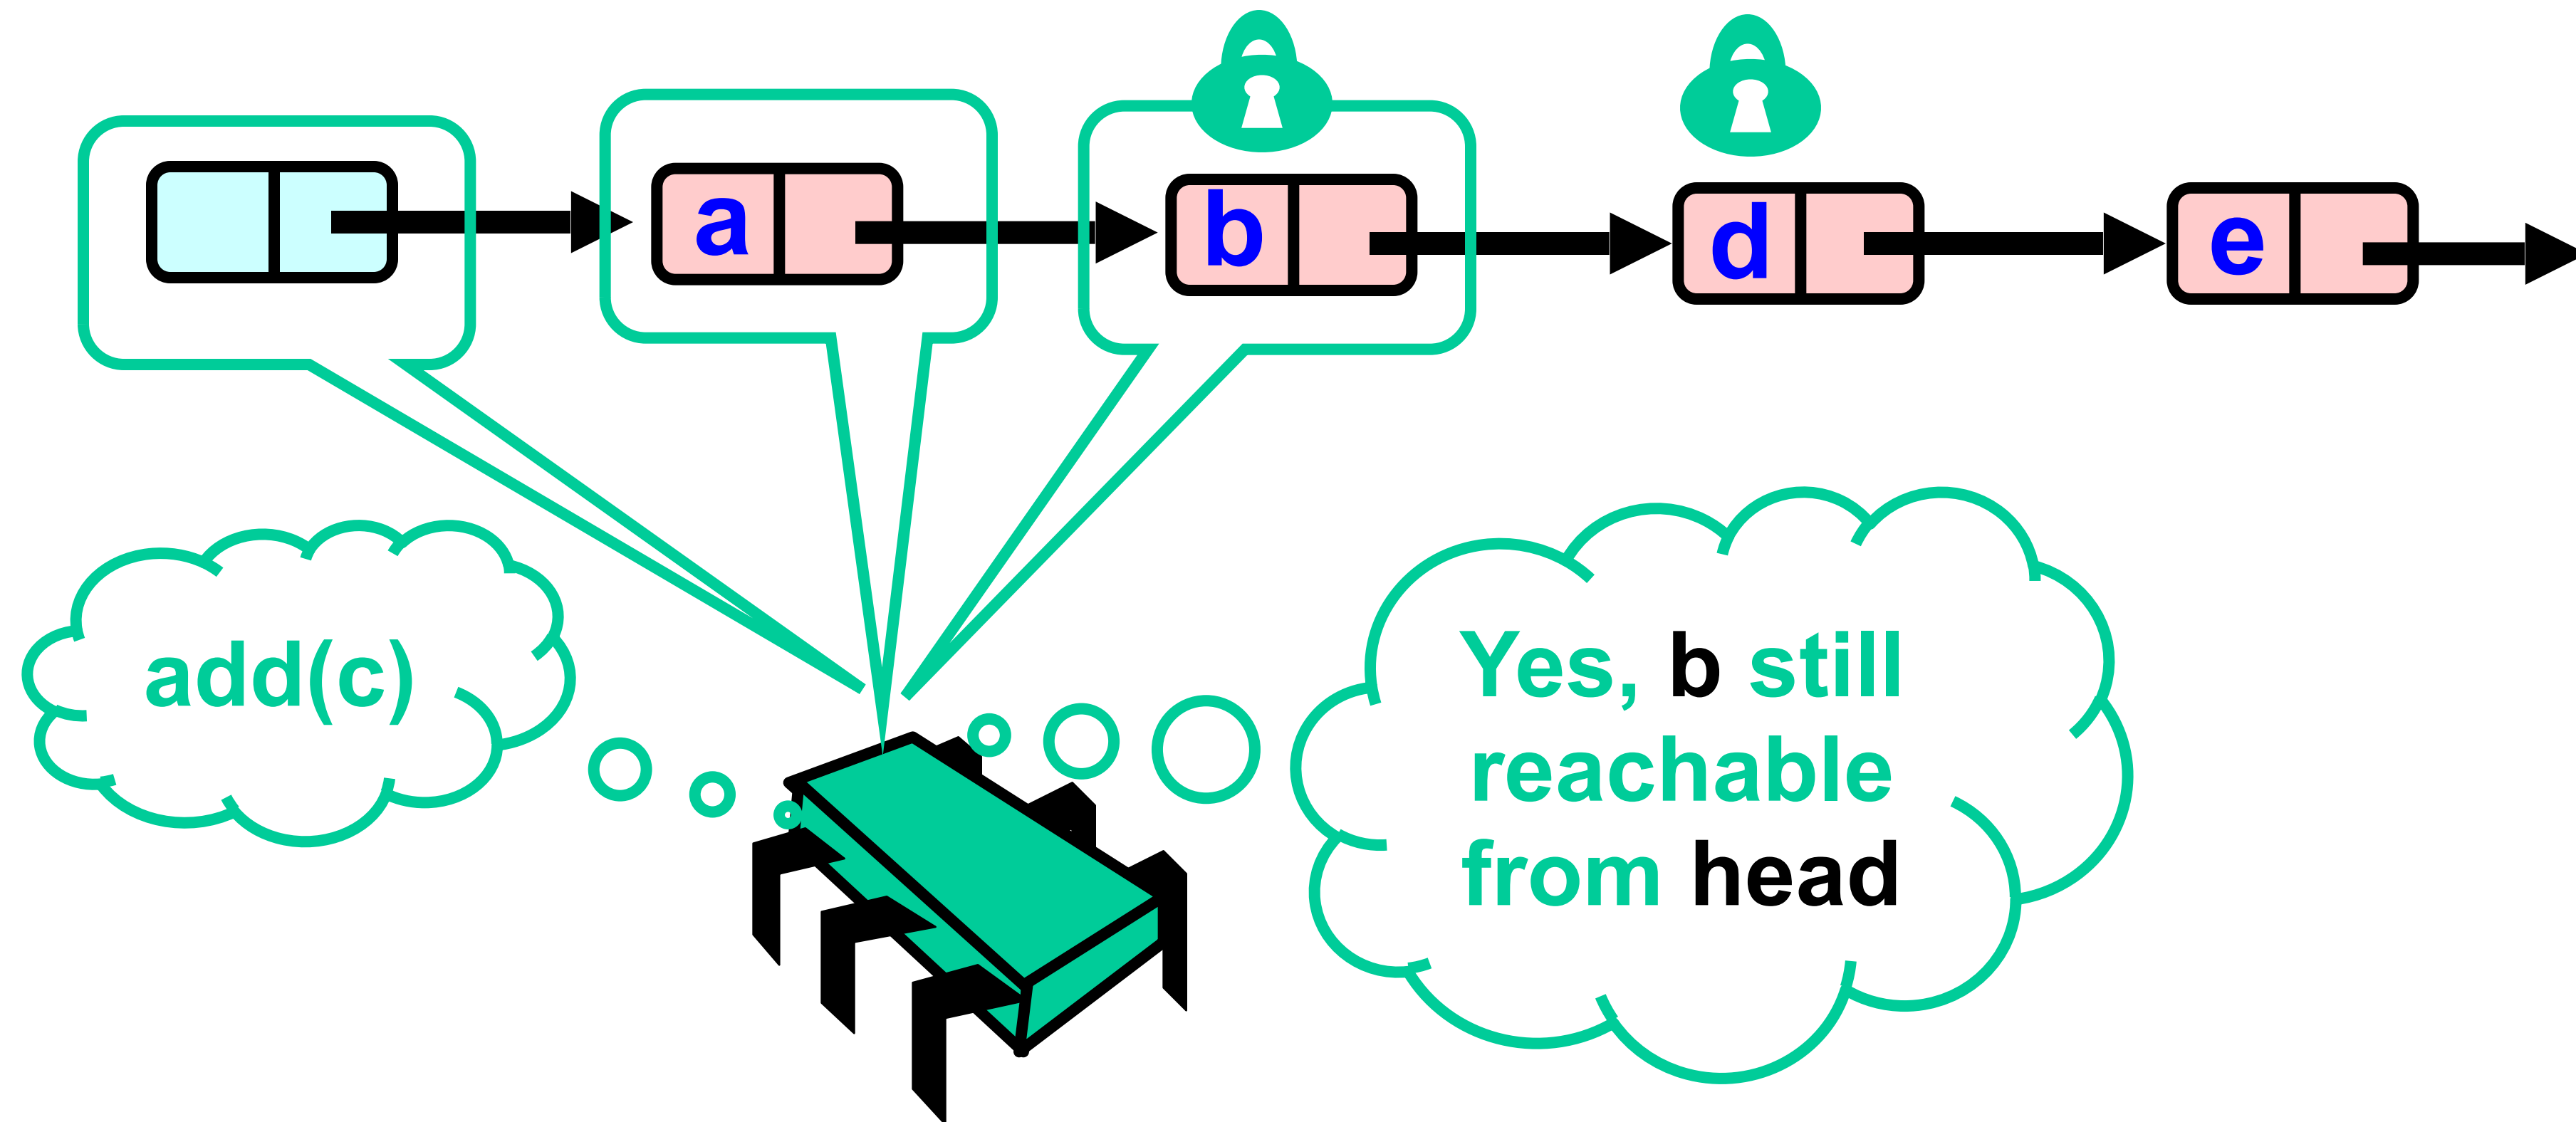
What could go wrong?




What could go wrong?




Validate – Part 1




What Else Could Go Wrong?




What Else Could Go Wrong?



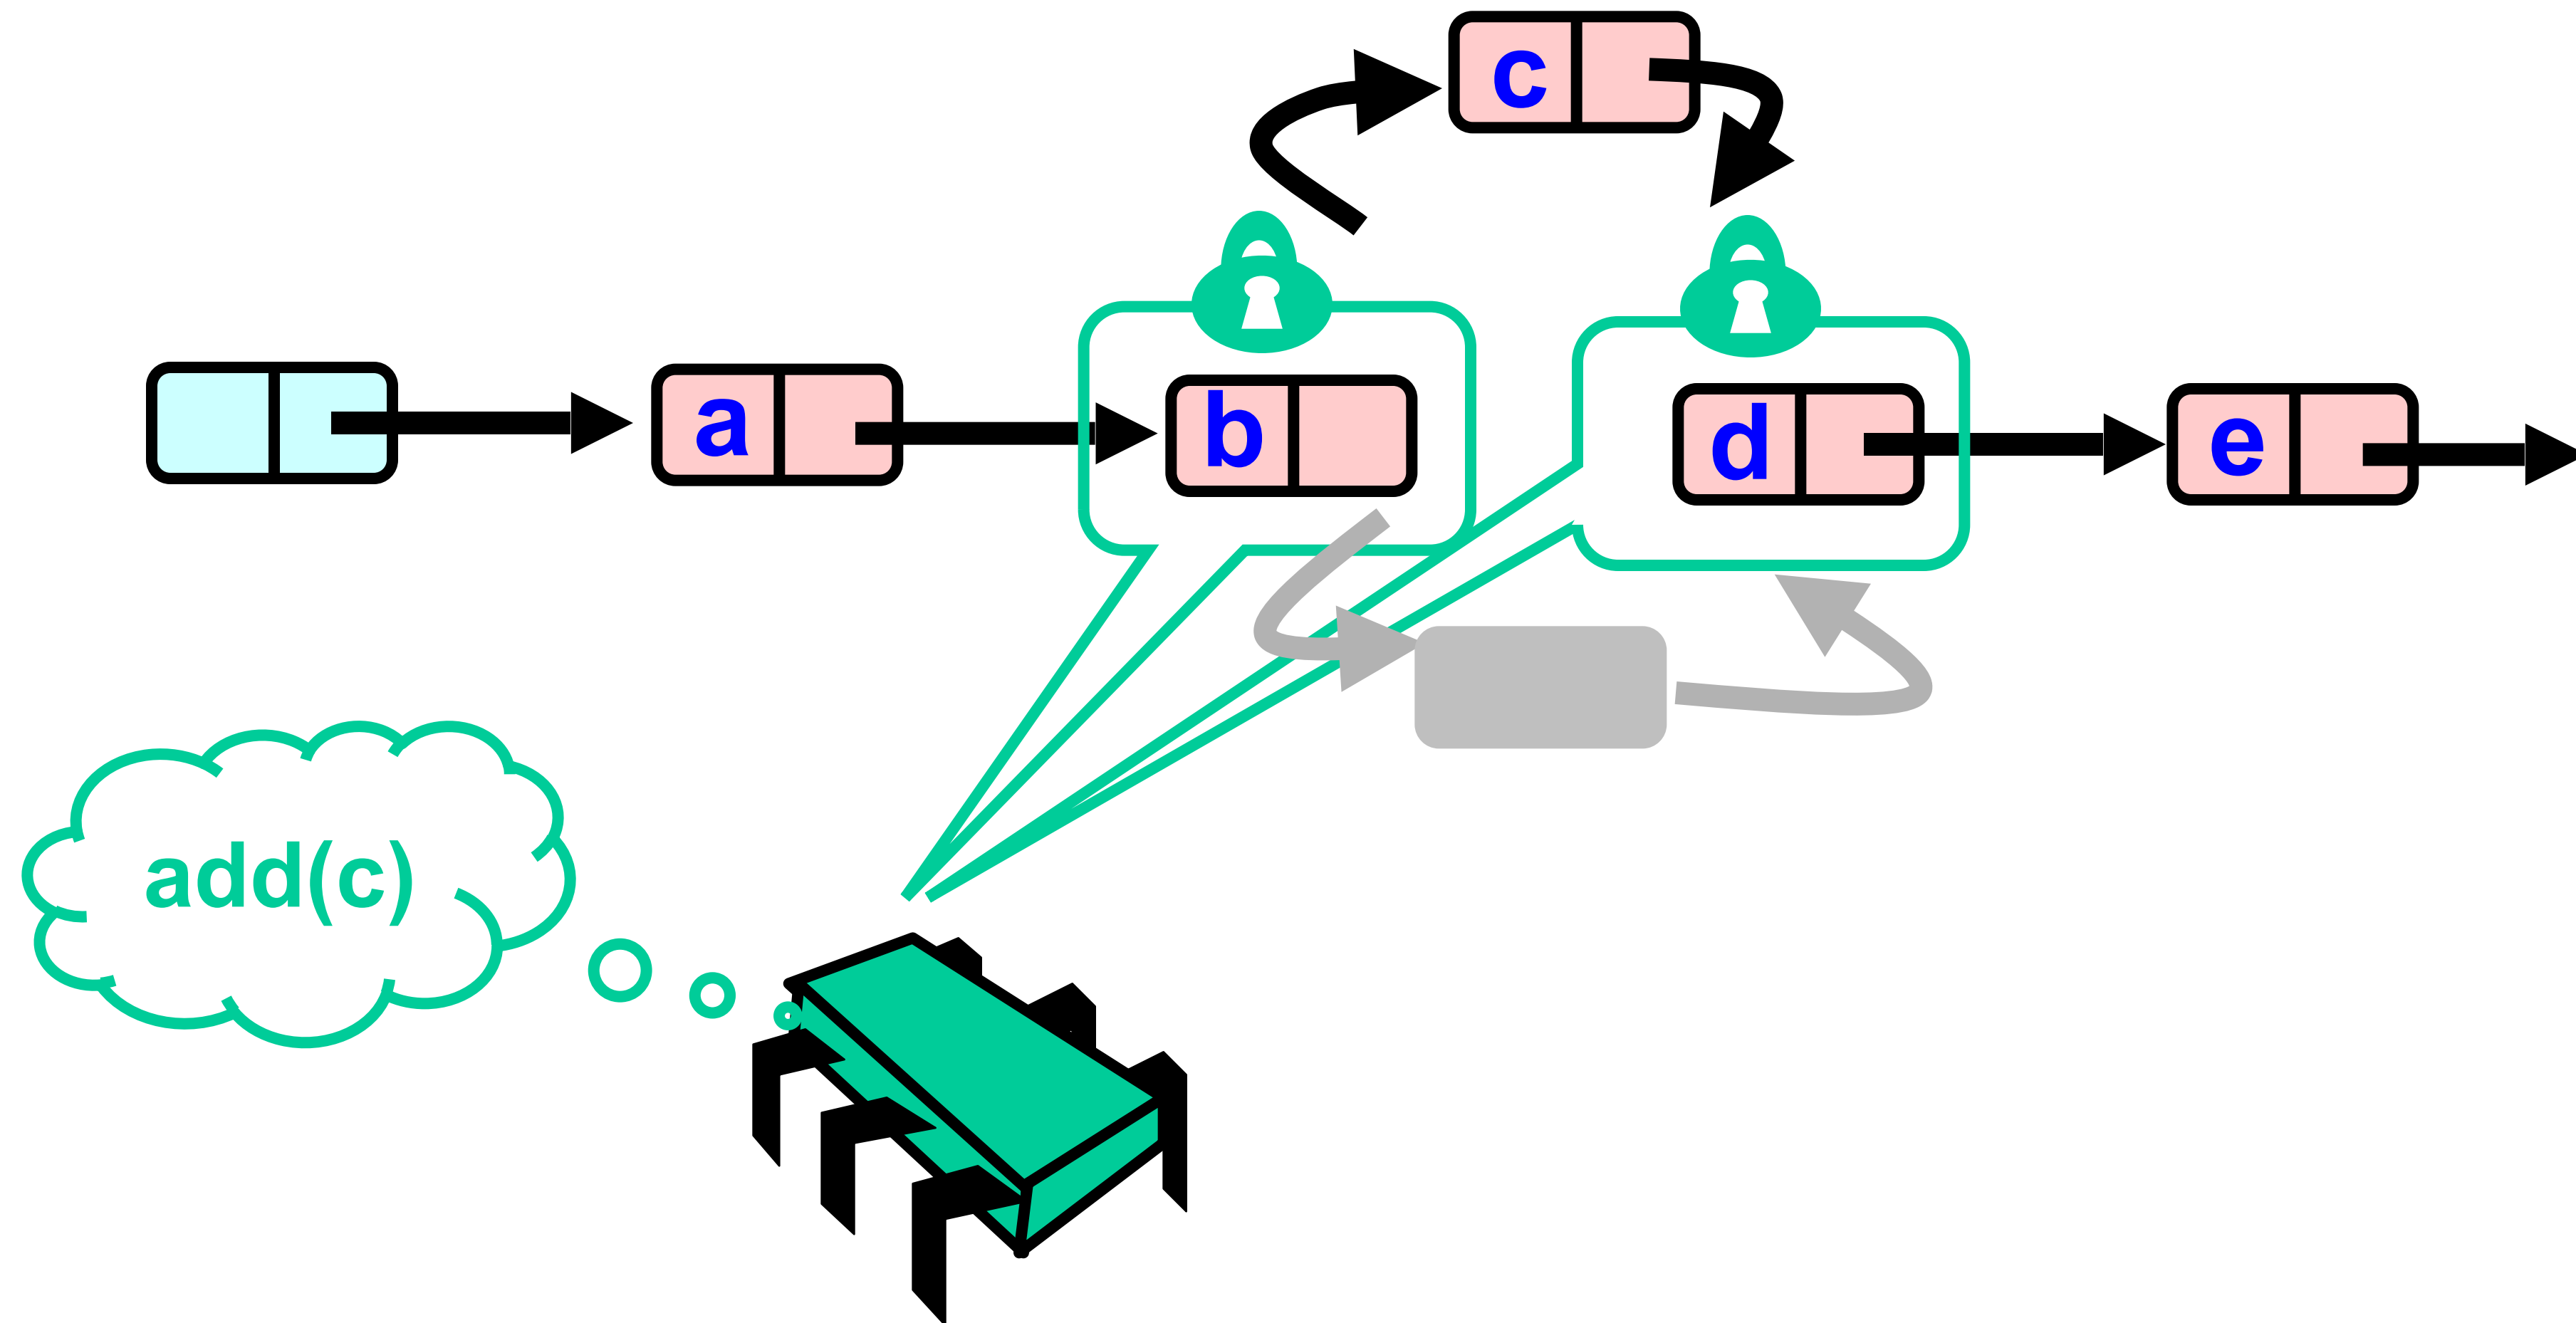
What Else Could Go Wrong?



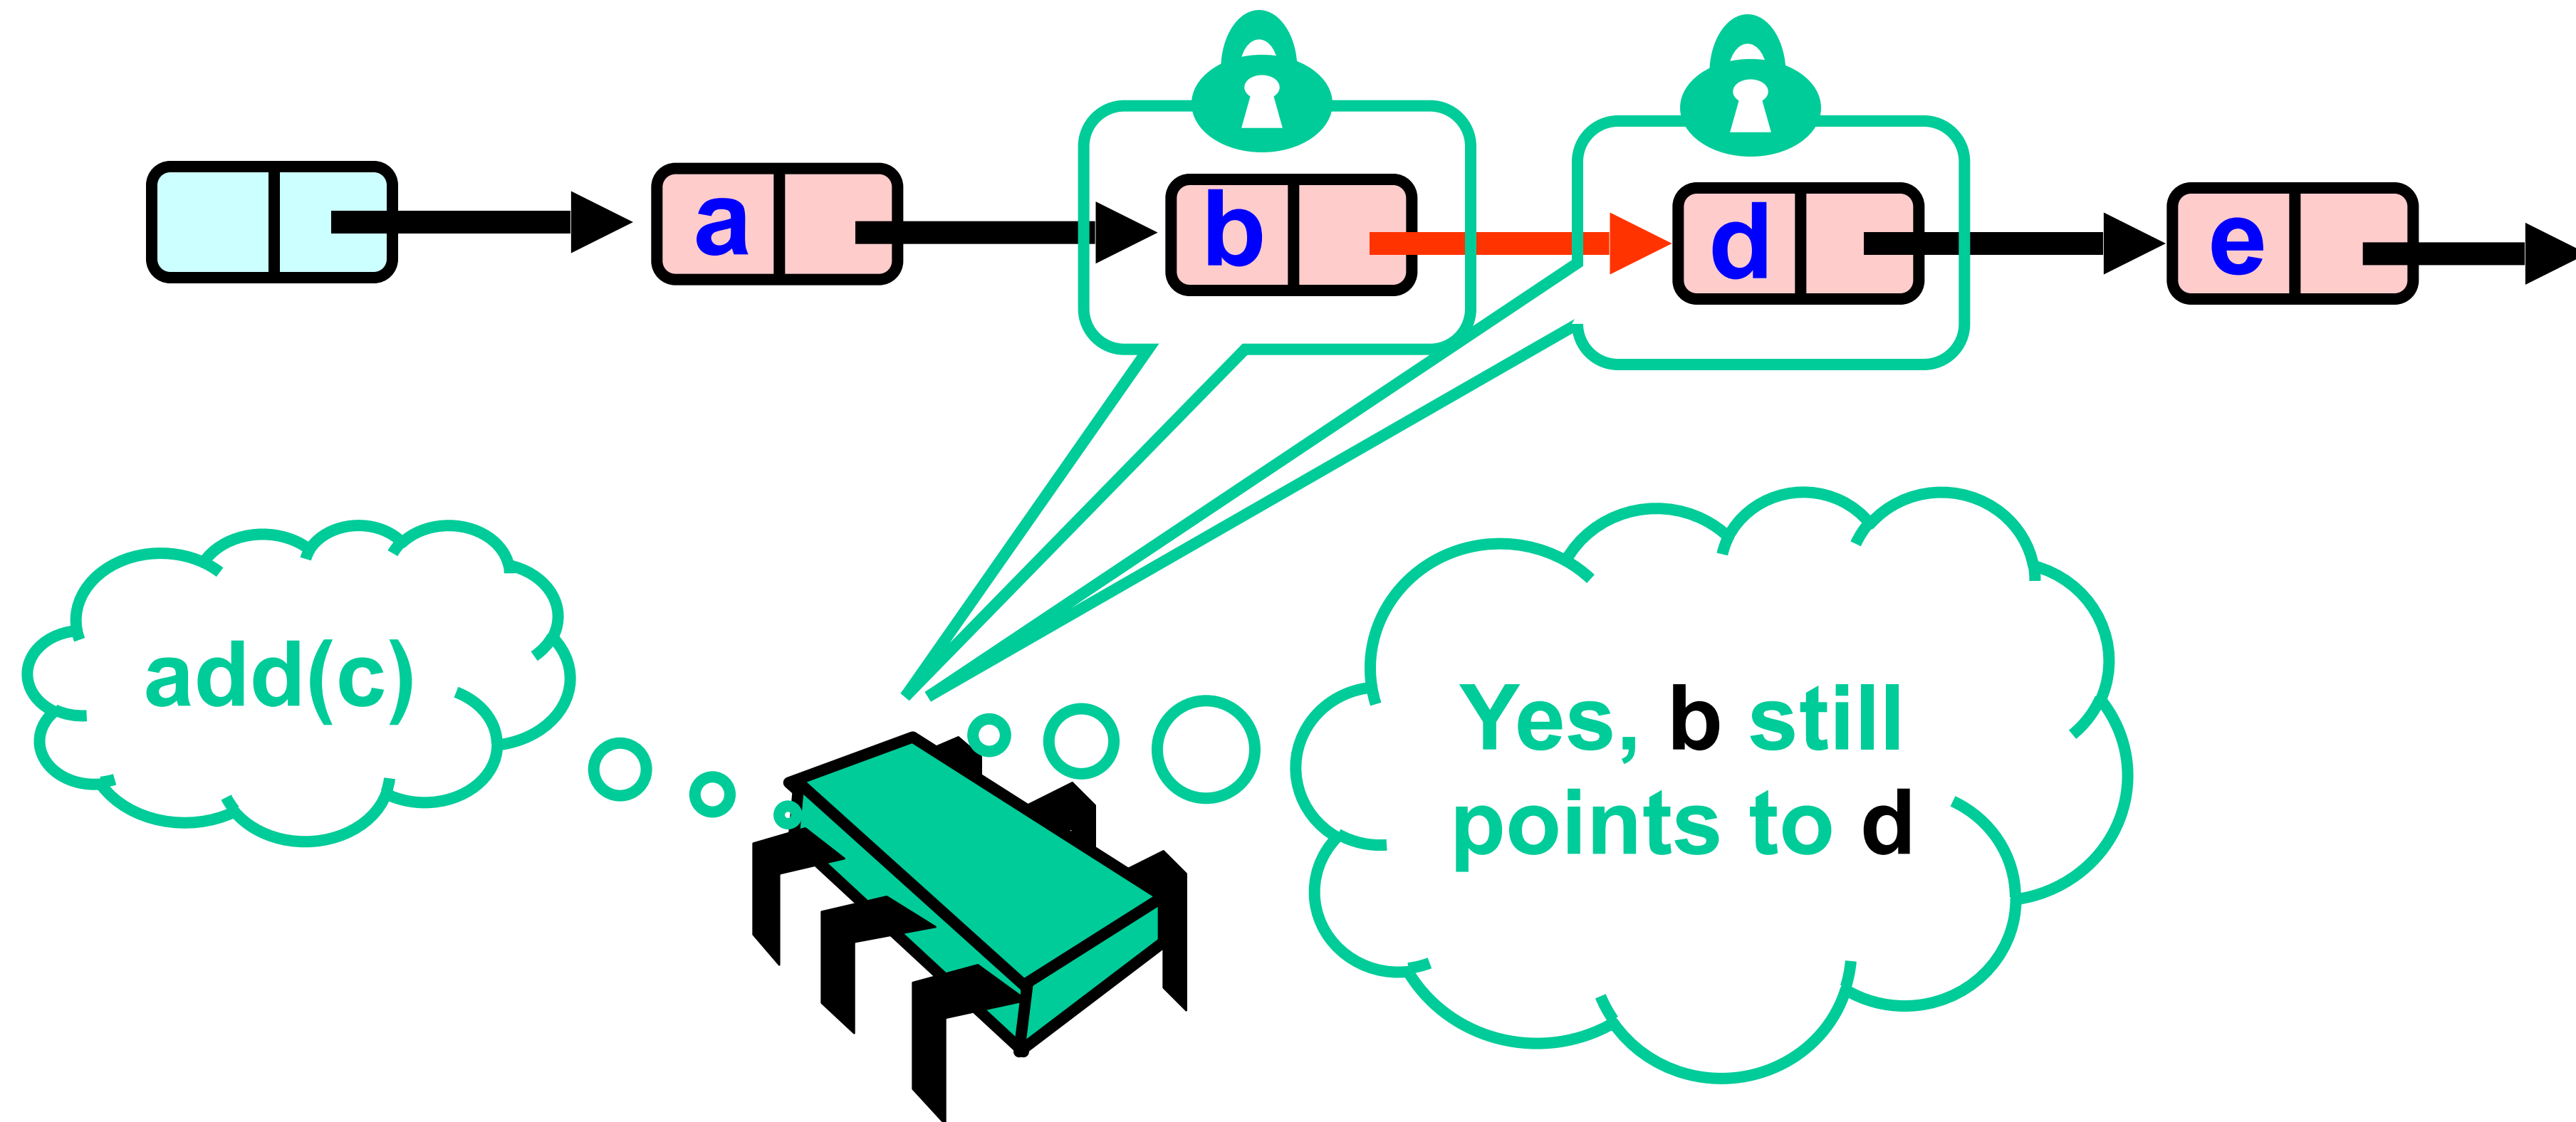
What Else Could Go Wrong?




What Else Could Go Wrong?



Validate Part 2 (while holding locks)



Optimistic: Linearization Point




Same Abstraction Map

- $S(\text{head}) =$
 $\{ x \mid \text{there exists } a \text{ such that}$
 - $a \text{ reachable from head and}$
 - $a.\text{item} = x$ $\}$

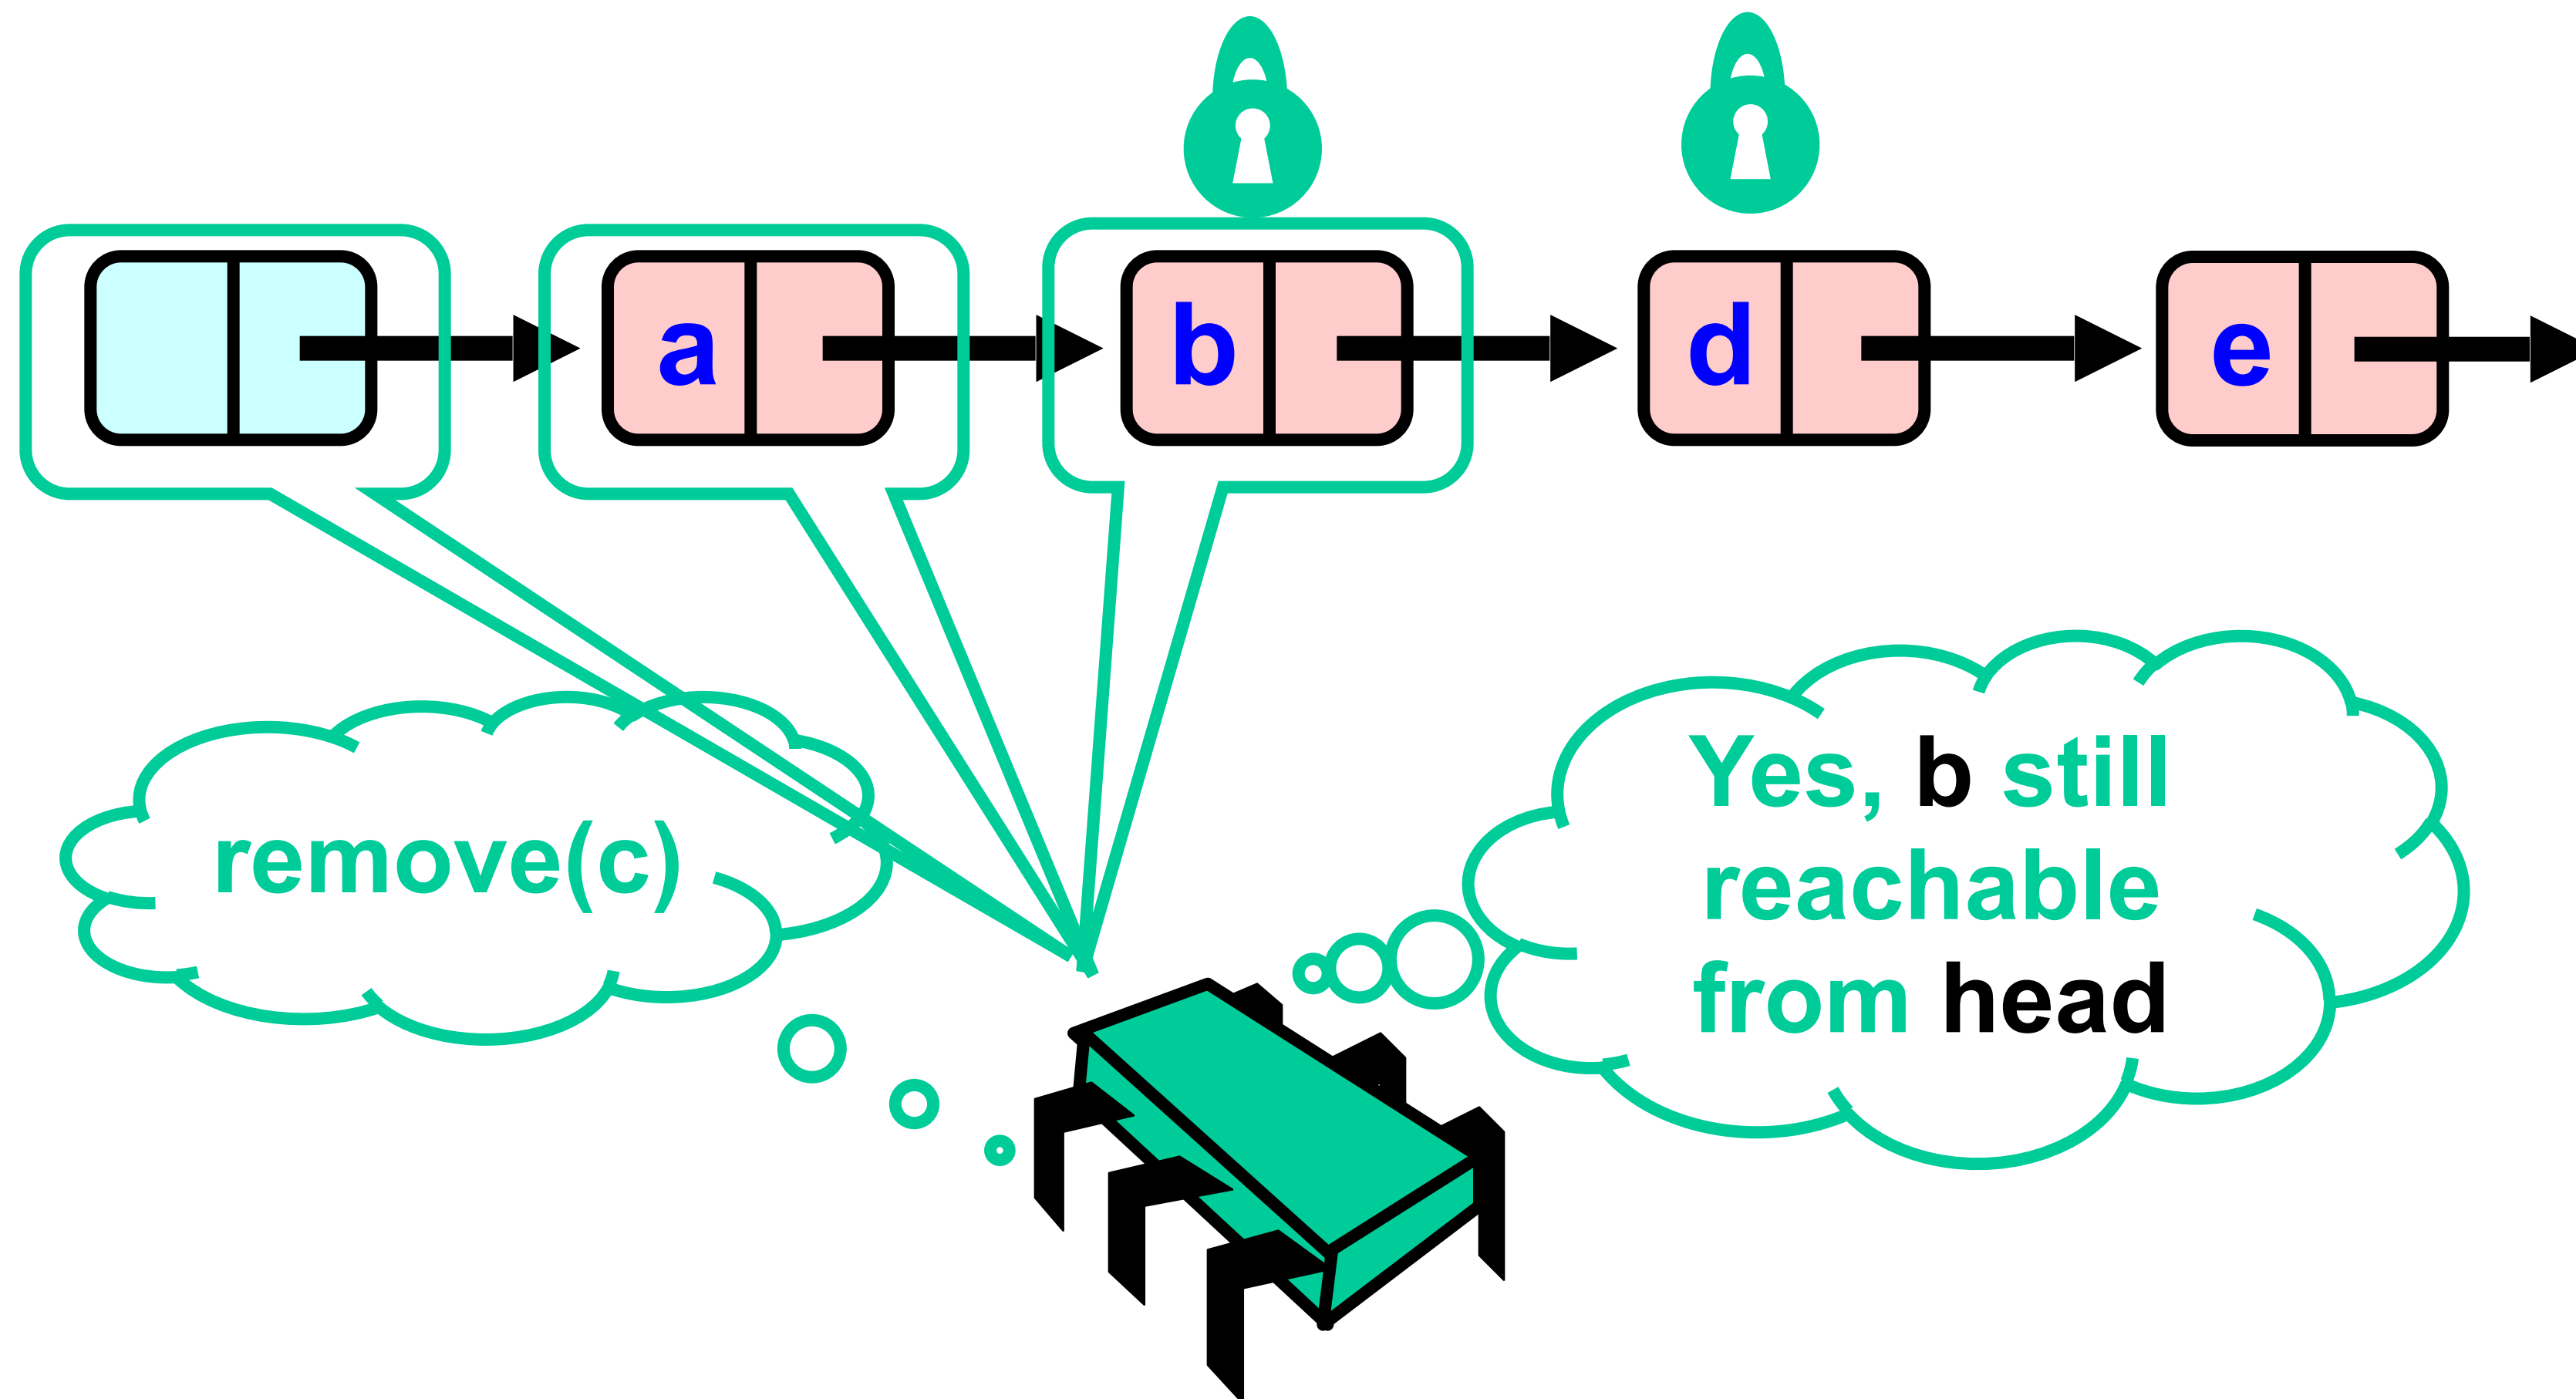
Invariants

- Careful: we may traverse deleted nodes
- But we establish properties by
 - Validation
 - After we lock target nodes

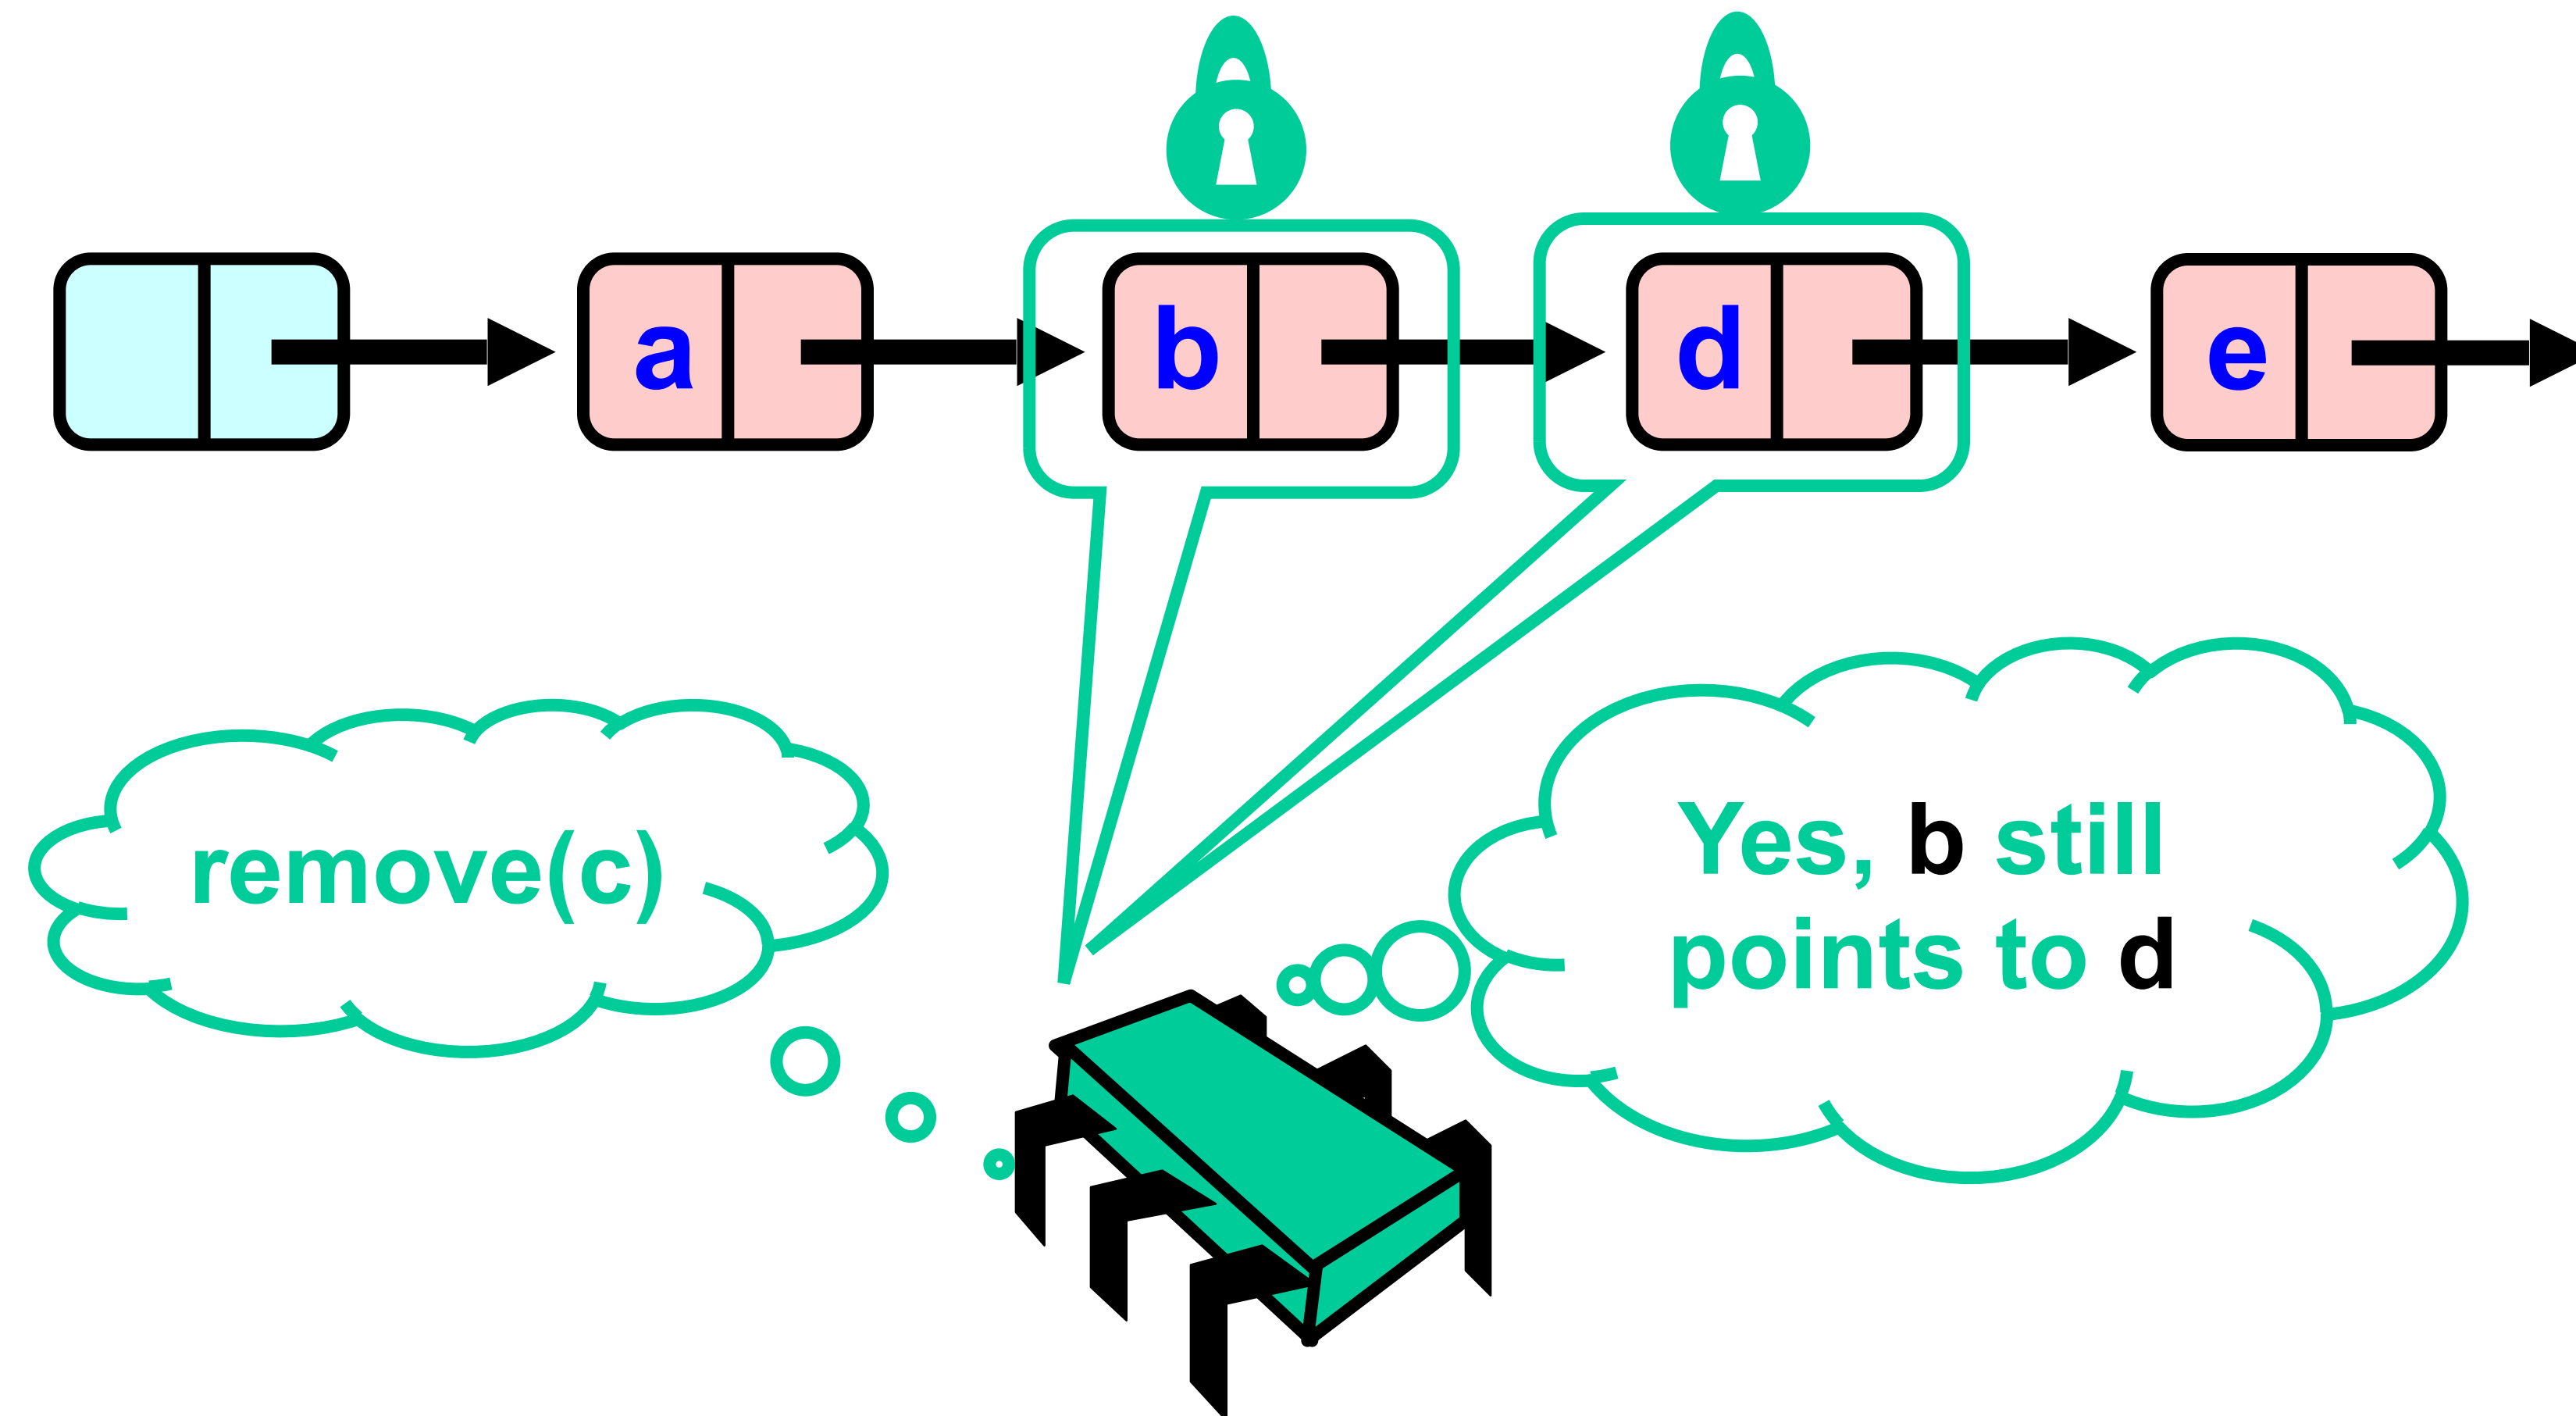
Unsuccessful Removal of c




Validate (1)



Validate (2)



OK Computer



Correctness

- If
 - Nodes **b** and **d** both locked
 - Node **b** still accessible
 - Node **d** still successor to **b**
- Then
 - Neither will be deleted
 - No thread can add **c** after **b**
 - OK to return false


Validation

```
def validate(pred: Node, curr: Node): Boolean = {  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

Validation

```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```


**Predecessor &
current nodes**




Validation

```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```

**Start at the
beginning**




Validation



```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```

Search range of keys


Validation



```
def validate(pred: Node, curr: Node):  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

Predecessor reachable

Validation




```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```

Is current node next?

Validation

Otherwise move on


```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```



Validation


Predecessor not reachable

```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```




Remove: searching

```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```



Remove: searching


```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```



Search key

Remove: searching


```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```



**Loop until no synchronization conflict
(see the code further)**

Remove: searching

```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```




Examine predecessor and current nodes

Remove: searching

```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }
```

Search by key



On Exit from While-True-Loop

- If item is present
 - curr holds item
 - pred just before curr
- If item is absent
 - curr has first higher key
 - pred just before curr
- Assuming no synchronization problems

Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) { // present in list
            pred.next = curr.next
            return true
        } else { // not present in list
            return false
        }
    }
} finally { // always unlock
    pred.unlock(); curr.unlock()
}
```

Remove Method

```
pred.lock(); curr.lock()
try {
    if (!validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally { // always unlock
    pred.unlock(); curr.unlock()
}
```

Always unlock

Remove Method

```
pred.lock(); curr.lock()
```

```
try {  
    if (validate(pred, curr)) {  
        if (curr.key == key) {  
            pred.next = curr.next  
            return true  
        } else {  
            return false  
        }  
    }  
} finally {  
    pred.unlock(); curr.unlock()  
}
```

Lock both nodes

Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally {
    pred.unlock(); curr.unlock()
}
```

Check for synchronization conflicts

Remove Method

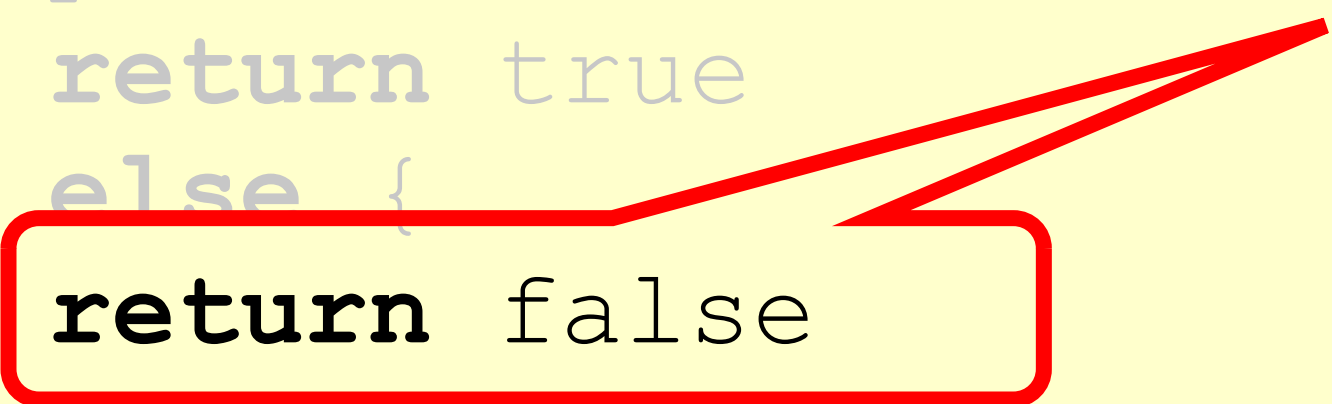
```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally {
    pred.unlock(); curr.unlock()
}
```

**target found, remove
node**

Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally {
    pred.unlock(); curr.unlock()
}
```

target not found



Optimistic List

- Limited hot-spots
 - Targets of `add()`, `remove()`, `contains()`
 - No contention on traversals
- Moreover
 - Traversals are wait-free
 - Food for thought ...

So Far, So Good

- Much less lock acquisition/release
 - Performance
 - Concurrency
- Problems
 - Need to traverse list twice
 - `contains()` method acquires locks

Evaluation

- Optimistic is effective if
 - cost of scanning twice without locks is less than
 - cost of scanning once with locks
- Drawback
 - `contains()` acquires locks
 - 90% of calls in many apps


Lazy List

- Like optimistic, except
 - Scan once
 - `contains(x)` never locks ...
- Key insight
 - Removing nodes causes trouble
 - Do it “lazily”


Lazy List

- **remove ()**
 - Scans list (as before)
 - Locks predecessor & current (as before)
- Logical delete
 - Marks current node as removed (new!)
- Physical delete
 - Redirects predecessor's next (as before)


Lazy Removal




Lazy Removal




Lazy Removal



Lazy Removal



Lazy Removal




Lazy List

- All Methods
 - Scan through locked and marked nodes
 - Removing a node doesn't slow down other method calls ...
- Must still lock `pred` and `curr` nodes.


Validation

- No need to rescan list!
- Check that `pred` is not marked
- Check that `curr` is not marked
- Check that `pred` points to `curr`


Business as Usual




Business as Usual




Business as Usual




Business as Usual




Business as Usual




Business as Usual




Business as Usual



Business as Usual



Business as Usual



New Abstraction Map

- $S(\text{head}) =$
 $\{ x \mid \text{there exists node } a \text{ such that}$
 - a reachable from head and
 - $a.\text{item} = x$ and
 - a is unmarked $\}$

Invariant

- If not marked then item in the set
- and is reachable from head
- and if not yet traversed it is reachable from pred

Validation

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```

List Validate Method

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```

**Predecessor not
Logically removed**

List Validate Method

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```



**Current not
Logically removed**

List Validate Method

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```

**Predecessor still
Points to current**

Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true  
            pred.next = curr.next  
            return true;  
        } else {  
            return false  
        }  
    } finally {  
        pred.unlock()  
        curr.unlock()  
    }  
}
```


Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true  
            pred.next = curr.next  
            return true  
        } else {  
            return false  
        }  
    }  
    finally {  
        pred.unlock()  
        curr.unlock()  
    }  
}
```

Validate as before

Remove

```
try {  
    pred.lock(); curr.lock();  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true;  
            pred.next = curr.next;  
            return true;  
        } else {  
            return false;  
        }  
    }  
} finally {  
    pred.unlock();  
    curr.unlock();  
}
```

Key found

Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true;  
            pred.next = curr.next  
            return true  
        } else {  
            return false  
        }  
    } finally {  
        pred.unlock()  
        curr.unlock()  
    }  
}
```

Logical remove

Remove

```
try {
    pred.lock(); curr.lock()
    if (validate(pred, curr) {
        if (curr.key == key) {
            curr.marked = true
            pred.next = curr.next;
            return true
        } else {
            return false
        }
    } finally {
        pred.unlock()
        curr.unlock()
    }
}
```

physical remove

Contains

```
def contains(item: T) = {  
    val key = item.hashCode  
    var curr = this.head  
    while (curr.key < key) curr = curr.next  
    curr.key == key && !curr.marked  
}
```

Contains

```
def contains(item: T) = {  
    val key = item.hashCode  
    var curr = this.head  
    while (curr.key < key) curr = curr.next  
    curr.key == key && !curr.marked  
}
```

Start at the head

Contains

```
def contains(item: T) = {  
    val key = item.hashCode  
    var curr = this.head  
    while (curr.key < key) curr = curr.next  
    curr.key == key && !curr.marked  
}
```

Search key range

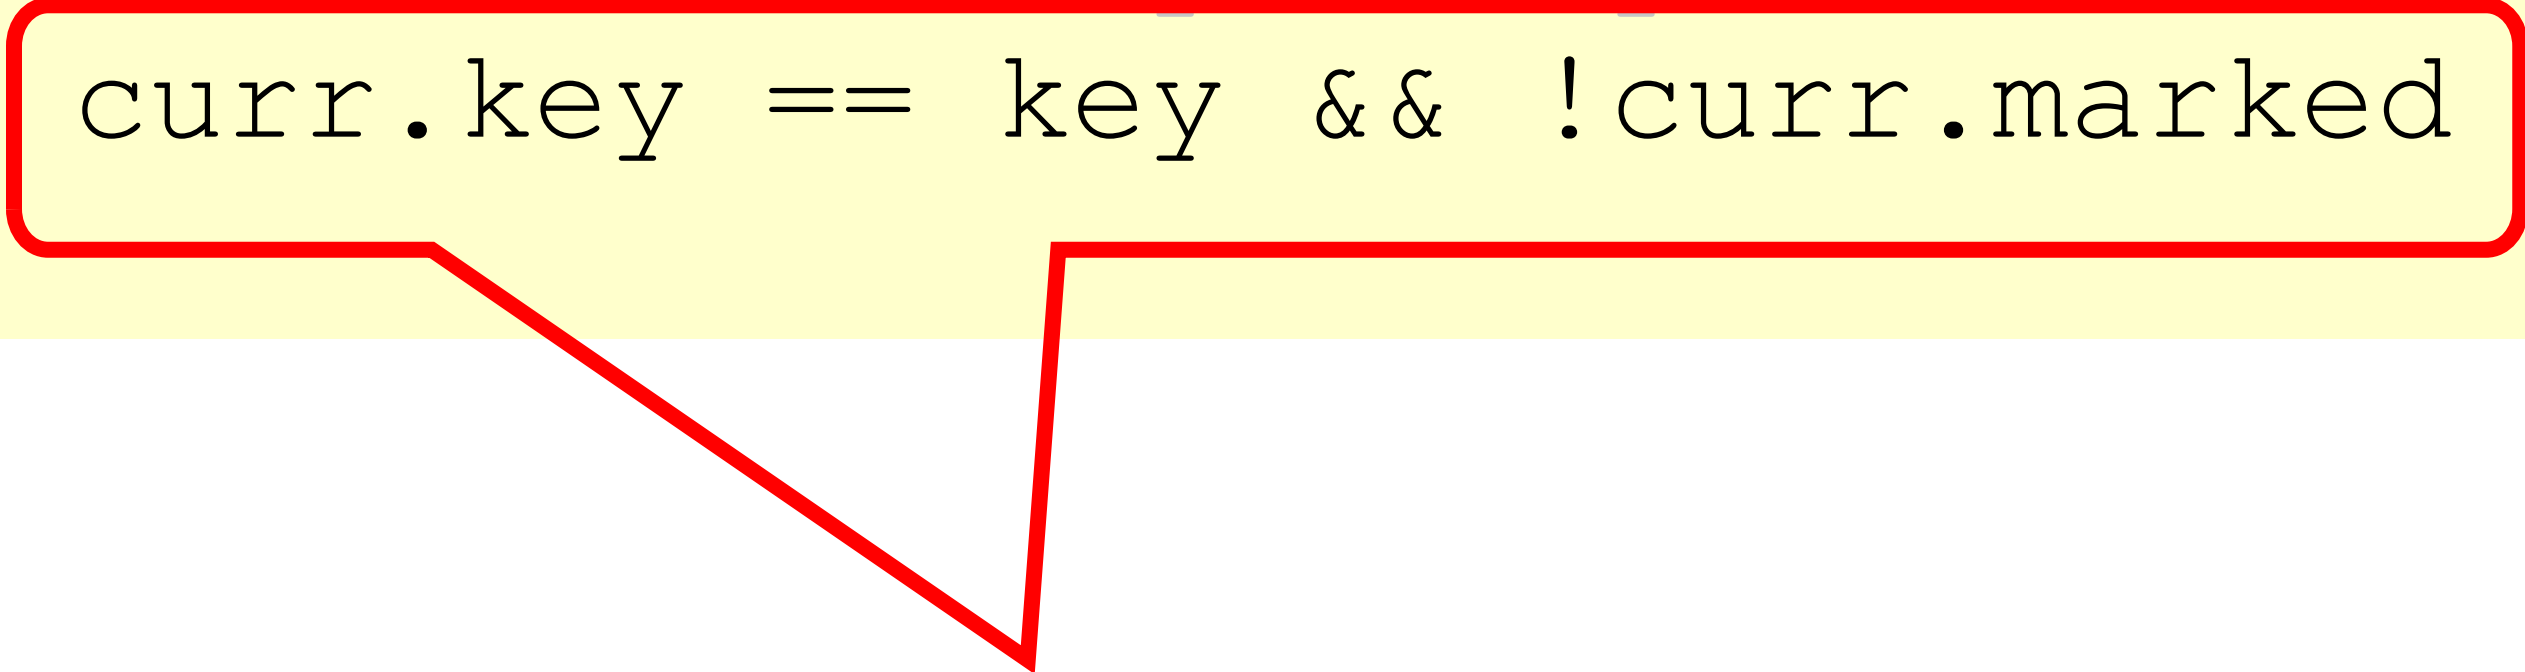
Contains

```
def contains(item: T) = {  
  val key = item.hashCode  
  var curr = this.head  
  while (curr.key < key) curr = curr.next  
  curr.key == key && !curr.marked  
}
```

Traverse *without locking*
(nodes may have been removed)


Contains

```
def contains(item: T) = {  
    val key = item.hashCode  
    var curr = this.head  
    while (curr.key < key) curr = curr.next  
    curr.key == key && !curr.marked  
}
```



Present and undeleted?


Summary: Wait-free Contains



Use Mark bit + list ordering

1. Not marked \rightarrow in the set
2. Marked or missing \rightarrow not in the set

Lazy List



Lazy add () and remove () + Wait-free contains ()

Evaluation

- Good:
 - **contains()** doesn't lock
 - In fact, its wait-free!
 - Good because typically high % contains()
 - Uncontended calls don't re-traverse
- Bad
 - Contended **add()** and **remove()** calls must re-traverse
 - Traffic jam if one thread delays

Traffic Jam

- Any concurrent data structure based on mutual exclusion has a weakness
- If one thread
 - Enters critical section
 - And “eats the big muffin”
 - Cache miss, page fault, descheduled ...
 - Everyone else using that lock is stuck!
 - Need to trust the scheduler....

Reminder: Lock-Free Data Structures




- No matter what ...
 - Guarantees minimal progress in any execution
 - i.e. Some thread will always complete a method call
 - Even if others halt at malicious times
 - Implies that implementation can't use locks

Lock-free Lists

- Next logical step
 - Wait-free `contains()`
 - lock-free `add()` and `remove()`
- Use only `compareAndSet()`
 - What could go wrong?

Lock-free Lists

Logical Removal




Physical Removal

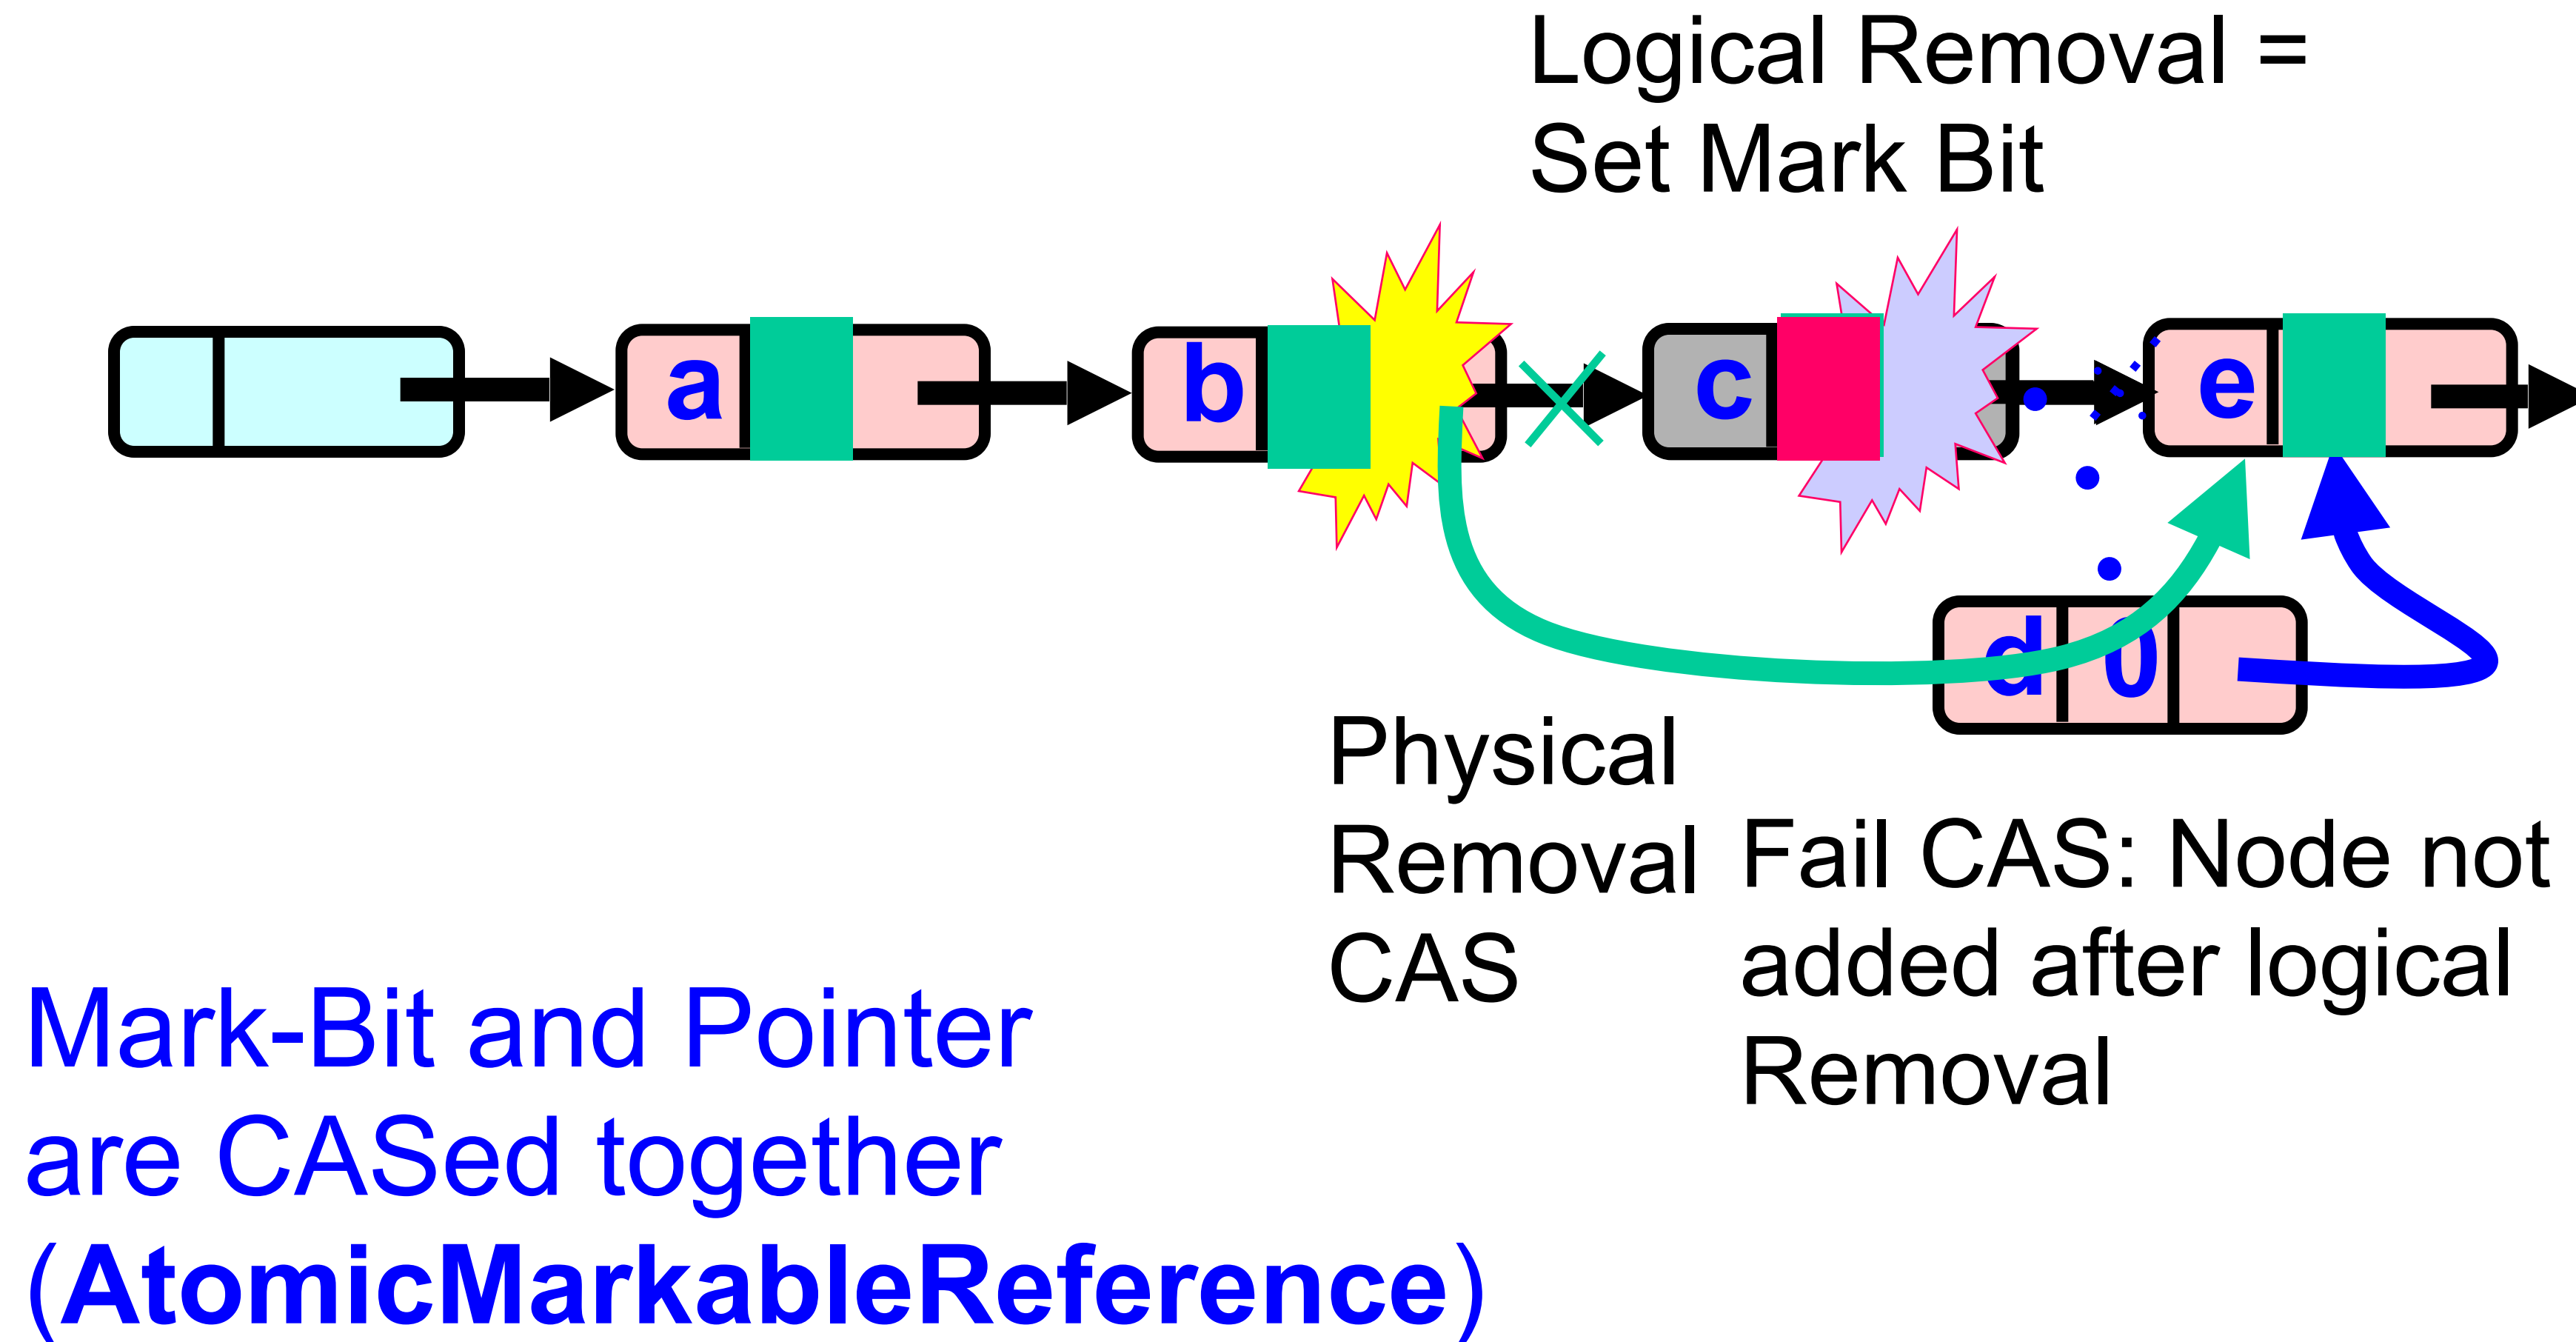
Use CAS to verify pointer
is correct

Not enough!

Problem...



The Solution: Combine Bit and Pointer




Solution

- Use AtomicMarkableReference
- Atomically
 - Swing reference and
 - Update flag
- Remove in two steps
 - Set mark bit in next field
 - Redirect predecessor's pointer

Marking a Node


- **AtomicMarkableReference** class
 - Java.util.concurrent.atomic package




Performance

- Different list-based set implementations
- 16-node machine
- Vary percentage of **contains()** calls


High Contains Ratio



Low Contains Ratio



As Contains Ratio Increases



Summary

- Coarse-grained locking
- Fine-grained locking (“hand-over-hand”)
- Optimistic synchronization
- Lazy synchronization
- Lock-free synchronization

“To Lock or Not to Lock”

- Locking vs. Non-blocking:
 - Extremist views on both sides
 - Locking: longs waits
 - Non-blocking: long “clean-ups”
- The answer: nobler to compromise
 - Example: Lazy list combines blocking **add()** and **remove()** and a wait-free **contains()**
 - Remember: Blocking/non-blocking is a property of a method



This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](https://creativecommons.org/licenses/by-sa/2.5/).

- **You are free:**
 - **to Share** — to copy, distribute and transmit the work
 - **to Remix** — to adapt the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
 - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
 - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.