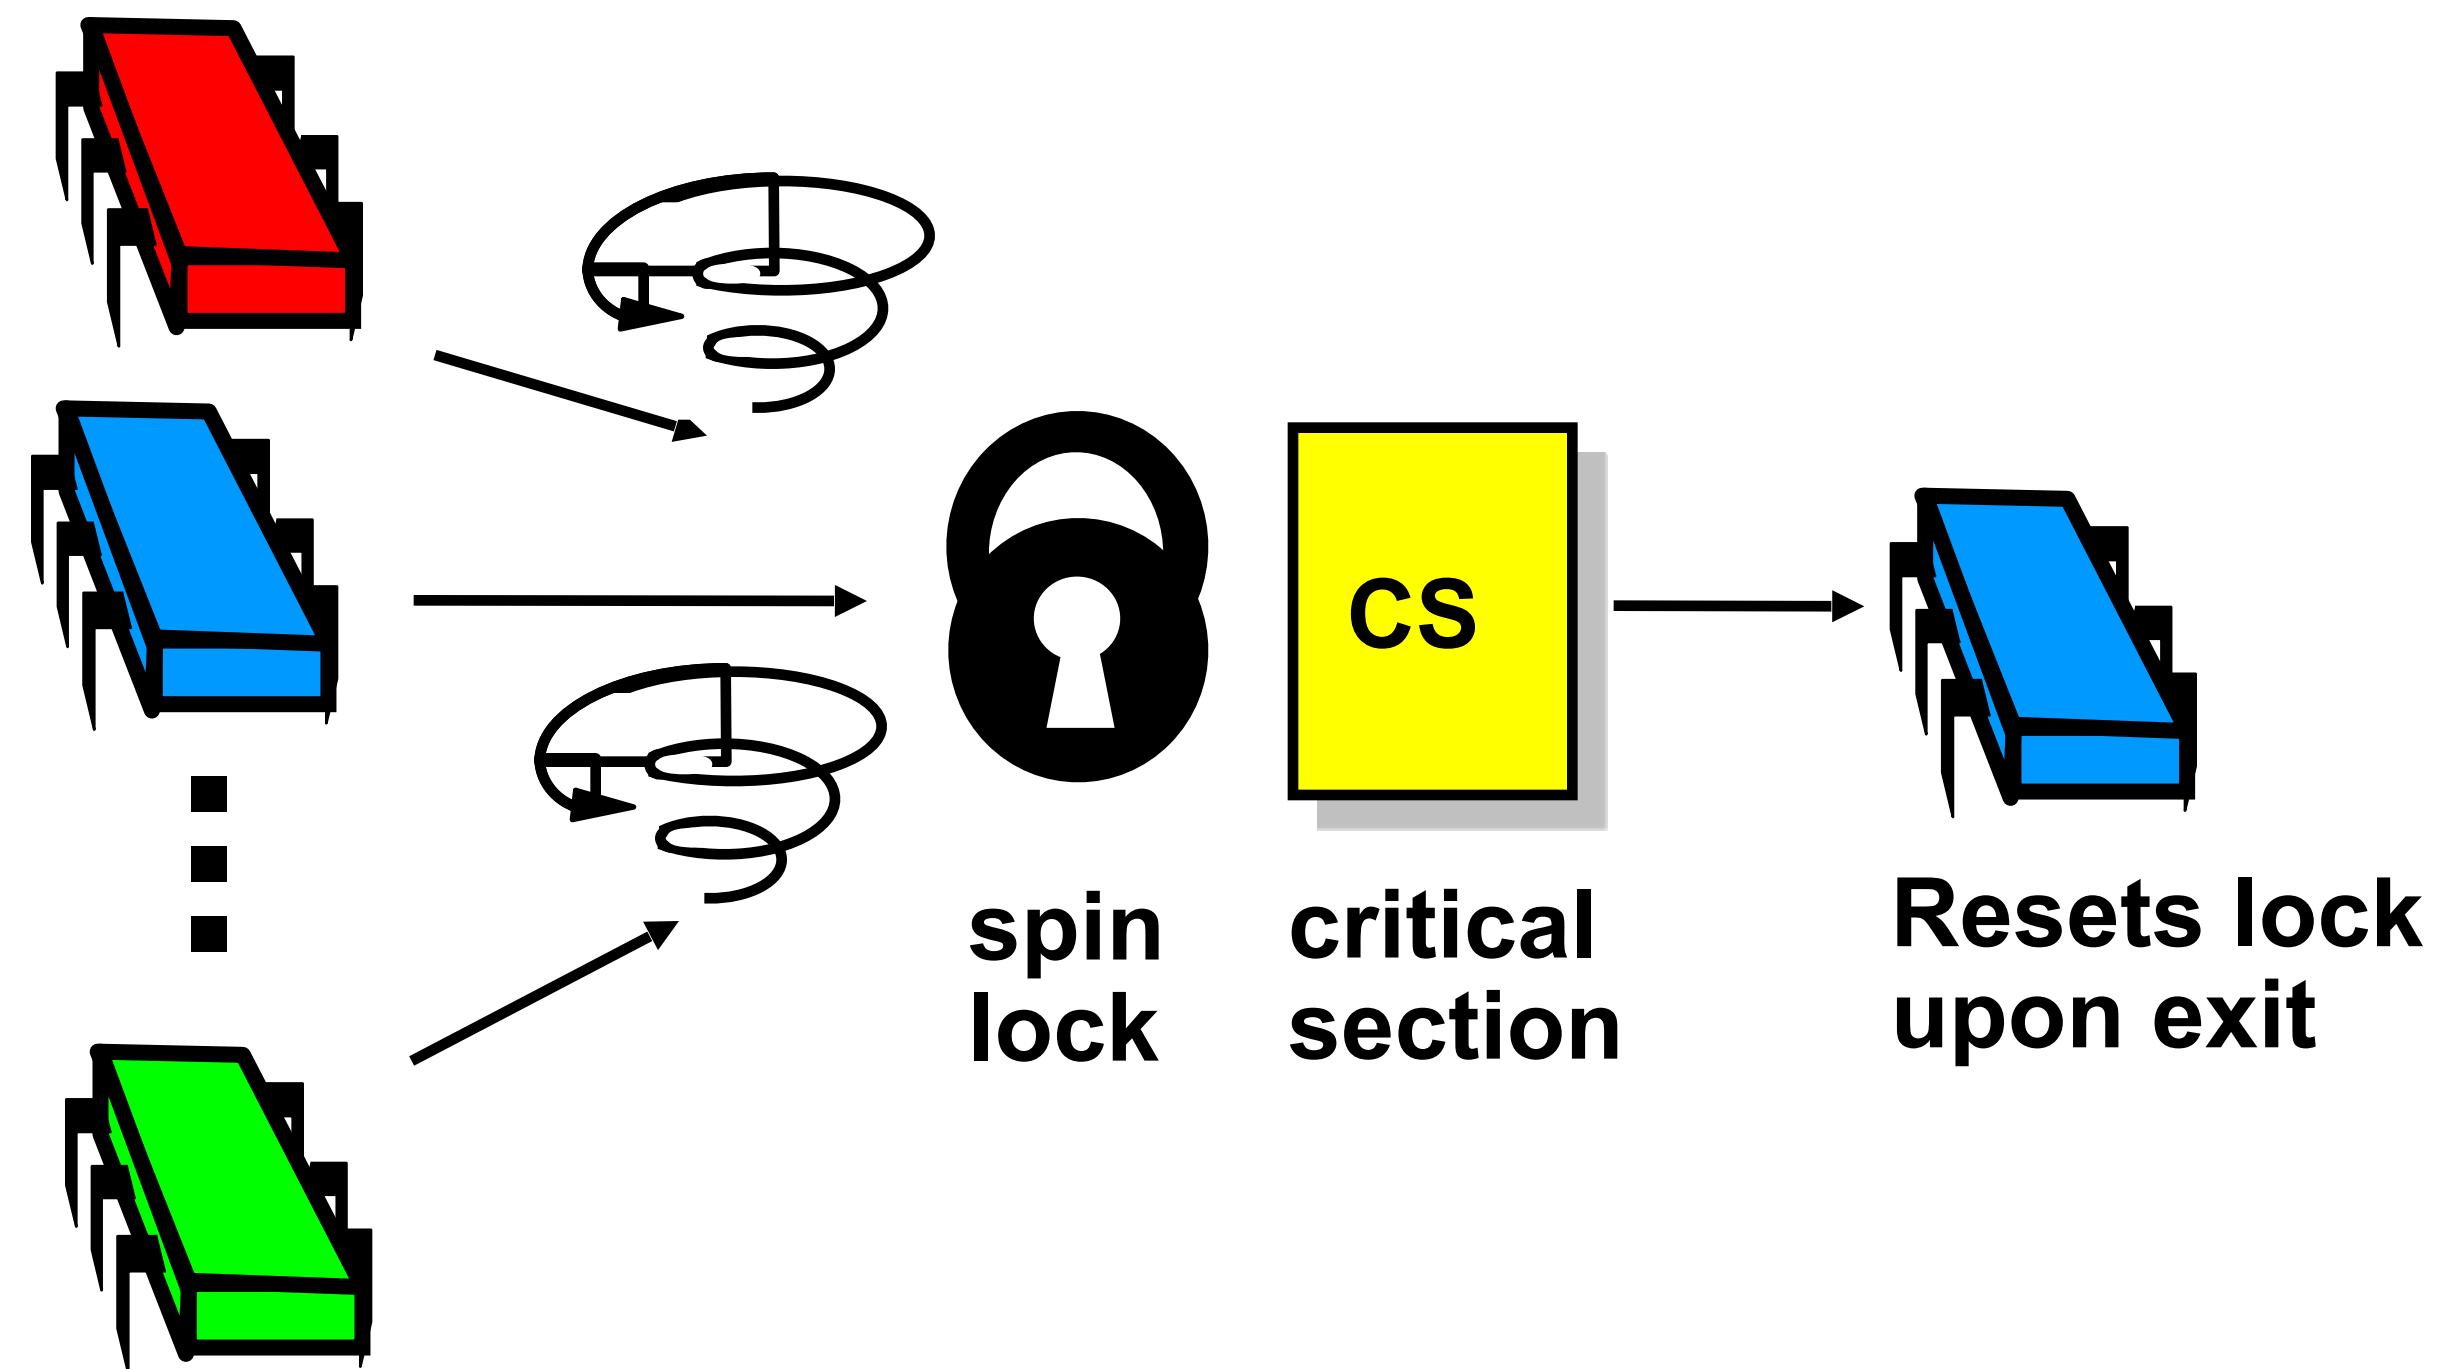


YSC3248: Parallel, Concurrent and Distributed Programming

Concurrent Linked Lists

Previous Lectures: Spin-Locks



Today: More Concurrent Objects

- Adding threads should not lower throughput
 - Contention effects
 - Can be mitigated by back-offs, arrays, etc.

Today: More Concurrent Objects

- Adding threads should not lower throughput
 - Contention effects
 - Can be mitigated by back-offs, arrays, etc.
- Should increase throughput
 - Not possible if inherently sequential
 - Surprising things are parallelizable

Coarse-Grained Synchronization

- Each method locks the object
 - Avoid contention using queue locks

Coarse-Grained Synchronization

- Each method locks the object
 - Avoid contention using locks
 - Easy to reason about
 - In simple cases

Coarse-Grained Synchronization

- Each method locks the object
 - Avoid contention using locks
 - Easy to reason about
 - In simple cases
- So, are we done?

Coarse-Grained Synchronization

- Sequential bottleneck
 - Threads “stand in line”

Coarse-Grained Synchronization

- Sequential bottleneck
 - Threads “stand in line”
- Adding more threads
 - Does not improve throughput
 - Struggle to keep it from getting worse

Coarse-Grained Synchronization

- Sequential bottleneck
 - Threads “stand in line”
- Adding more threads
 - Does not improve throughput
 - Struggle to keep it from getting worse
- So why even use a multiprocessor?
 - Well, some apps inherently parallel ...

This Lecture

- Introduce several “patterns”
 - Bag of tricks ...
 - Methods that work more than once ...

This Lecture

- Introduce several “patterns”
 - Bag of tricks ...
 - Methods that work more than once ...
- For highly-concurrent objects
 - Concurrent access
 - More threads, more throughput

First:

Fine-Grained Synchronization

- Instead of using a single lock ...

First:

Fine-Grained Synchronization

- Instead of using a single lock ...
- Split object into
 - Independently-synchronized components

First:

Fine-Grained Synchronization

- Instead of using a single lock ...
- Split object into
 - Independently-synchronized components
- Methods conflict when they access
 - The same component ...
 - At the same time

Second: Optimistic Synchronization

- Search without locking ...

Second: Optimistic Synchronization

- Search without locking ...
- If you find it, lock and check ...
 - OK: we are done
 - Oops: start over

Second: Optimistic Synchronization

- Search without locking ...
- If you find it, lock and check ...
 - OK: we are done
 - Oops: start over
- Evaluation
 - Usually cheaper than locking, but
 - Mistakes are expensive

Third: Lazy Synchronization

- Postpone hard work

Third:

Lazy Synchronization

- Postpone hard work
- Removing components is tricky

Third:

Lazy Synchronization

- Postpone hard work
- Removing components is tricky
 - Logical removal
 - Mark component to be deleted

Third:

Lazy Synchronization

- Postpone hard work
- Removing components is tricky
 - Logical removal
 - Mark component to be deleted
 - Physical removal
 - Do what needs to be done

Fourth:

Lock-Free Synchronization

- Don't use locks at all
 - Use `compareAndSet()` & relatives ...

Fourth:

Lock-Free Synchronization

- Don't use locks at all
 - Use `compareAndSet()` & relatives ...
- Advantages
 - No Scheduler Assumptions/Support

Fourth:

Lock-Free Synchronization

- Don't use locks at all
 - Use `compareAndSet()` & relatives ...
- Advantages
 - No Scheduler Assumptions/Support
- Disadvantages
 - Complex
 - Sometimes high overhead

Linked List

- Illustrate these patterns ...
- Using a list-based Set
 - Common application
 - Building block for other apps

Set Interface

- Unordered collection of items

Set Interface

- Unordered collection of items
- No duplicates

Set Interface

- Unordered collection of items
- No duplicates
- Methods
 - **add (x)** put **x** in set
 - **remove (x)** take **x** out of set
 - **contains (x)** tests if **x** in set

Warm-up: Testing Concurrent Sets

List-Based Sets

```
trait ConcurrentSet[T] {  
  def add(item: T): Boolean  
  def remove(item: T): Boolean  
  def contains(item: T): Boolean  
}
```

List-Based Sets

```
trait ConcurrentSet[T] {  
  def add(item: T): Boolean  
  def remove(item: T): Boolean  
  def contains(item: T): Boolean  
}
```

Add item to set

List-Based Sets

```
trait ConcurrentSet[T] {  
  def add(item: T): Boolean  
  def remove(item: T): Boolean  
  def contains(item: T): Boolean  
}
```

Remove item from set

List-Based Sets

```
trait ConcurrentSet[T] {  
  def add(item: T): Boolean  
  def remove(item: T): Boolean  
  def contains(item: T): Boolean  
}
```

Is item in set?

List Node

```
class Node (val item: T) {  
    def key : Int  
    @volatile var next: Node = _  
}
```

List Node

```
class Node (val item: T) {  
    def key : Int  
    @volatile var next: Node = _  
}
```

item of interest

List Node

```
class Node (val item: T) {  
  def key : Int  
  @volatile var next: Node = _  
}
```

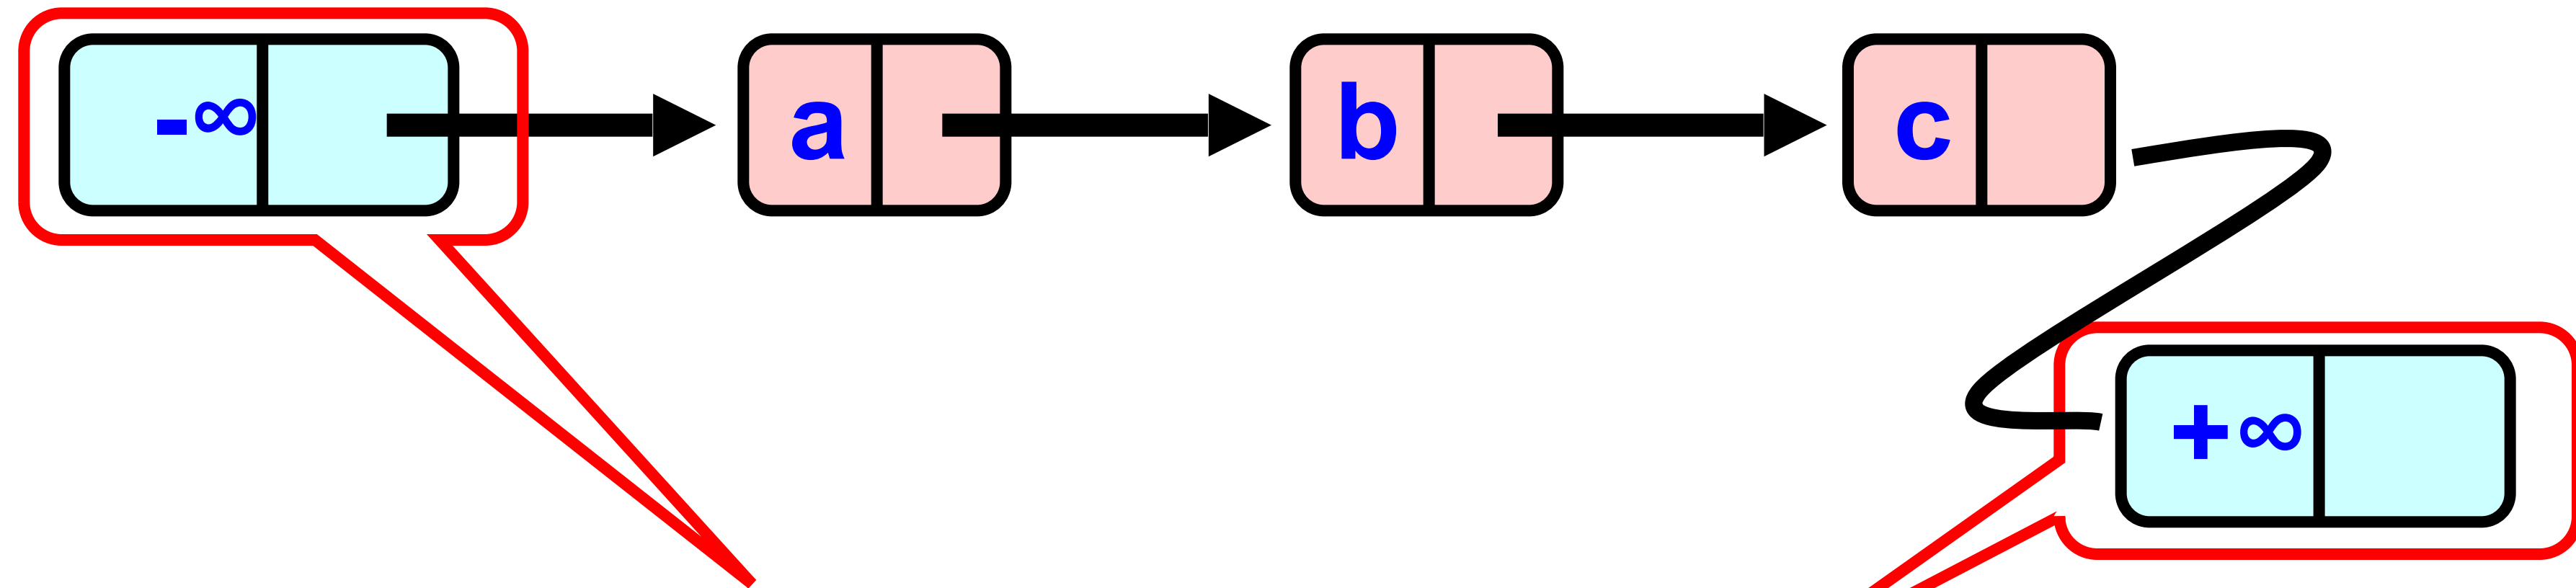
Usually hash code

List Node

```
class Node (val item: T) {  
    def key : Int  
    @volatile var next: Node = _  
}
```

Reference to next node

The List-Based Set



Sorted with Sentinel nodes
(min & max possible keys)

Reasoning about Concurrent Objects

- Invariant
 - Property that always holds

Reasoning about Concurrent Objects

- Invariant
 - Property that always holds
- Established because
 - True when object is **created**
 - Truth **preserved** by each method
 - Each **step** of each method

Specifically ...

- Invariants preserved by
 - `add()`
 - `remove()`
 - `contains()`

Specifically ...

- Invariants preserved by
 - `add()`
 - `remove()`
 - `contains()`
- Most steps are trivial
 - Usually one step tricky
 - Often it is the linearization point

Interference

- Invariants make sense only if
 - methods considered
 - are the only modifiers

Interference

- Invariants make sense only if
 - methods considered
 - are the only modifiers
- Language encapsulation helps
 - List nodes *not visible* outside class

Interference


- Invariants make sense only if
 - methods considered
 - are the only modifiers
- Language encapsulation helps
 - List nodes *not visible* outside class
- Similar to *loop invariants*
 - Each method must preserve the invariant (same as each loop iteration)

Interference

- Freedom from interference needed even for removed nodes
 - Some algorithms traverse removed nodes
 - Careful with `malloc()` & `free()`!
- We rely on garbage collection

Recap: Abstract Data Types

- Concrete representation:



- Abstract Type:
 $\{a, b\}$

Abstract Data Types

- Meaning of rep given by abstraction map

$$S(\text{[]} \rightarrow \text{[a]} \rightarrow \text{[b]} \rightarrow \text{[]}) = \{a, b\}$$

Representation Invariant

- Which concrete values meaningful?
 - Sorted?
 - Duplicates?
- Representation invariant
 - Characterises legal *concrete representations*
 - Preserved by methods
 - Relied on by methods

Blame Game

- Rep invariant is a **contract**
- Suppose
 - **add()** leaves behind 2 copies of *x*
 - **remove()** removes only 1
- Which is incorrect?

Blame Game

- Suppose
 - **add()** leaves behind 2 copies of x
 - **remove()** removes only 1

Blame Game

- Suppose
 - **add()** leaves behind 2 copies of x
 - **remove()** removes only 1
- Which is incorrect?
 - If rep invariant says *no duplicates*
 - **add()** is incorrect
 - Otherwise
 - **remove()** is incorrect

Lists' Rep Invariant (partly)

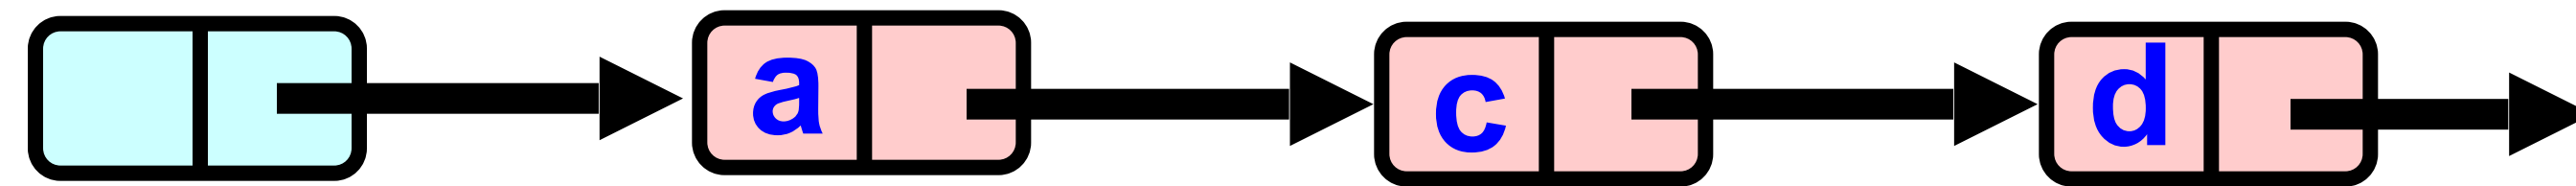
- Sentinel nodes
 - tail reachable from head
- Sorted
- No duplicates

Abstraction Map

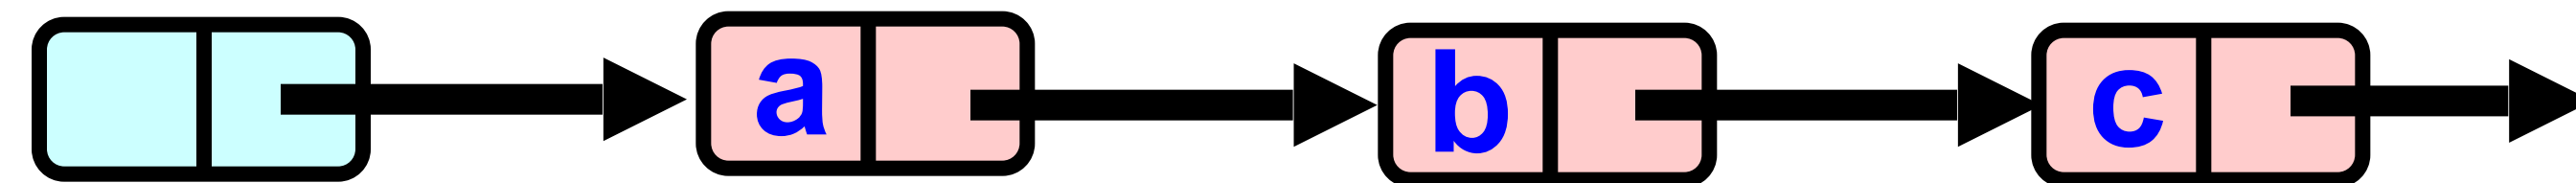
- $S(\text{head}) =$
 $\{ x \mid \text{there exists } a \text{ such that}$
 - $a \text{ reachable from head and}$
 - $a.\text{item} = x$ $\}$

Sequential List Based Set

`add()`

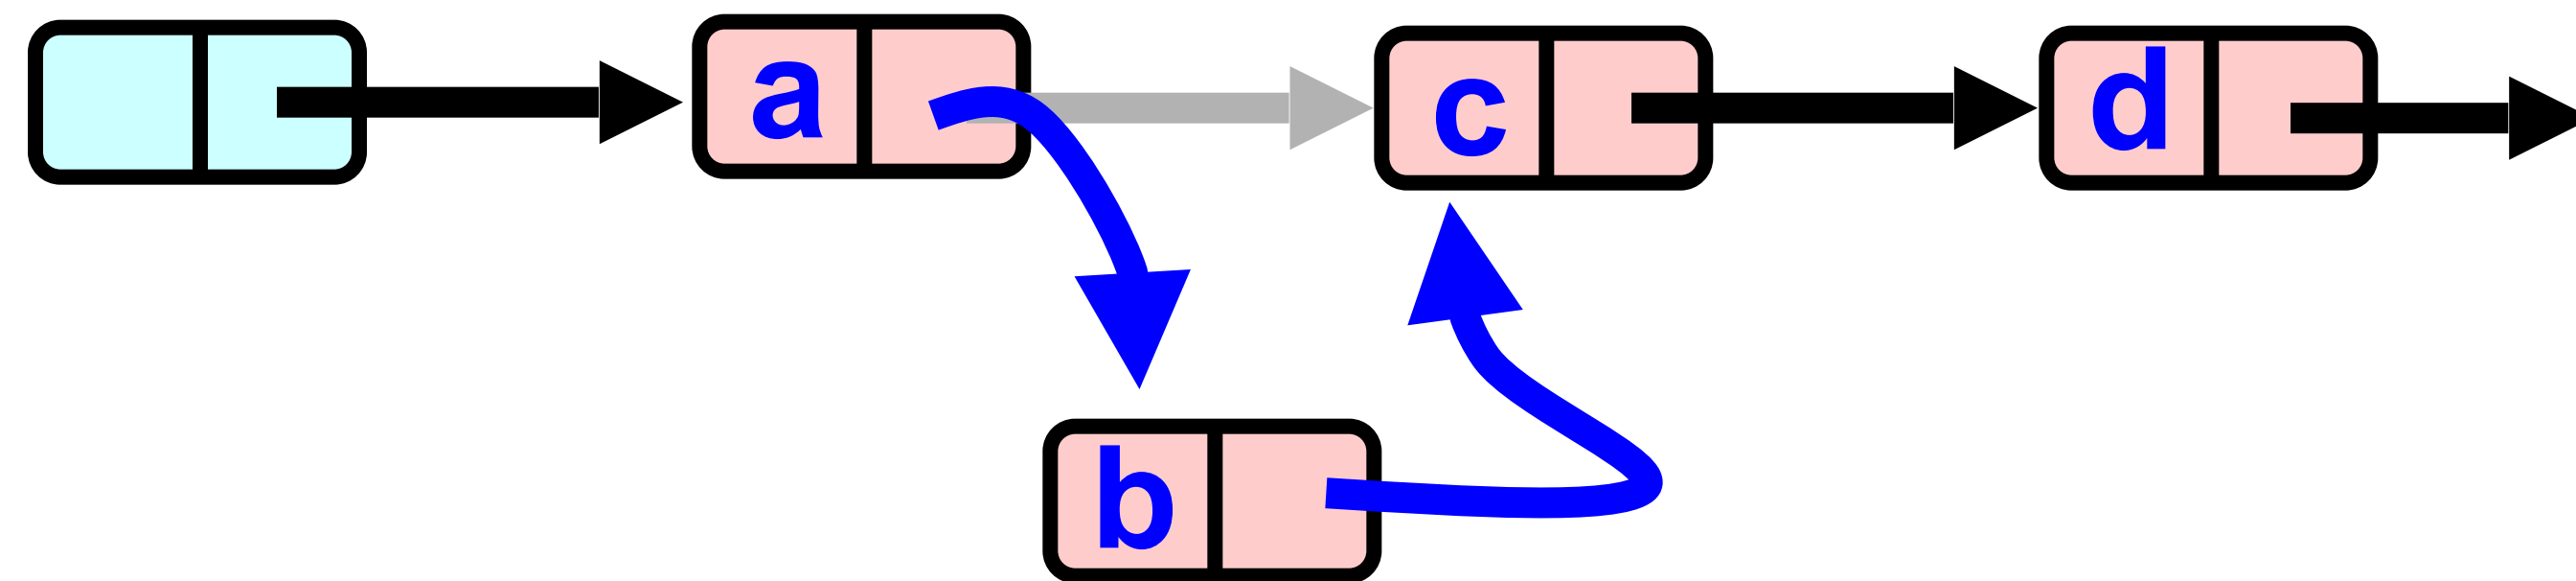


`remove()`

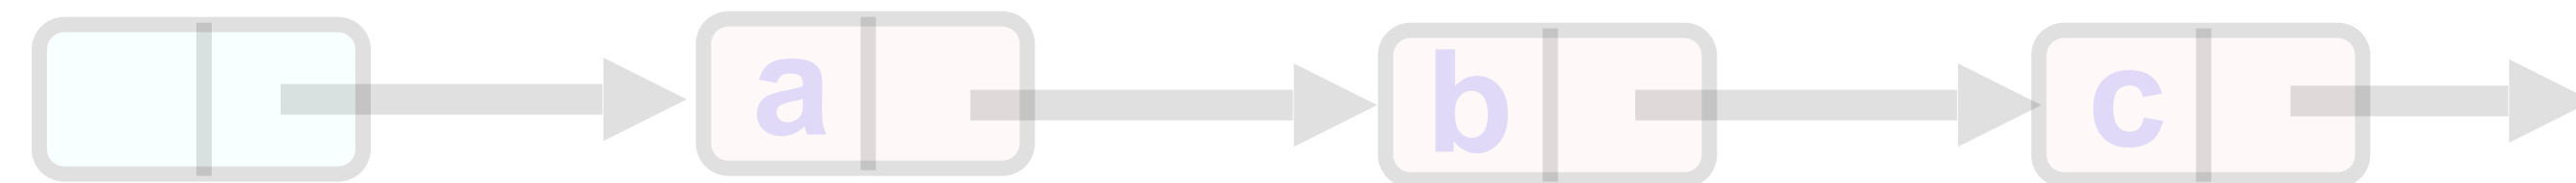


Sequential List Based Set

add()

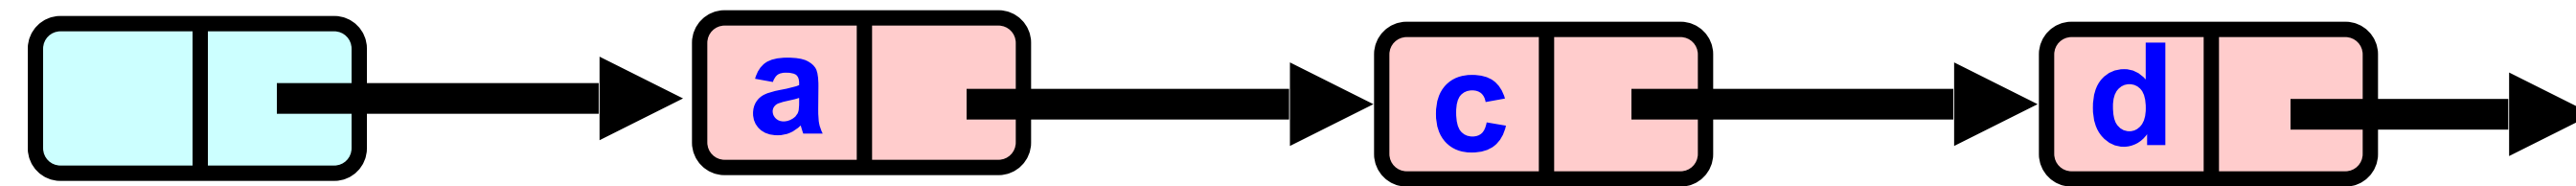


remove()

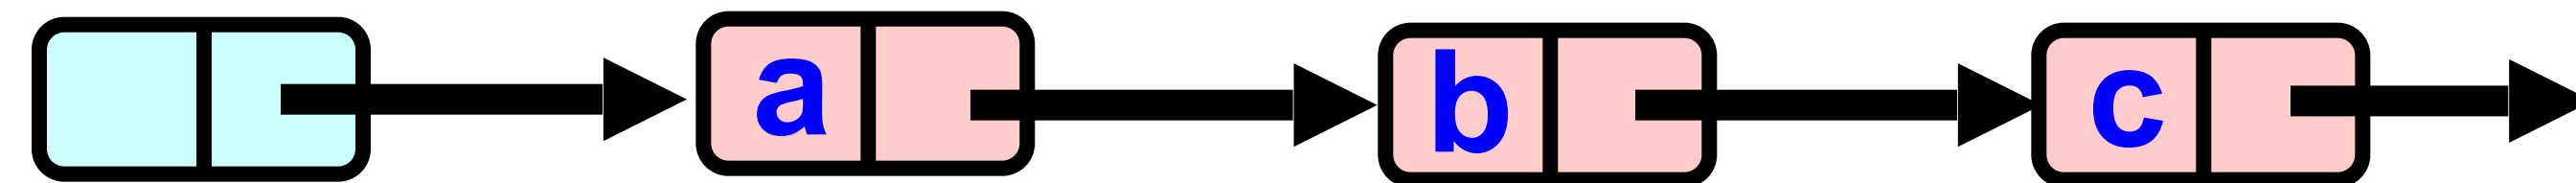


Sequential List Based Set

`add()`

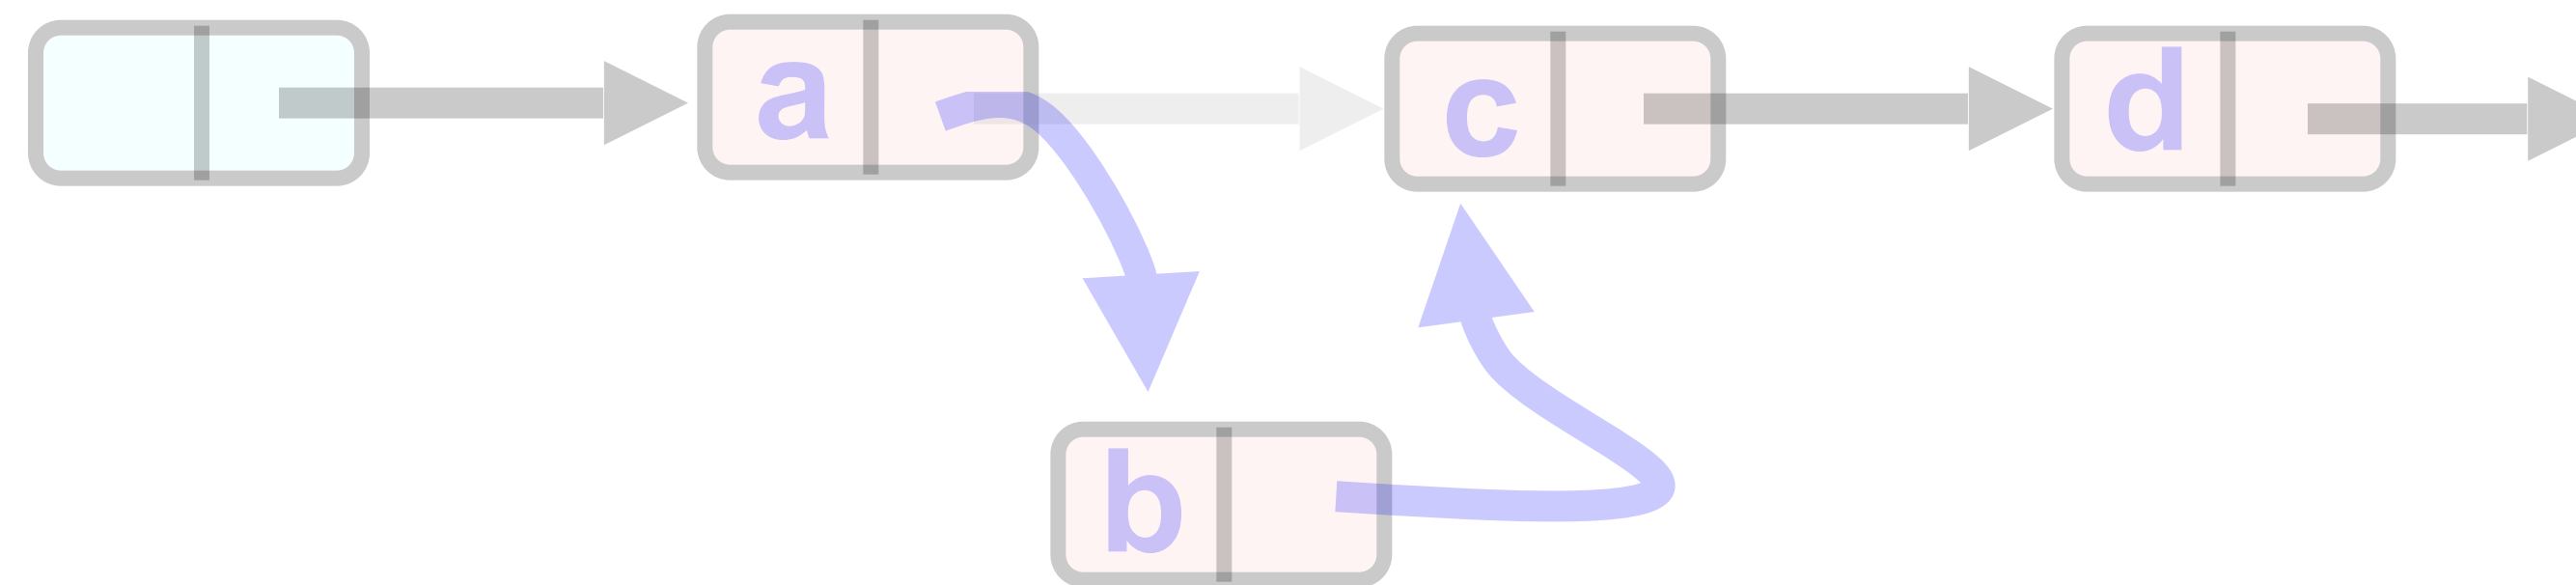


`remove()`

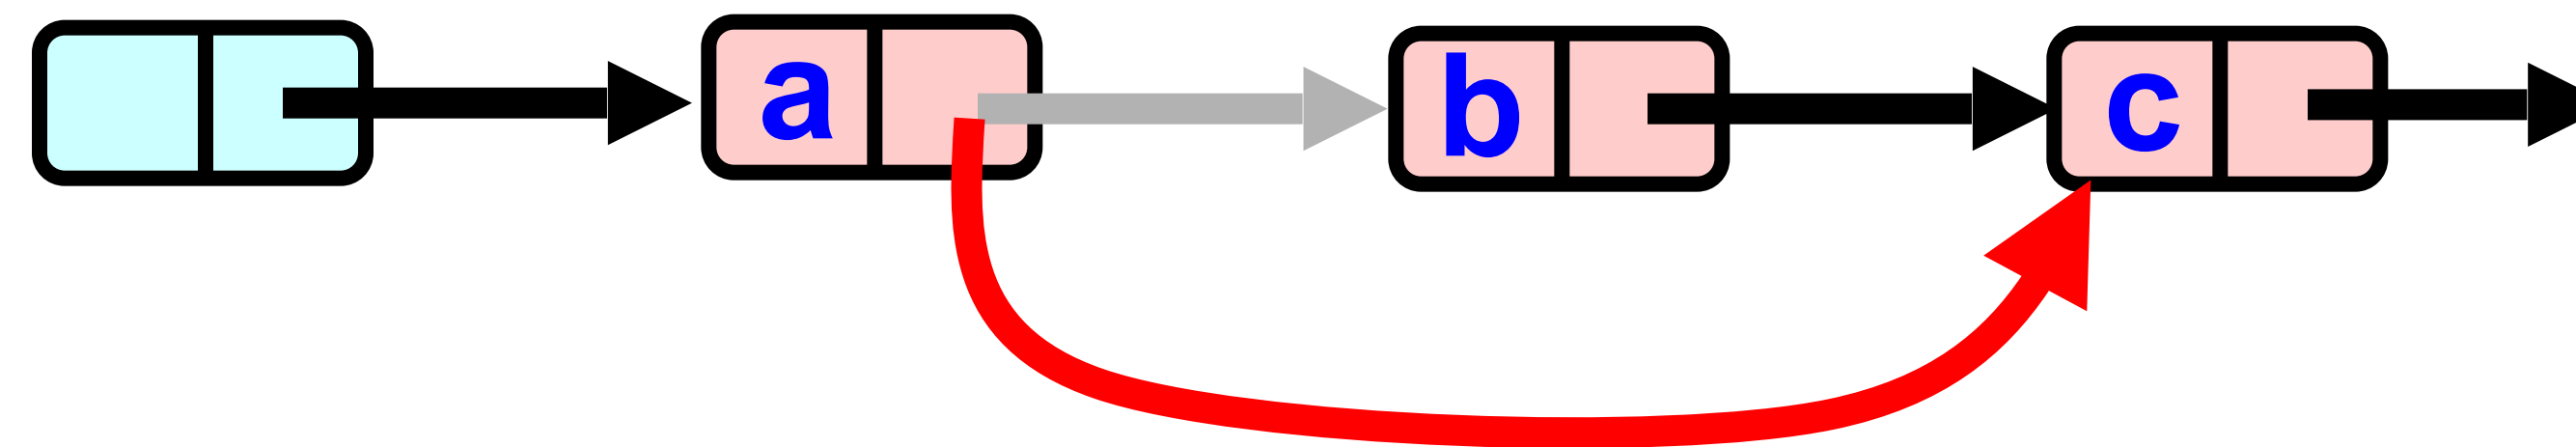


Sequential List Based Set

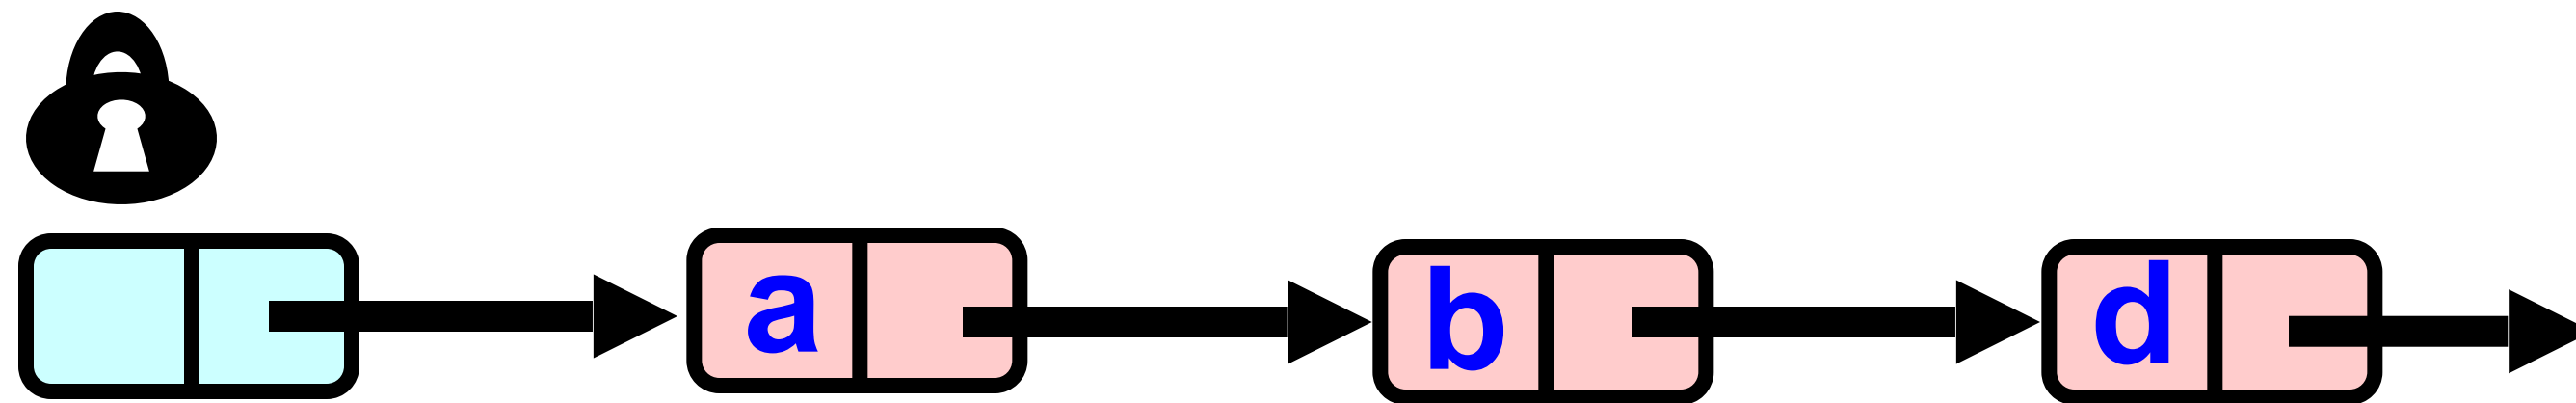
add()



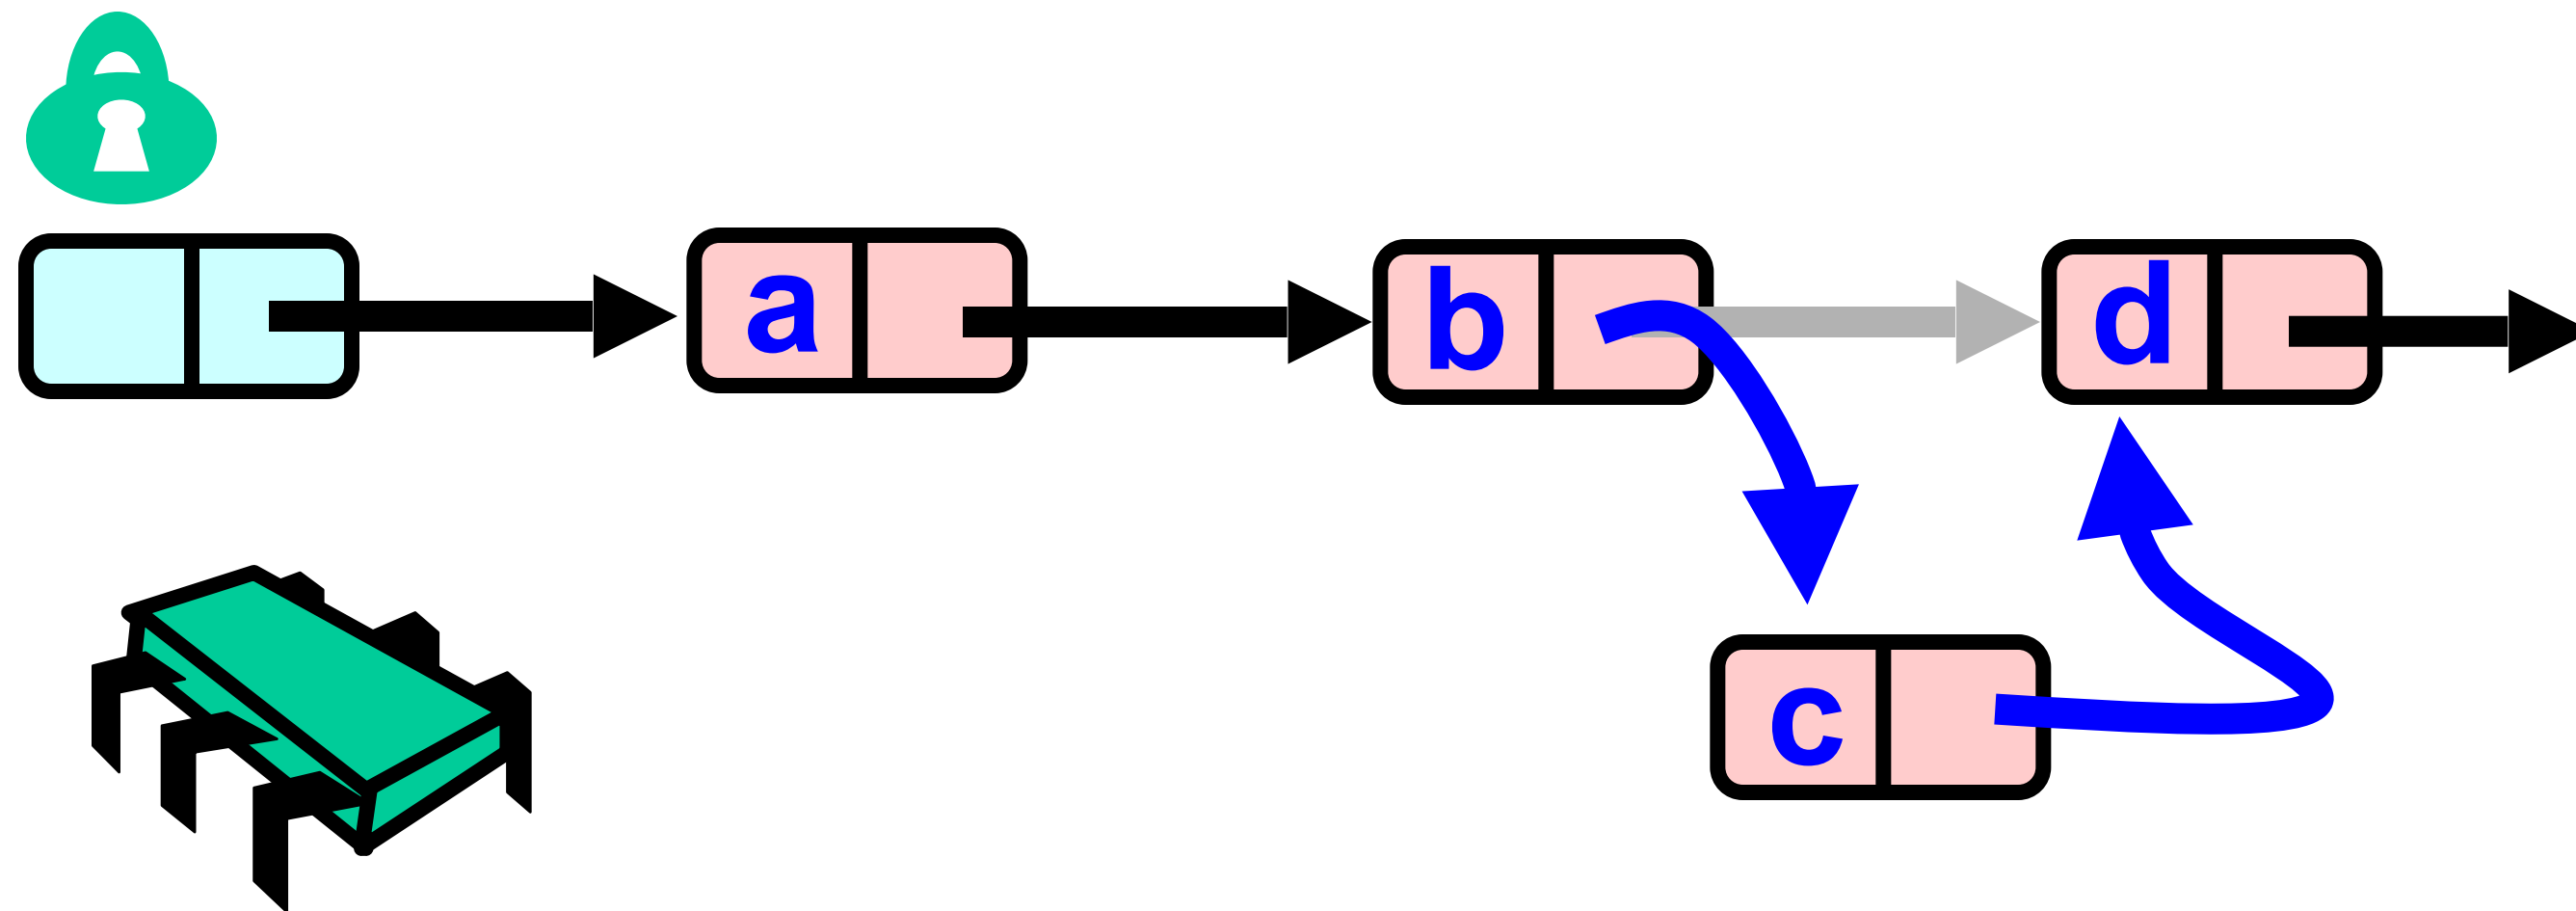
remove()




Coarse-Grained Locking



Coarse-Grained Locking



Coarse-Grained Locking



Simple but hotspot + bottleneck

Coarse-Grained Locking

- Easy, same as synchronized methods
 - “One lock to rule them all ...”

Coarse-Grained Locking

- Easy, same as synchronized methods
 - “One lock to rule them all ...”
- Simple, clearly correct
 - Deserves respect!
- Works poorly with contention
 - Queue locks help
 - But bottleneck still an issue

Demo: Benchmarking Concurrent Lists



Fine-grained Locking

- Requires **careful thought**
 - “Do not meddle in the affairs of wizards, for they are subtle and quick to anger”


Fine-grained Locking

- Requires **careful** thought
 - “Do not meddle in the affairs of wizards, for they are subtle and quick to anger”
- Split object into pieces
 - Each piece has own lock
 - Methods that work on disjoint pieces need not exclude each other


Hand-over-Hand locking




Hand-over-Hand locking




Hand-over-Hand locking





Hand-over-Hand locking




Hand-over-Hand locking




Removing a Node




Removing a Node




Removing a Node




Removing a Node



Removing a Node




Removing a Node




<*A good place to pause*>


Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes



Concurrent Removes




Uh, Oh




Uh, Oh

Bad news, c not removed




Problem

- To delete node c
 - Swing node b's next field to d




- Problem is,
 - Someone deleting b concurrently could direct a pointer to c




Hand-over-Hand Locking: Insight

- If a node is locked
 - No one can delete node's *successor*
- If a thread locks
 - Node to be deleted
 - And its predecessor
 - Then it works


Hand-Over-Hand Again




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node



Removing a Node



Removing a Node



Remove method

```
def remove(item: T): Boolean = {  
    var pred, curr: Node = null  
    val key = item.hashCode  
  
    try { ... } finally {  
        curr.unlock()  
        pred.unlock()  
    }  
}
```

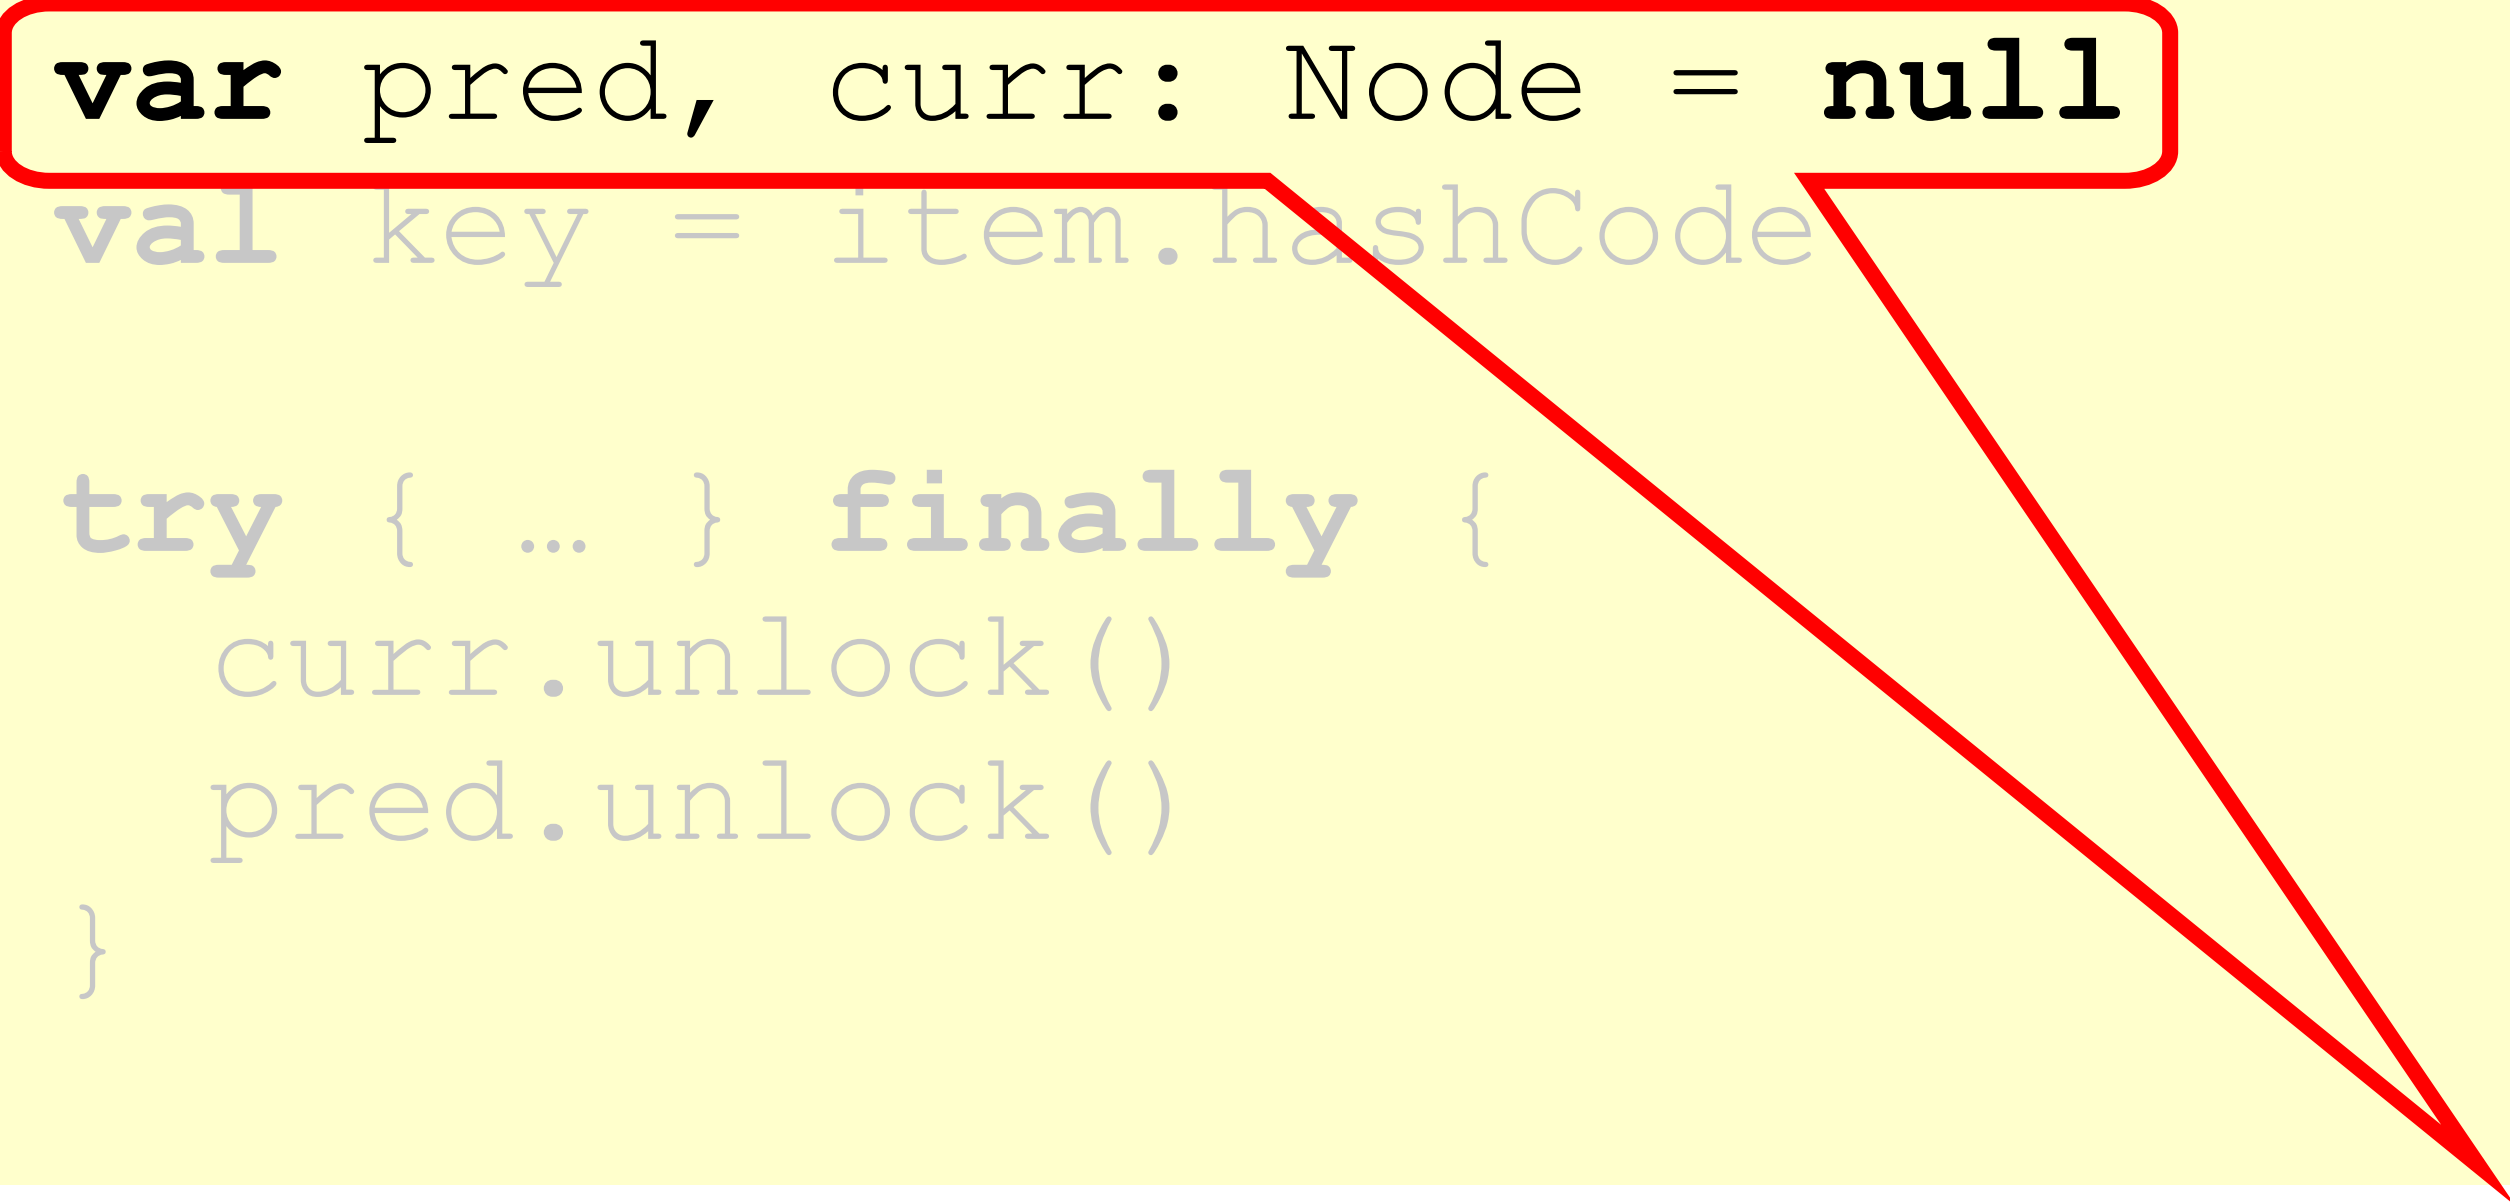
Remove method

```
def remove(item: T): Boolean = {  
    var pred, curr: Node = null  
    val key = item.hashCode  
  
    try { ... } finally {  
        curr.unlock()  
        pred.unlock()  
    }  
}
```

Key used to order node

Remove method

```
def remove(item: T): Boolean = {  
  var pred, curr: Node = null  
  val key = item.hashCode  
  
  try { ... } finally {  
    curr.unlock()  
    pred.unlock()  
  }  
}
```



Predecessor and current nodes

Remove method

```
def remove(item: T): Boolean = {  
  var pred, curr: Node = null  
  val key = item.hashCode  
  
  try { ... } finally {  
    curr.unlock()  
    pred.unlock()  
  }  
}
```

**Make sure
locks released**

Remove method

```
def remove(item: T): Boolean = {  
    var pred, curr: Node = null  
    val key = item.hashCode  
  
    try { ... } finally {  
        curr.unlock()  
        pred.unlock()  
    }  
}
```

Everything else


Remove method

```
try {  
    pred = head  
    pred.lock()  
    curr = pred.next  
    curr.lock()  
    ...  
} finally { ... }
```

Remove method

lock pred == head


```
try {  
    pred = head  
    pred.lock()  
    curr = pred.next  
    curr.lock()  
    ...  
} finally { ... }
```



Remove method

```
try {  
    pred = head;  
    pred.lock();  
    curr = pred.next  
    curr.lock()  
    ...  
} finally { ... }
```


Lock current



Remove method


```
try {  
    pred = head  
    pred.lock()  
    curr = pred.next  
    curr.lock()  
    ...  
} finally { ... }
```

Traversing list



Remove: searching


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```


Search key range



Remove: searching

```
while (curr.key <= key) :  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```


**At start of each loop:
curr and pred locked**



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

If item found, remove node




Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

pred.unlock()


Unlock predecessor



Remove: searching

Only one node locked!


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

demote current


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

Find and lock new current


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = currNode  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

Loop invariant restored

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = currNode  
    curr = curr.next  
    curr.lock()  
}  
return false
```



Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}
```

Otherwise, not present

return false

Why does this work?

- To remove node e
 - Must lock e
 - Must lock e 's predecessor
- Therefore, if you lock a node
 - It can't be removed
 - And neither can its successor

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

- **pred** reachable from head
- **curr** is pred.next
- So **curr.item** is in the set

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

**Linearization point if
item is present**

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}  
return false
```

**Node locked, so no other
thread can remove it**

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}
```

return false;

Item not present

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next  
        return true  
    }  
    pred.unlock()  
    pred = curr  
    curr = curr.next  
    curr.lock()  
}
```

return false

- pred **reachable** from head
- curr **is** pred.next
- pred.key < **key**
- **key** < curr.key

Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

Linearization point



Adding Nodes

- To add node e
 - Must lock predecessor
 - Must lock successor
- Neither can be deleted
 - (Is successor lock actually required?)

Same Abstraction Map

- $S(\text{head}) =$
 $\{ x \mid \text{there exists } a \text{ such that}$
 - $a \text{ reachable from head and}$
 - $a.\text{item} = x$ $\}$

Rep Invariant

- Easy to check that
 - tail always reachable from head
 - Nodes sorted, no duplicates

Demo: Benchmarking Fine-Grained Lists


Drawbacks

- Better than coarse-grained lock
 - Threads can traverse in parallel
- Still not ideal
 - Long chain of acquire/release
 - Inefficient


Optimistic Synchronization

- Find nodes without locking
- Lock nodes
- Check that everything is OK


Optimistic: Traverse without Locking




Optimistic: Lock and Load




Optimistic: Lock and Load




What could go wrong?




What could go wrong?




What could go wrong?




What could go wrong?



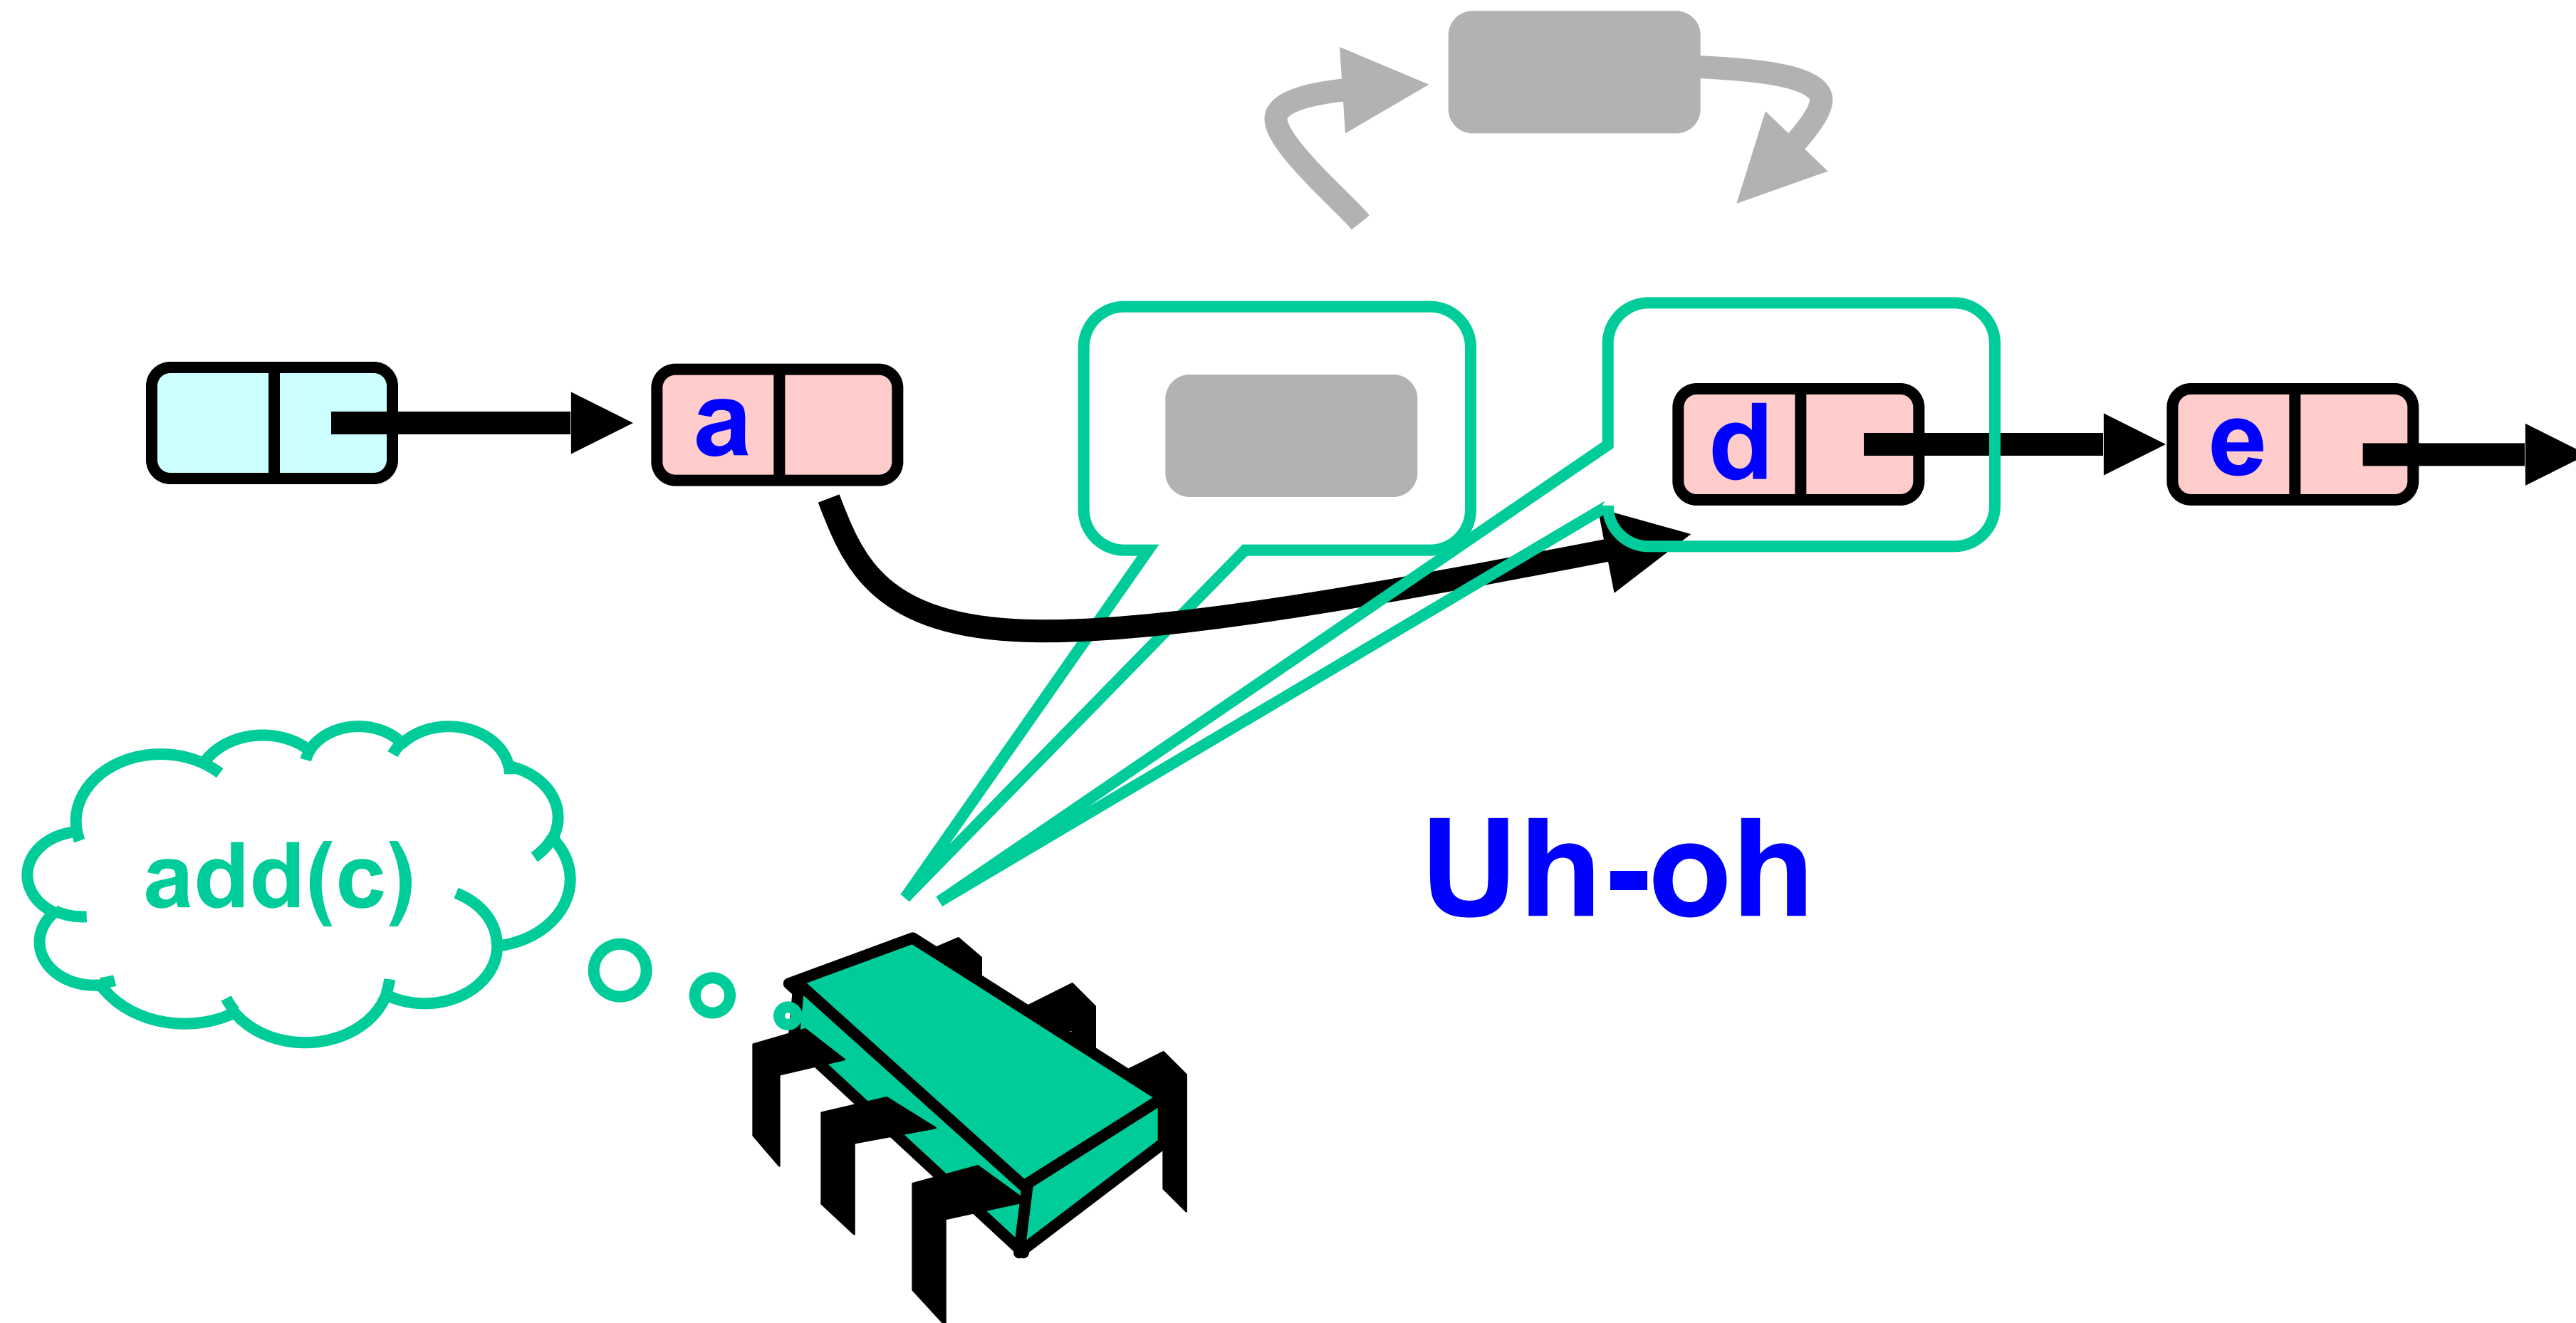
What could go wrong?



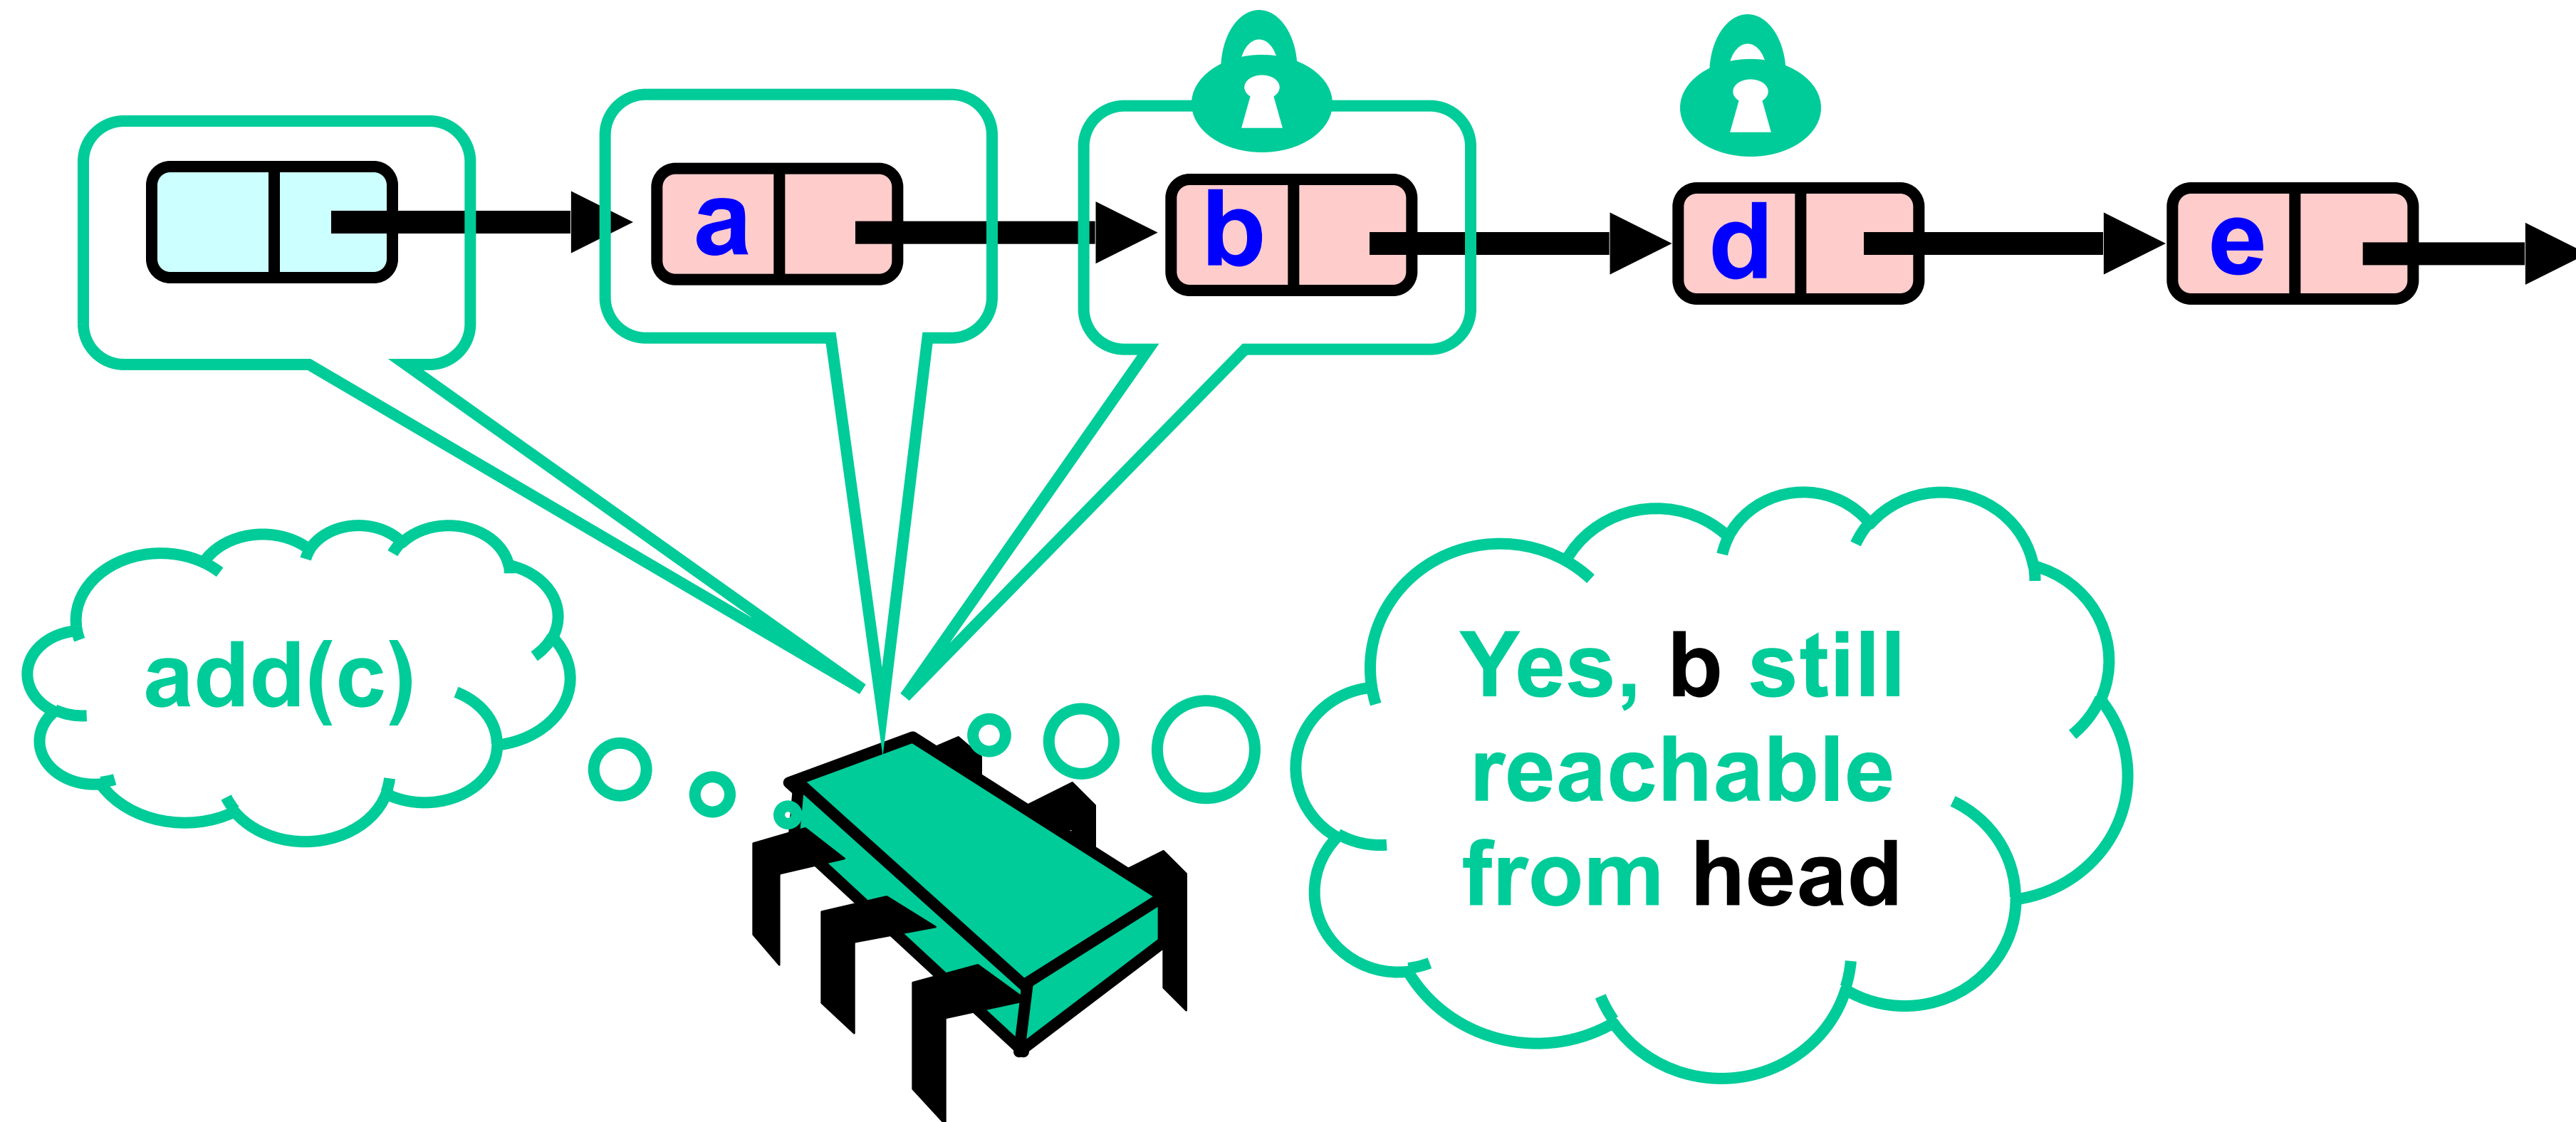
What could go wrong?




What could go wrong?




Validate – Part 1




What Else Could Go Wrong?




What Else Could Go Wrong?



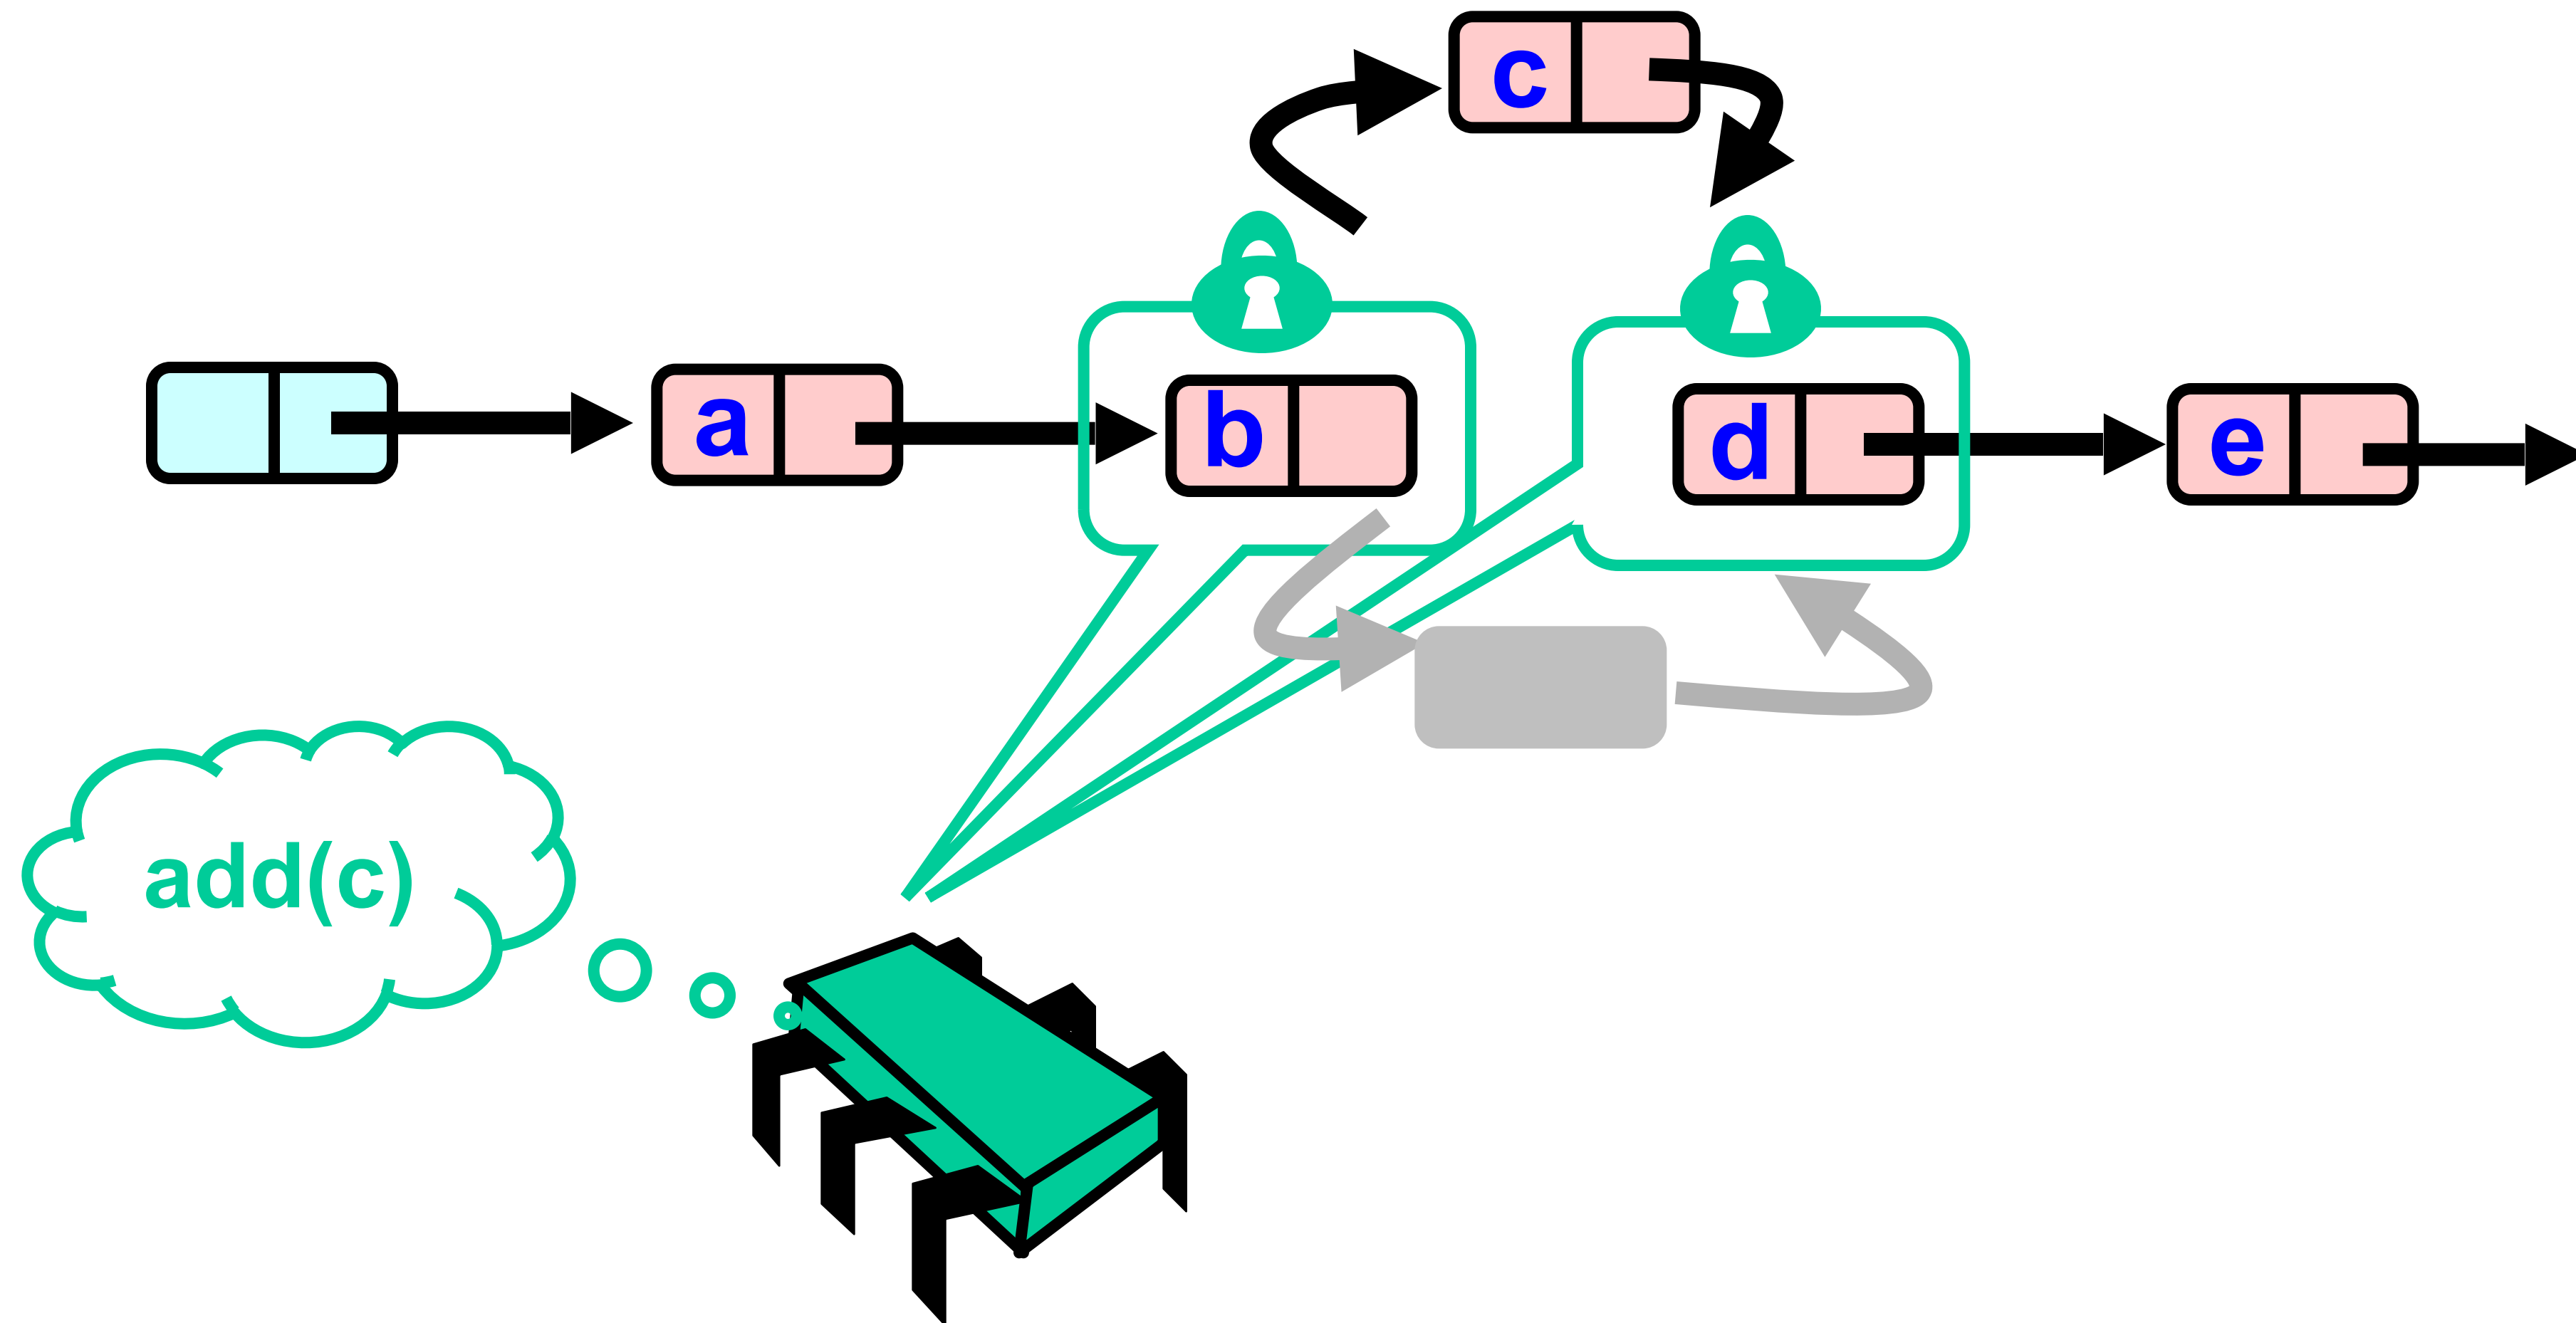
What Else Could Go Wrong?



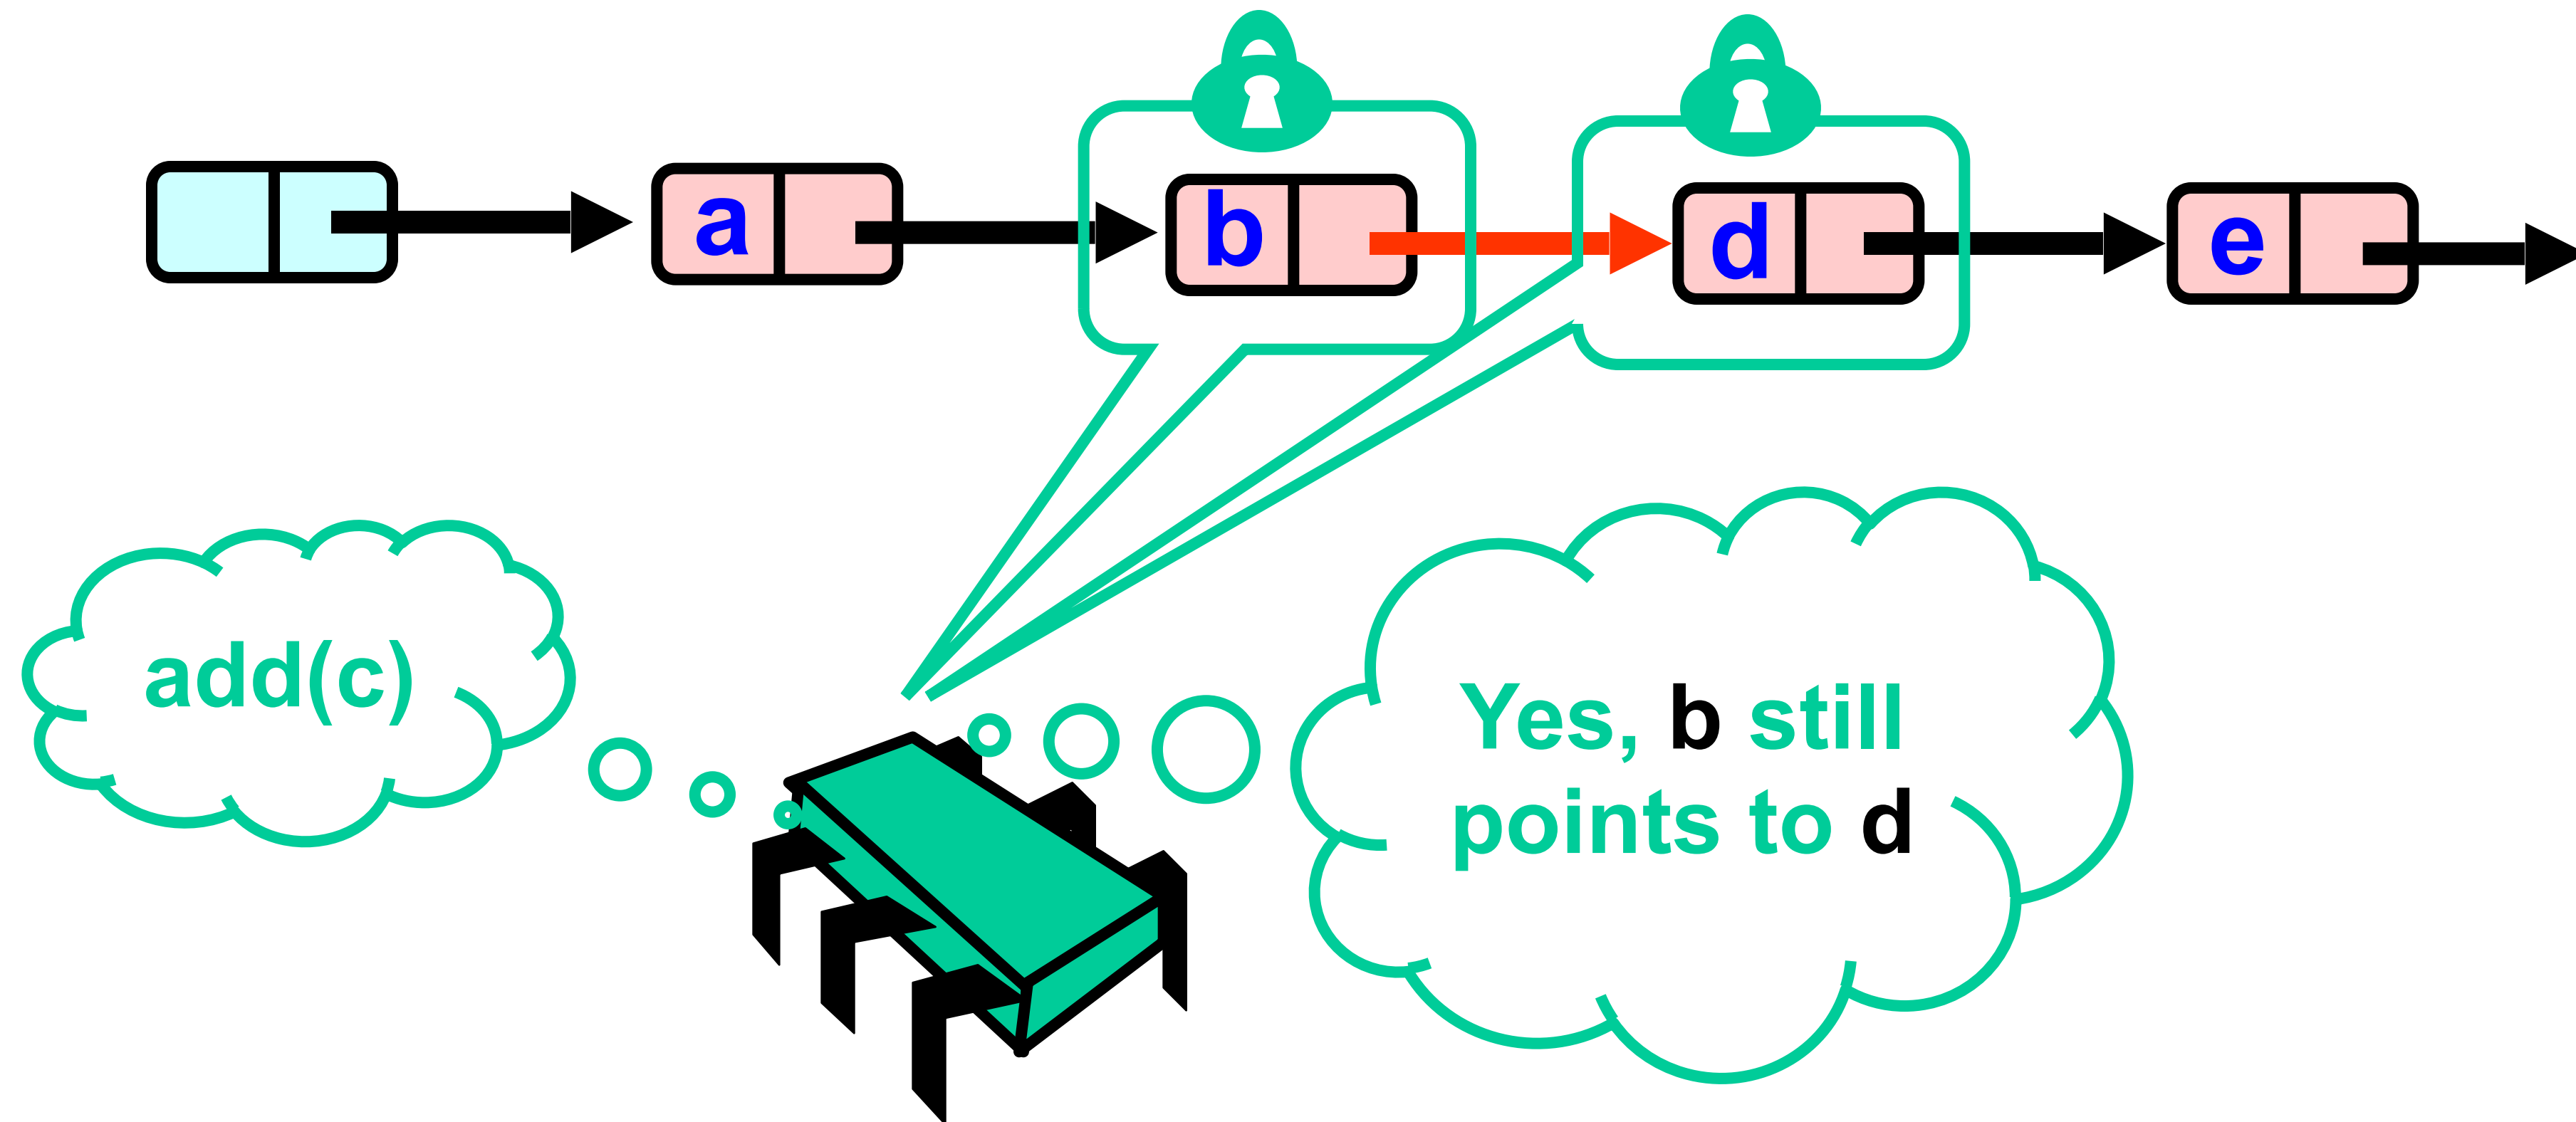
What Else Could Go Wrong?




What Else Could Go Wrong?



Validate Part 2 (while holding locks)



Optimistic: Linearization Point



Same Abstraction Map

- $S(\text{head}) =$
 $\{ x \mid \text{there exists } a \text{ such that}$
 - $a \text{ reachable from head and}$
 - $a.\text{item} = x$ $\}$


Invariants

- Careful: we may traverse deleted nodes
- But we establish properties by
 - Validation
 - After we lock target nodes

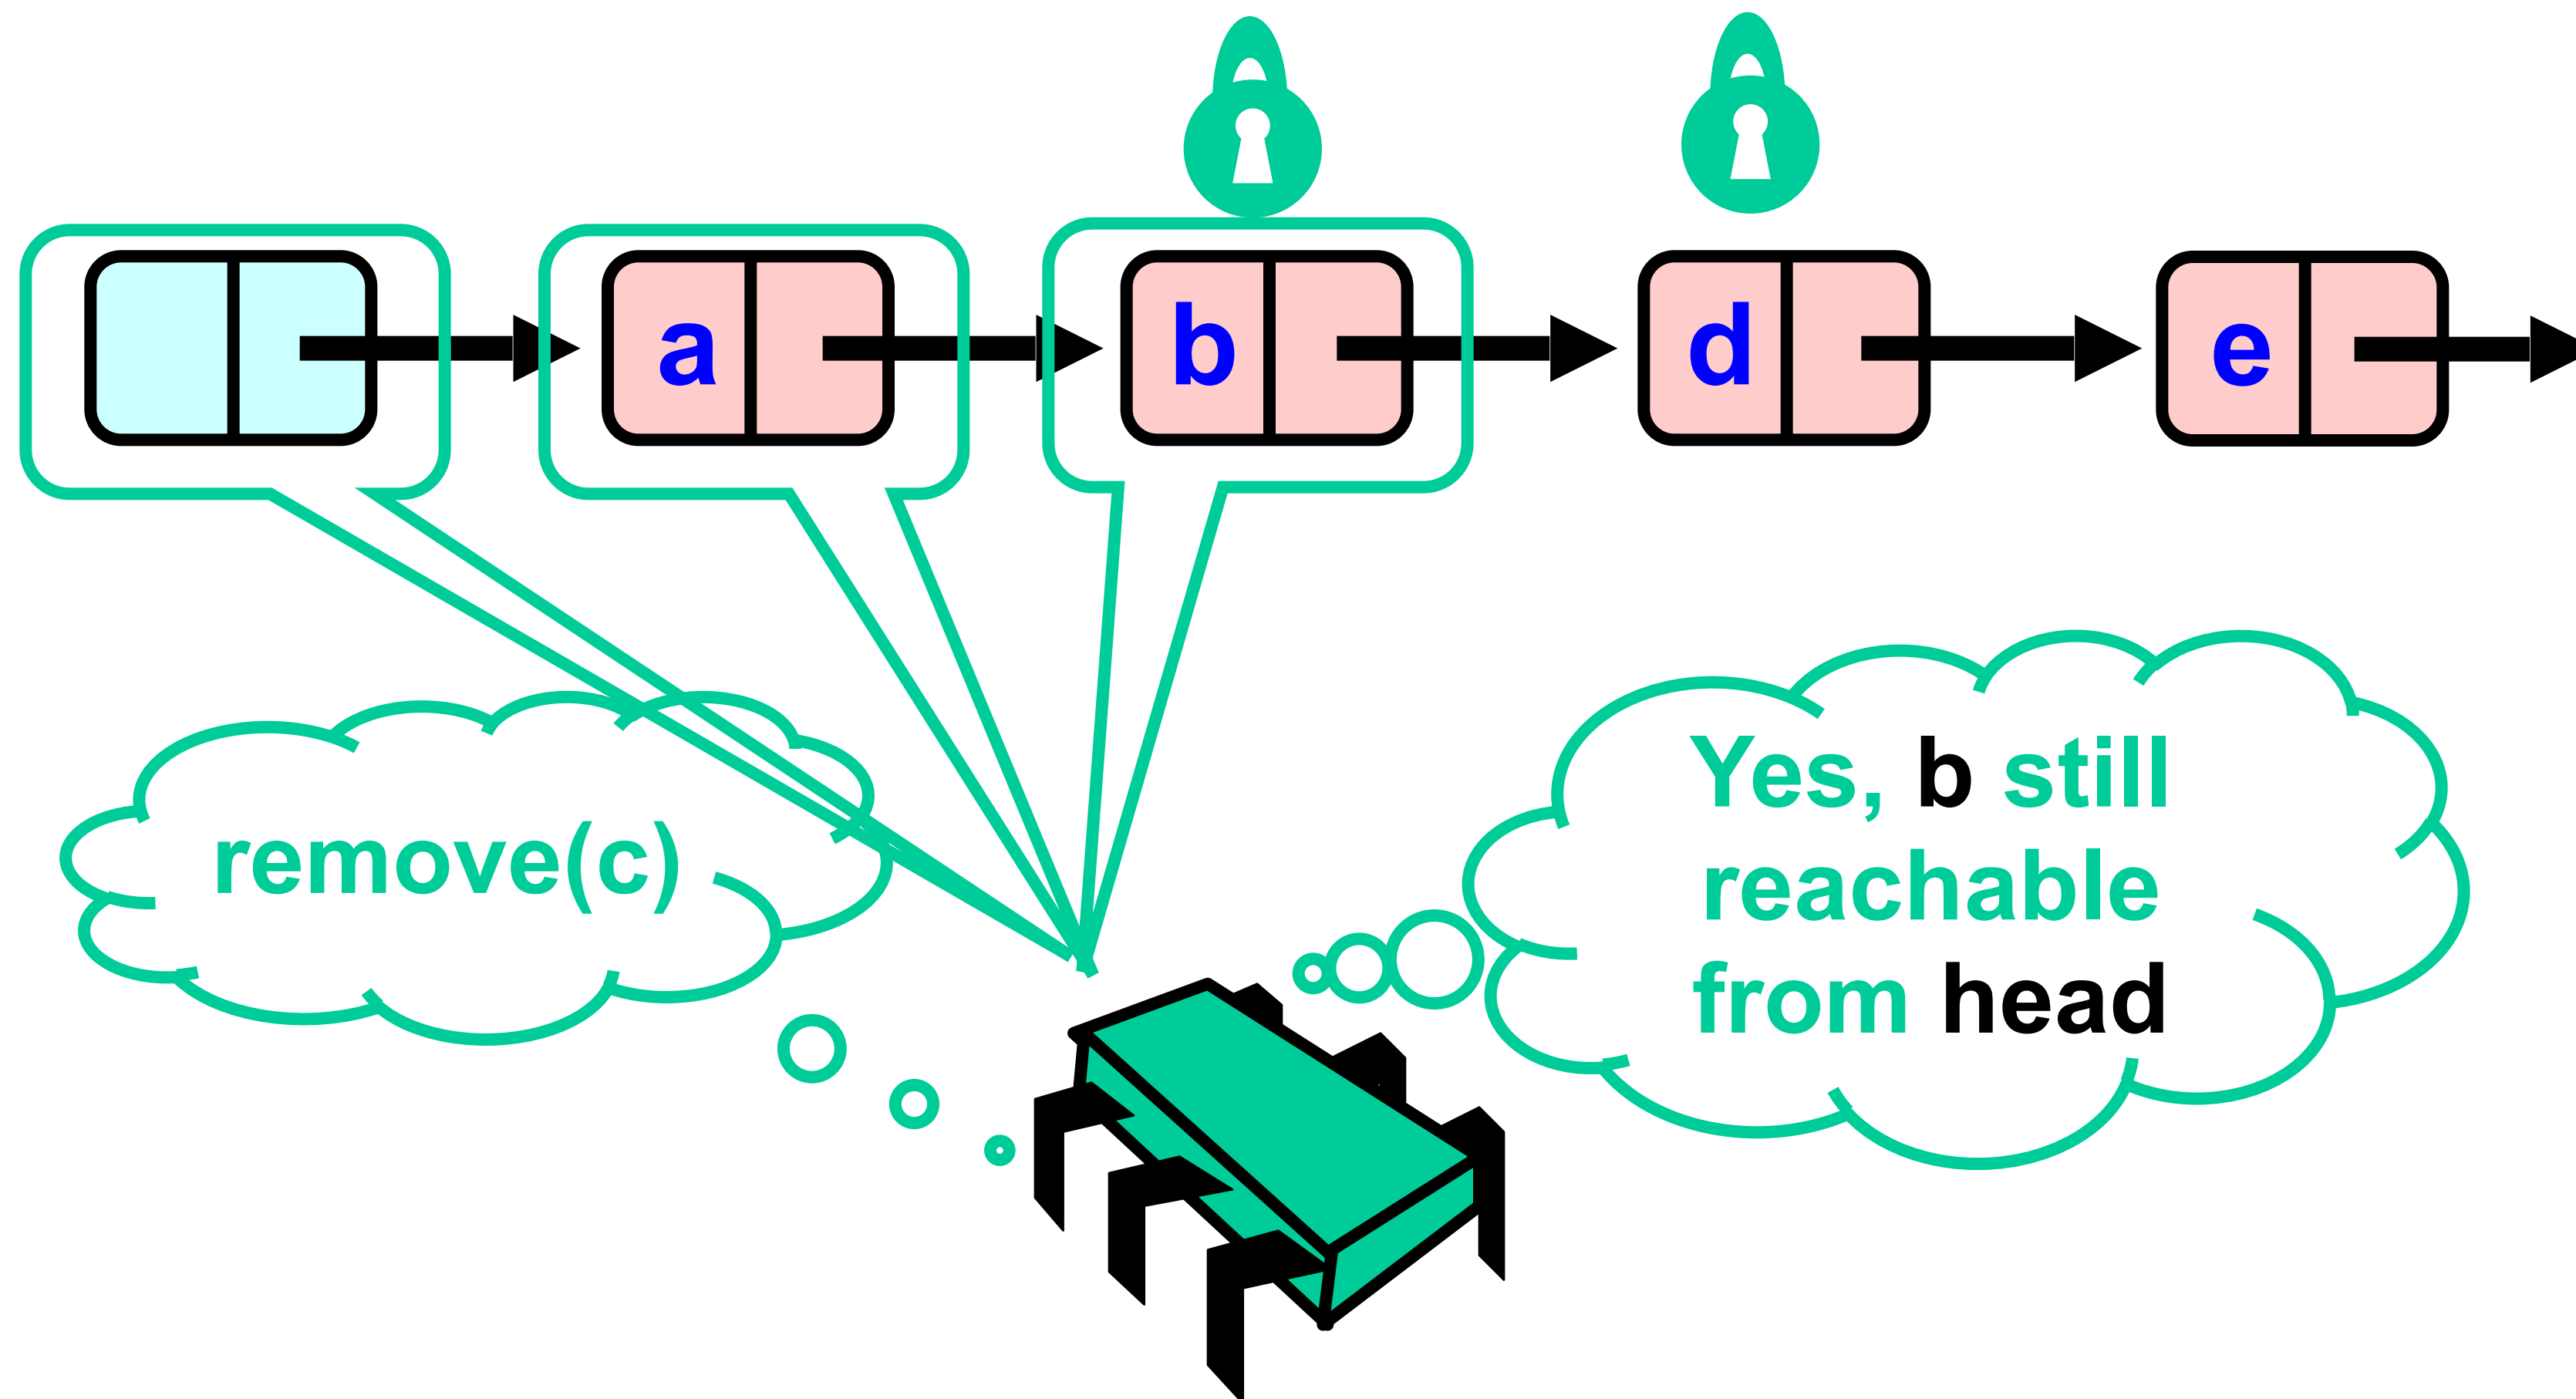
Removal

- If
 - Nodes **b** and **c** both locked
 - Node **b** still accessible
 - Node **c** still successor to **b**
- Then
 - Neither will be deleted
 - OK to delete and return **true**

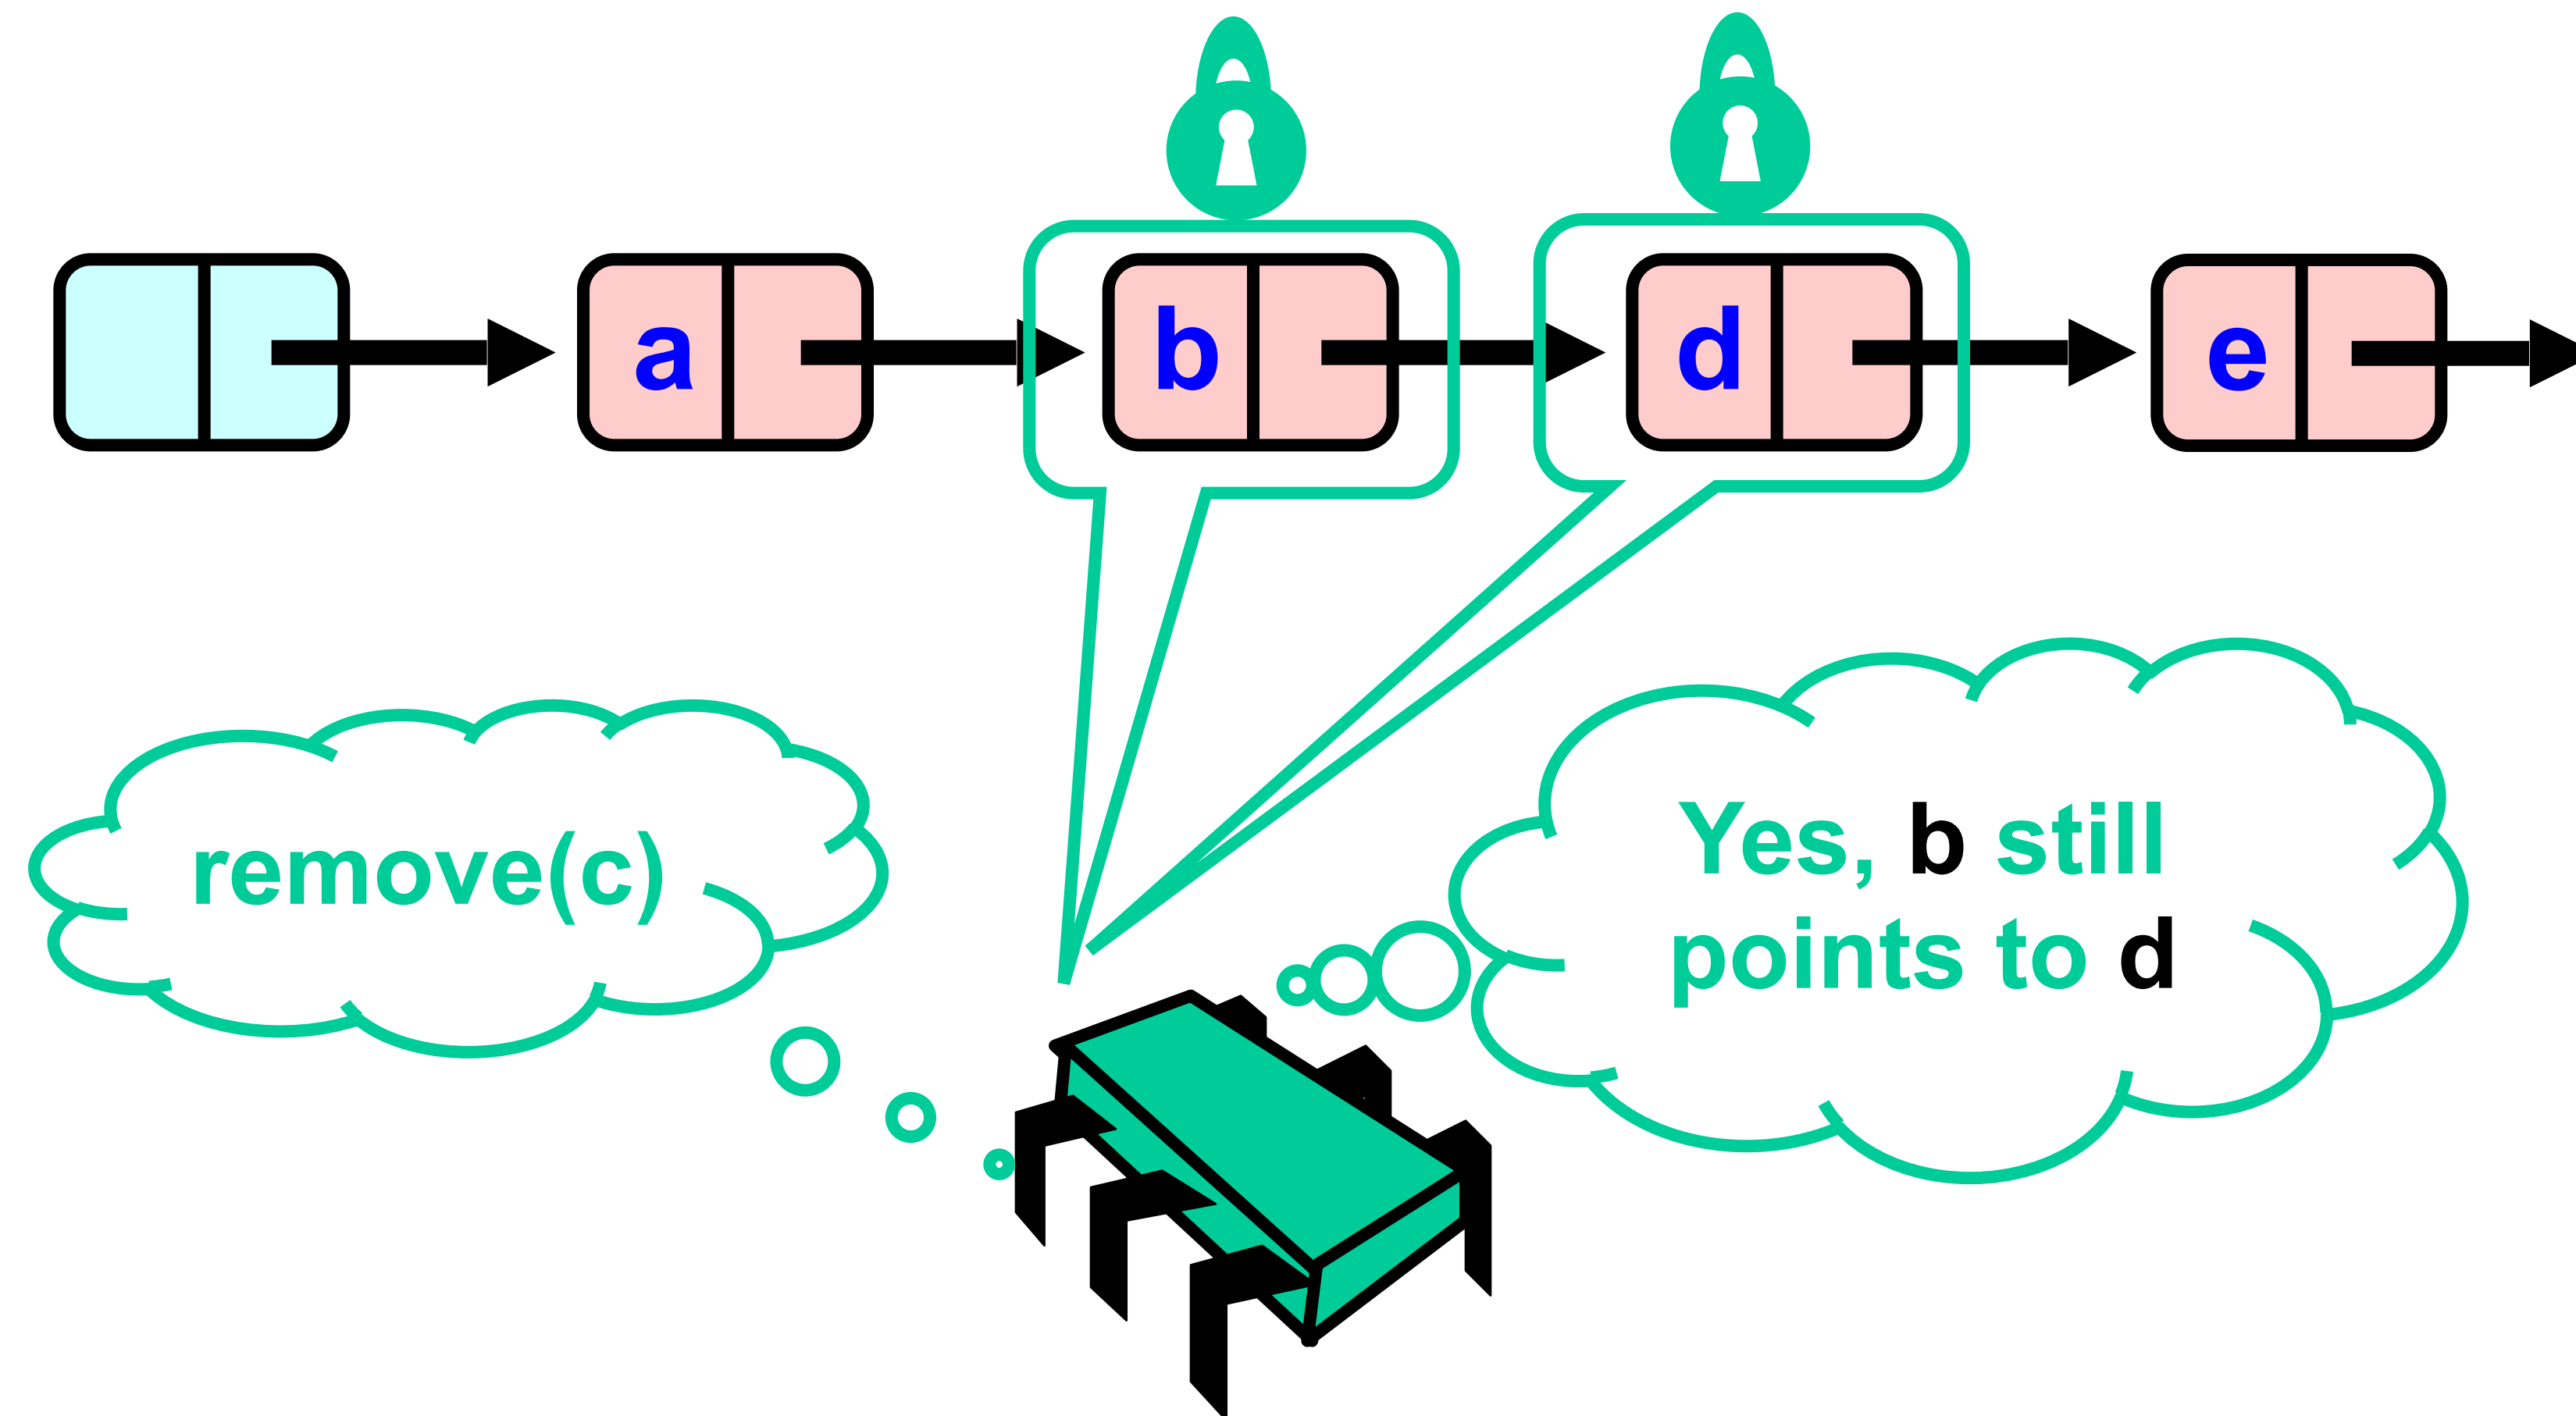
Unsuccessful Remove




Validate (1)



Validate (2)



OK Computer



Correctness

- If
 - Nodes **b** and **d** both locked
 - Node **b** still accessible
 - Node **d** still successor to **b**
- Then
 - Neither will be deleted
 - No thread can add **c** after **b**
 - OK to return false


Validation

```
def validate(pred: Node, curr: Node): Boolean = {  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```


Validation

```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```


**Predecessor &
current nodes**




Validation

```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```

**Start at the
beginning**




Validation



```
def validate(pred: Node, curr: Node):  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

Search range of keys


Validation



```
def validate(pred: Node, curr: Node):  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

Predecessor reachable

Validation




```
def validate(pred: Node, curr: Node):  
    var entry = head  
    while (entry.key <= pred.key) {  
        // Checking for reference equality  
        if (entry eq pred) {  
            return pred.next eq curr  
        }  
        entry = entry.next  
    }  
    false  
}
```

Is current node next?

Validation

Otherwise move on


```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```



Validation


Predecessor not reachable

```
def validate(pred: Node, curr: Node): Boolean = {  
  var entry = head  
  while (entry.key <= pred.key) {  
    // Checking for reference equality  
    if (entry eq pred) {  
      return pred.next eq curr  
    }  
    entry = entry.next  
  }  
  false  
}
```




Remove: searching

```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```



Remove: searching


```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```



Search key

Remove: searching


```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```



**Loop until no synchronization conflict
(see the code further)**

Remove: searching

```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```




Examine predecessor and current nodes

Remove: searching

```
def remove(item: T): Boolean = {  
  val key = item.hashCode()  
  while (true) {  
    var pred = this.head  
    var curr = pred.next  
    while (curr.key < key) {  
      pred = curr  
      curr = curr.next  
    }  
    ...  
  }  
}
```

Search by key



On Exit from While-True-Loop

- If item is present
 - curr holds item
 - pred just before curr
- If item is absent
 - curr has first higher key
 - pred just before curr
- Assuming no synchronization problems

Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) { // present in list
            pred.next = curr.next
            return true
        } else { // not present in list
            return false
        }
    }
} finally { // always unlock
    pred.unlock(); curr.unlock()
}
```

Remove Method

```
pred.lock(); curr.lock()
```

```
try {  
    if (validate(pred, curr)) {  
        if (curr.key == key) {  
            pred.next = curr.next  
            return true  
        } else {  
            return false  
        }  
    }  
} finally {  
    pred.unlock(); curr.unlock()  
}
```

Lock both nodes

Remove Method

```
pred.lock(); curr.lock()
try {
    if (!validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally { // always unlock
    pred.unlock(); curr.unlock()
}
```

Always unlock

Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally {
    pred.unlock(); curr.unlock()
}
```

Check for synchronization conflicts

Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally {
    pred.unlock(); curr.unlock()
}
```

**target found, remove
node**

Remove Method

```
pred.lock(); curr.lock()
try {
    if (validate(pred, curr)) {
        if (curr.key == key) {
            pred.next = curr.next
            return true
        } else {
            return false
        }
    }
} finally {
    pred.unlock(); curr.unlock()
}
```

target not found

Optimistic List

- Limited hot-spots
 - Targets of `add()`, `remove()`, `contains()`
 - No contention on traversals
- Moreover
 - Traversals are wait-free
 - Food for thought ...

So Far, So Good

- Much less lock acquisition/release
 - Performance
 - Concurrency
- Problems
 - Need to traverse list twice
 - `contains()` method acquires locks

Evaluation

- Optimistic is effective if
 - cost of scanning twice without locks is less than
 - cost of scanning once with locks
- Drawback
 - `contains()` acquires locks
 - 90% of calls in many apps

Demo: Benchmarking Optimistic Lists

<*A good place to pause*>


Lazy List

- Like optimistic, except
 - Scan once
 - `contains(x)` never locks ...
- Key insight
 - Removing nodes causes trouble
 - Do it “lazily”


Lazy List

- **remove ()**
 - Scans list (as before)
 - Locks predecessor & current (as before)
- Logical delete
 - Marks current node as removed (new!)
- Physical delete
 - Redirects predecessor's next (as before)


Lazy Removal



Lazy Removal




Lazy Removal




Logically deleted

Lazy Removal



Lazy Removal




Lazy List

- All Methods
 - Scan through locked and marked nodes
 - Removing a node doesn't slow down other method calls ...
- Must still lock pred and curr nodes.


Validation

- No need to rescan list!
- Check that `pred` is not marked
- Check that `curr` is not marked
- Check that `pred` points to `curr`


Business as Usual




Business as Usual




Business as Usual




Business as Usual




Business as Usual




Business as Usual




Business as Usual



Business as Usual



Business as Usual



New Abstraction Map

- $S(\text{head}) =$
 $\{ x \mid \text{there exists node } a \text{ such that}$
 - a reachable from head and
 - $a.\text{item} = x$ and
 - a is unmarked $\}$

Invariant

- If not marked then item in the set
- and is reachable from head
- and if not yet traversed it is reachable from pred

Validation

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```

List Validate Method

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```

**Predecessor not
Logically removed**

List Validate Method

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```



**Current not
Logically removed**

List Validate Method

```
def validate(pred: Node, curr: Node) =  
    !pred.marked &&  
    !curr.marked &&  
    (pred.next eq curr)
```

**Predecessor still
Points to current**

Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true  
            pred.next = curr.next  
            return true;  
        } else {  
            return false  
        }  
    }  
} finally {  
    pred.unlock()  
    curr.unlock()  
}
```

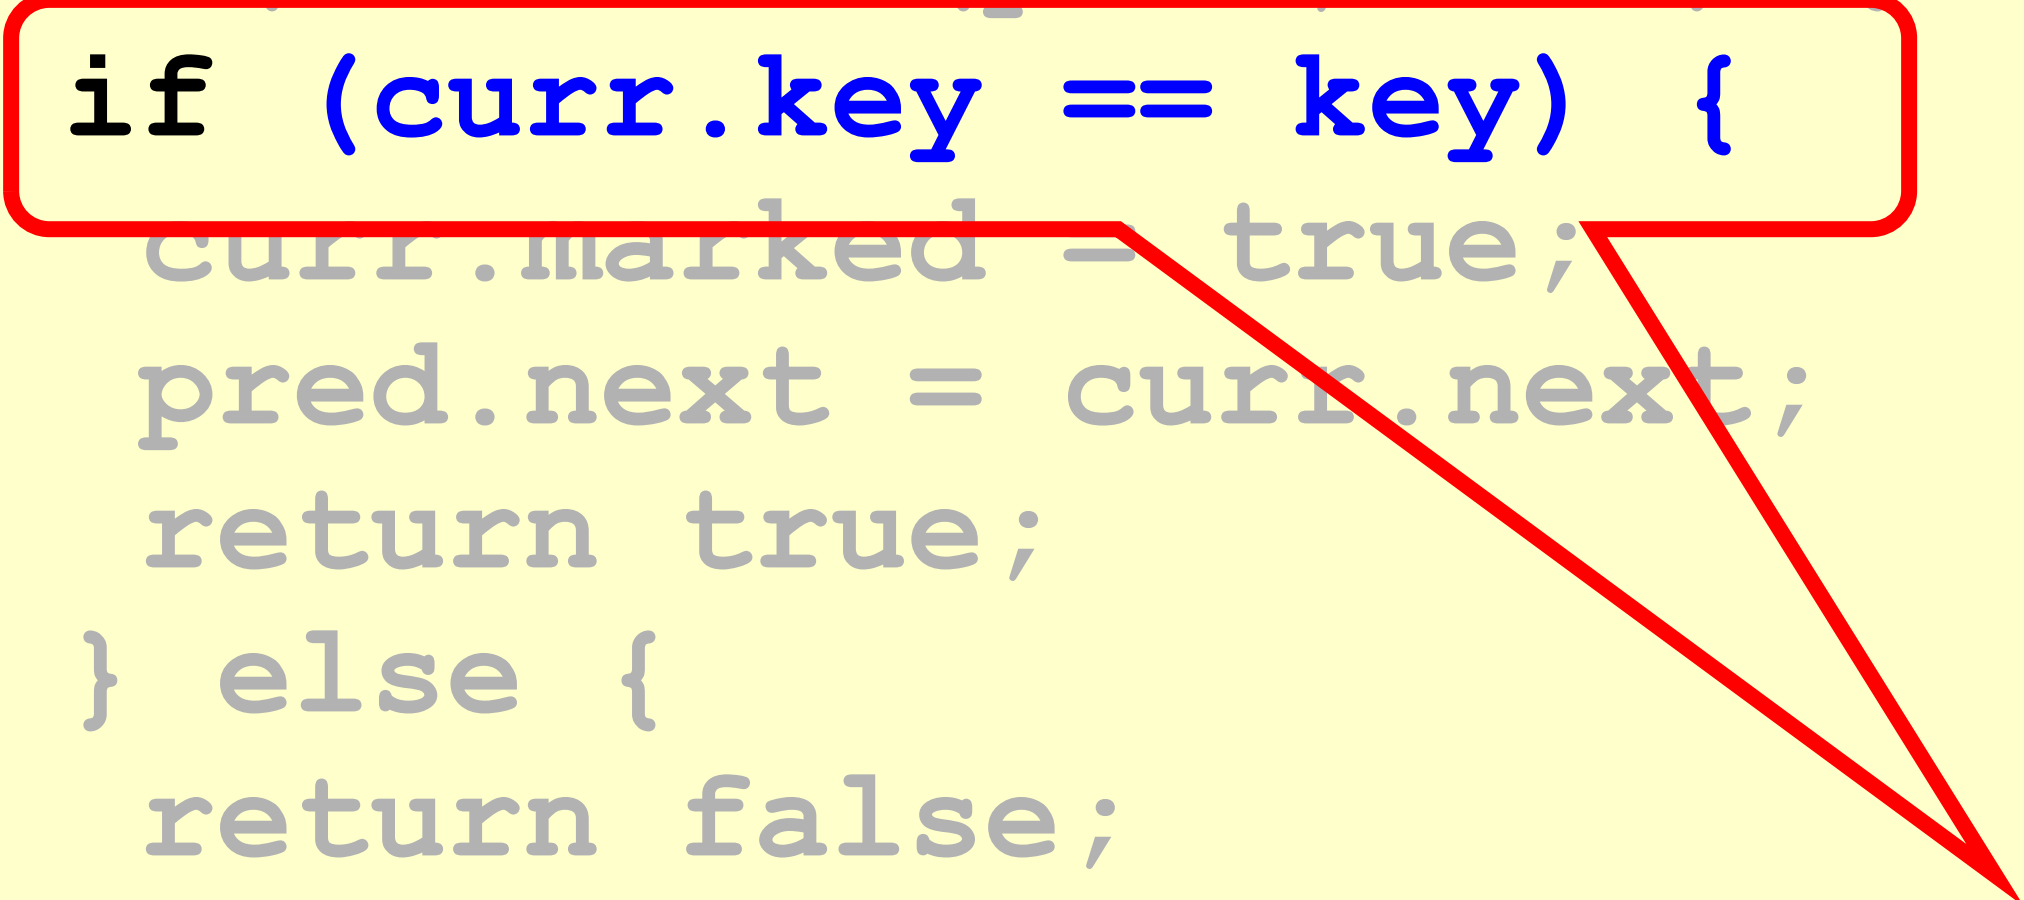
Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true  
            pred.next = curr.next  
            return true  
        } else {  
            return false  
        }  
    }  
    finally {  
        pred.unlock()  
        curr.unlock()  
    }  
}
```

Validate as before

Remove

```
try {  
    pred.lock(); curr.lock();  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true;  
            pred.next = curr.next;  
            return true;  
        } else {  
            return false;  
        }  
    }  
} finally {  
    pred.unlock();  
    curr.unlock();  
}
```



Key found

Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true;  
            pred.next = curr.next  
            return true  
        } else {  
            return false  
        }  
    } finally {  
        pred.unlock()  
        curr.unlock()  
    }  
}
```

Logical remove

Remove

```
try {  
    pred.lock(); curr.lock()  
    if (validate(pred, curr) {  
        if (curr.key == key) {  
            curr.marked = true  
            pred.next = curr.next;  
            return true  
        } else {  
            return false  
        }  
    }  
} finally {  
    pred.unlock()  
    curr.unlock()  
}
```

physical remove

Contains

```
def contains(item: T) = {  
    val key = item.hashCode  
    var curr = this.head  
    while (curr.key < key) curr = curr.next  
    curr.key == key && !curr.marked  
}
```

Contains

```
def contains(item: T) = {  
    val key = item.hashCode  
    var curr = this.head  
    while (curr.key < key) curr = curr.next  
    curr.key == key && !curr.marked  
}
```

Start at the head

Contains

```
def contains(item: T) = {  
  val key = item.hashCode  
  var curr = this.head  
  while (curr.key < key) curr = curr.next  
  curr.key == key && !curr.marked  
}
```

Search key range

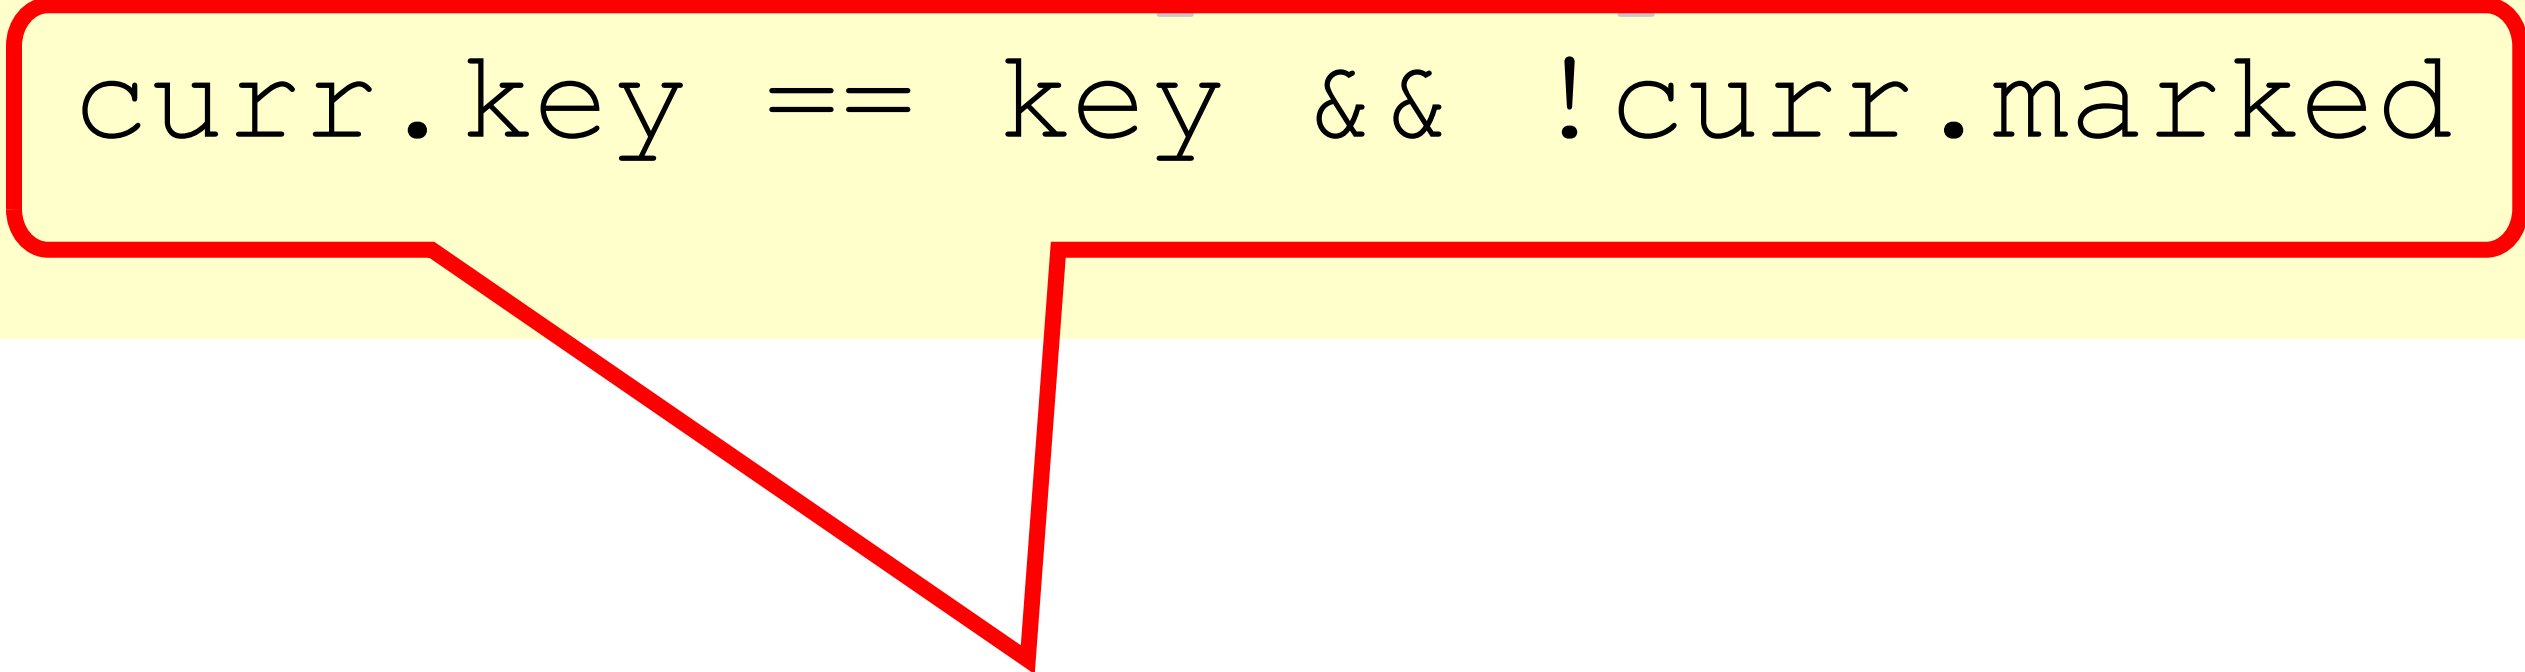
Contains

```
def contains(item: T) = {  
  val key = item.hashCode  
  var curr = this.head  
  while (curr.key < key) curr = curr.next  
  curr.key == key && !curr.marked  
}
```

Traverse *without locking*
(nodes may have been removed)


Contains

```
def contains(item: T) = {  
  val key = item.hashCode  
  var curr = this.head  
  while (curr.key < key) curr = curr.next  
  curr.key == key && !curr.marked  
}
```



Present and undeleted?


Summary: Wait-free Contains



Use Mark bit + list ordering

1. Not marked \rightarrow in the set
2. Marked or missing \rightarrow not in the set

Lazy List



Lazy add () and remove () + Wait-free contains ()

Evaluation

- Good:
 - **contains ()** doesn't lock
 - In fact, it's wait-free!
 - Good because typically high % contains()
 - Uncontended calls don't re-traverse
- Bad
 - Contended **add ()** and **remove ()** calls must re-traverse
 - Traffic jam if one thread delays

Traffic Jam

- Any concurrent data structure based on mutual exclusion has a weakness
- If one thread
 - Enters critical section
 - And “eats the big muffin”
 - Cache miss, page fault, descheduled ...
 - Everyone else using that lock is stuck!
 - Need to trust the scheduler....

Reminder: Lock-Free Data Structures




- No matter what ...
 - Guarantees minimal progress in any execution
 - i.e. Some thread will always complete a method call
 - Even if others halt at malicious times
 - Implies that implementation can't use locks

Lock-free Lists

- Next logical step
 - Wait-free `contains()`
 - lock-free `add()` and `remove()`
- Use only `compareAndSet()`
 - What could go wrong?

Lock-free Lists

Logical Removal




Use CAS to verify pointer
is correct

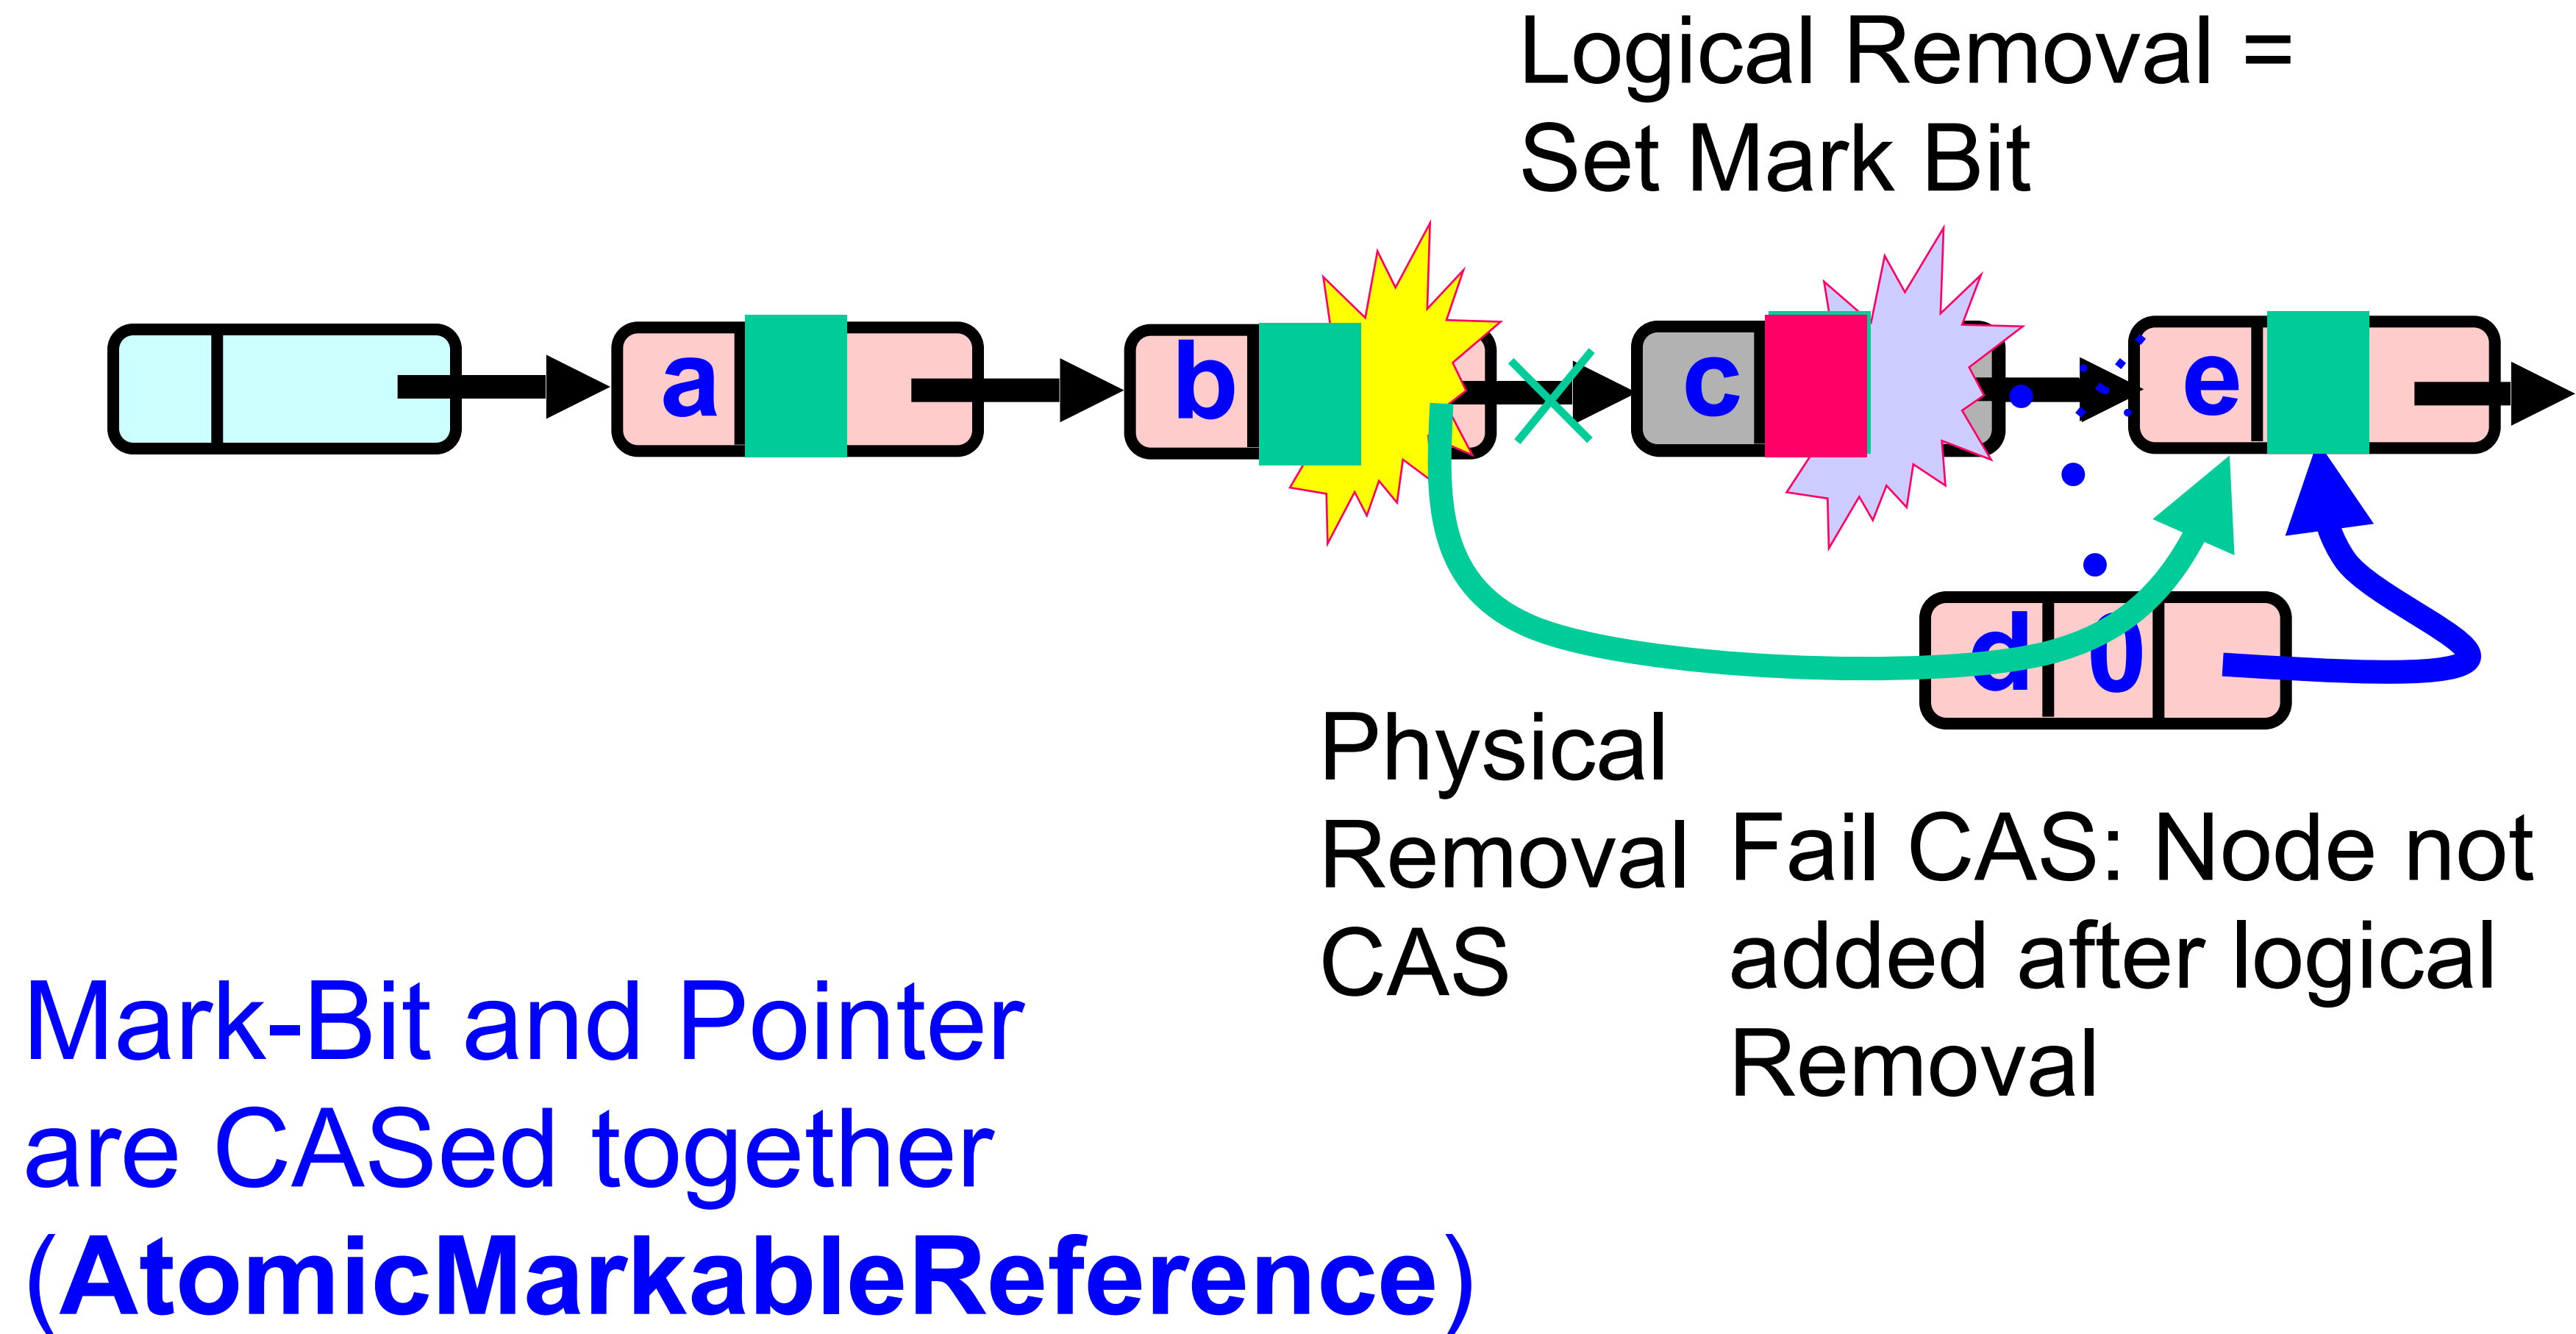
Physical Removal

Not enough!

Problem...



The Solution: Combine Bit and Pointer




Solution

- Use AtomicMarkableReference
- Atomically
 - Swing reference and
 - Update flag
- Remove in two steps
 - Set mark bit in next field
 - Redirect predecessor's pointer

Marking a Node


- **AtomicMarkableReference** class
 - Java.util.concurrent.atomic package




Performance

- Different list-based set implementations
- 16-node machine
- Vary percentage of **contains()** calls

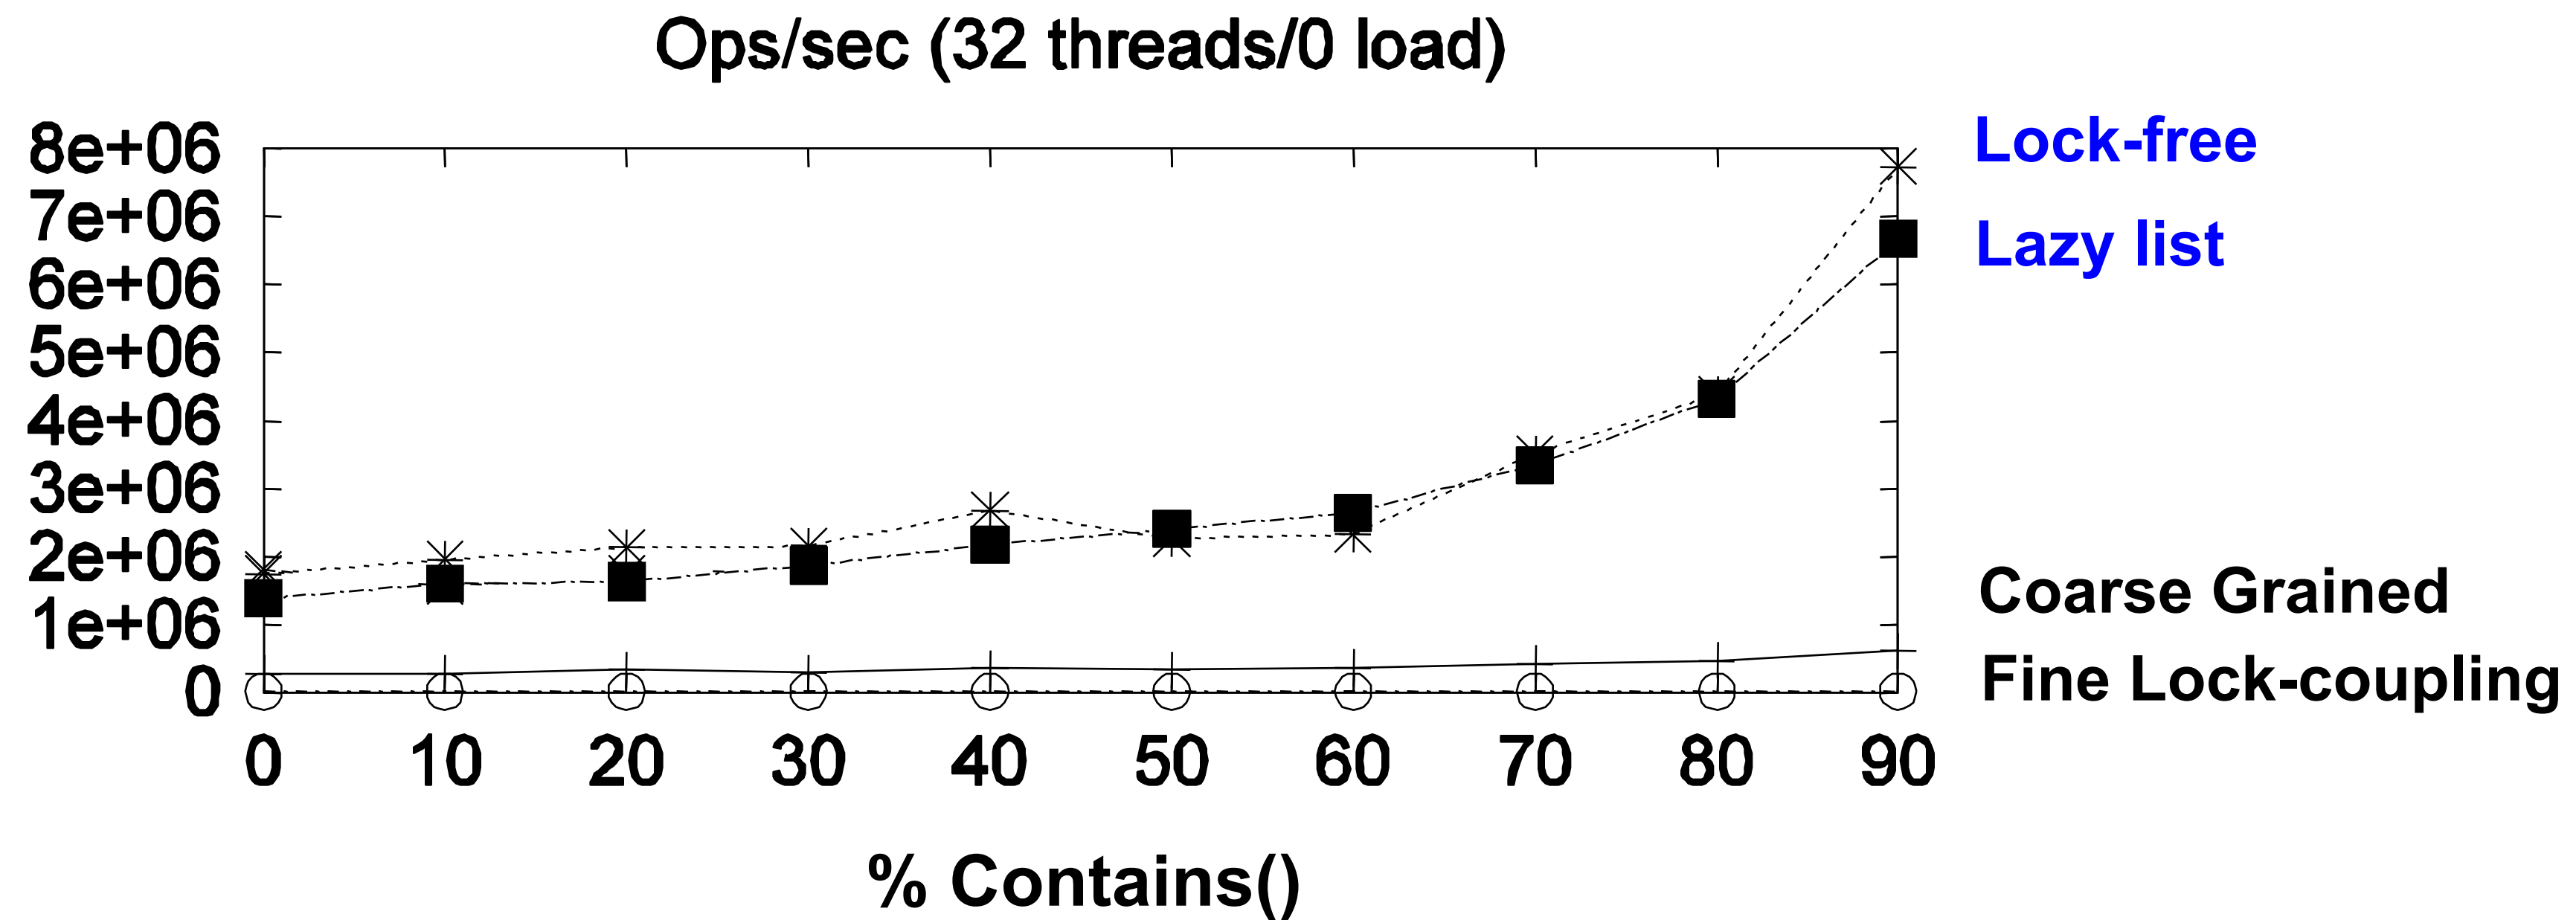
High Contains Ratio



Low Contains Ratio



As Contains Ratio Increases



Summary

- Coarse-grained locking
- Fine-grained locking (“hand-over-hand”)
- Optimistic synchronization
- Lazy synchronization
- Lock-free synchronization

“To Lock or Not to Lock”

- Locking vs. Non-blocking:
 - Extremist views on both sides
 - Locking: longs waits
 - Non-blocking: long “clean-ups”
- The answer: nobler to compromise
 - Example: Lazy list combines blocking `add()` and `remove()` and a wait-free `contains()`
 - Remember: Blocking/non-blocking is a property of a method



This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](https://creativecommons.org/licenses/by-sa/2.5/).

- **You are free:**
 - **to Share** — to copy, distribute and transmit the work
 - **to Remix** — to adapt the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
 - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
 - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.