

# YSC3242: Parallel, Concurrent and Distributed Programming

Concurrent Objects  
Part 2

# Formal Model of Executions

- Define precisely what we mean
  - Ambiguity is bad when intuition is weak
- Allow reasoning
  - Formal
  - But mostly informal
    - In the long run, actually more important

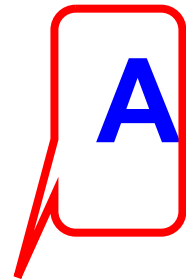
# Split Method Calls into Two Events

- Invocation
  - method name & args
  - `q.enq(x)`
- Response
  - result or exception
  - `q.enq(x)` returns `void`
  - `q.deq()` returns `x`
  - `q.deq()` throws `empty`

# Invocation Notation

**A q.enq(x)**

# Invocation Notation

 **A.q.enq(x)**

**thread**


# Invocation Notation

**A.q.enqueue(x)**


**thread**                      **method**

The diagram illustrates the components of the invocation notation **A.q.enqueue(x)**. The expression is written in blue. The **A** is enclosed in a red rounded rectangle, with a red line pointing from the word **thread** (in red) below it. The **q.enqueue(x)** is enclosed in another red rounded rectangle, with a red line pointing from the word **method** (in red) below it.

# Invocation Notation



# Invocation Notation

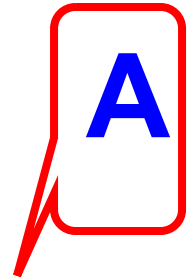




# Response Notation

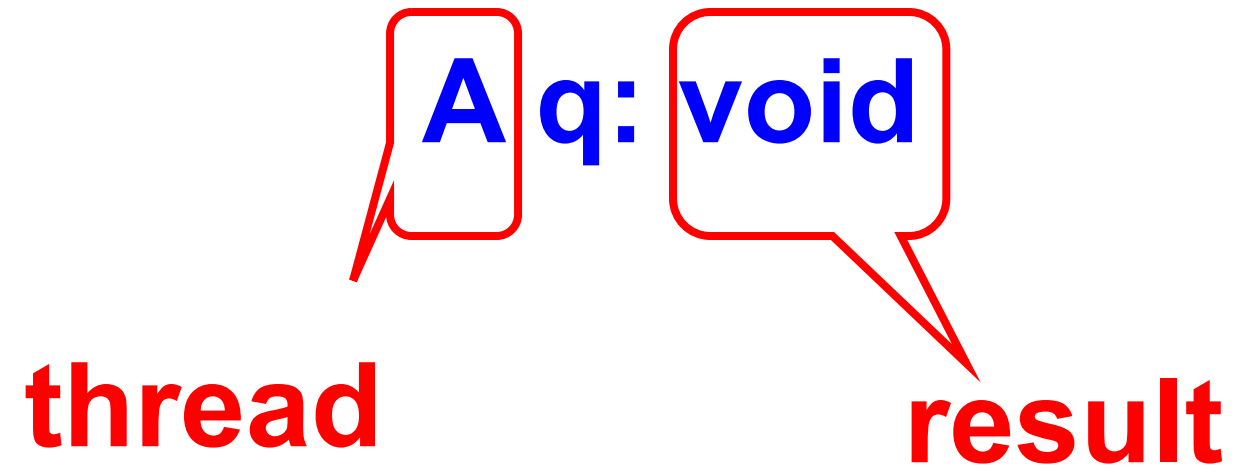
**A q: void**

# Response Notation

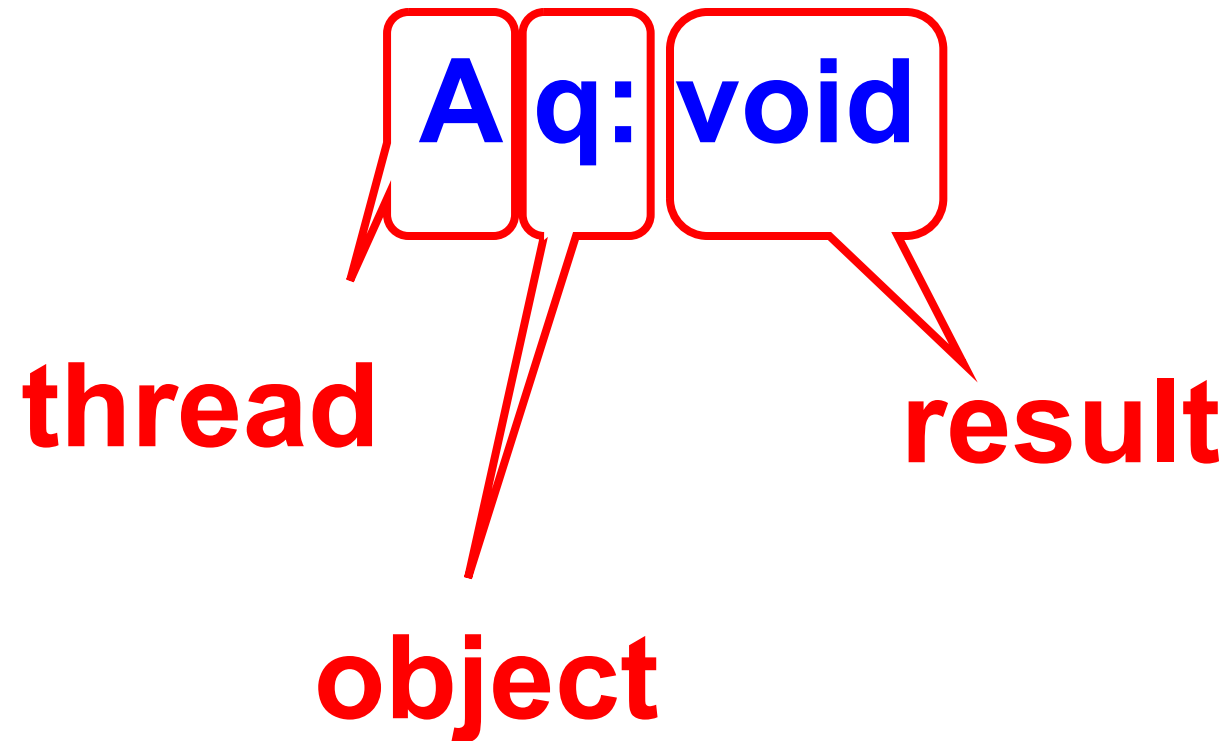
 **Aq: void**

**thread**

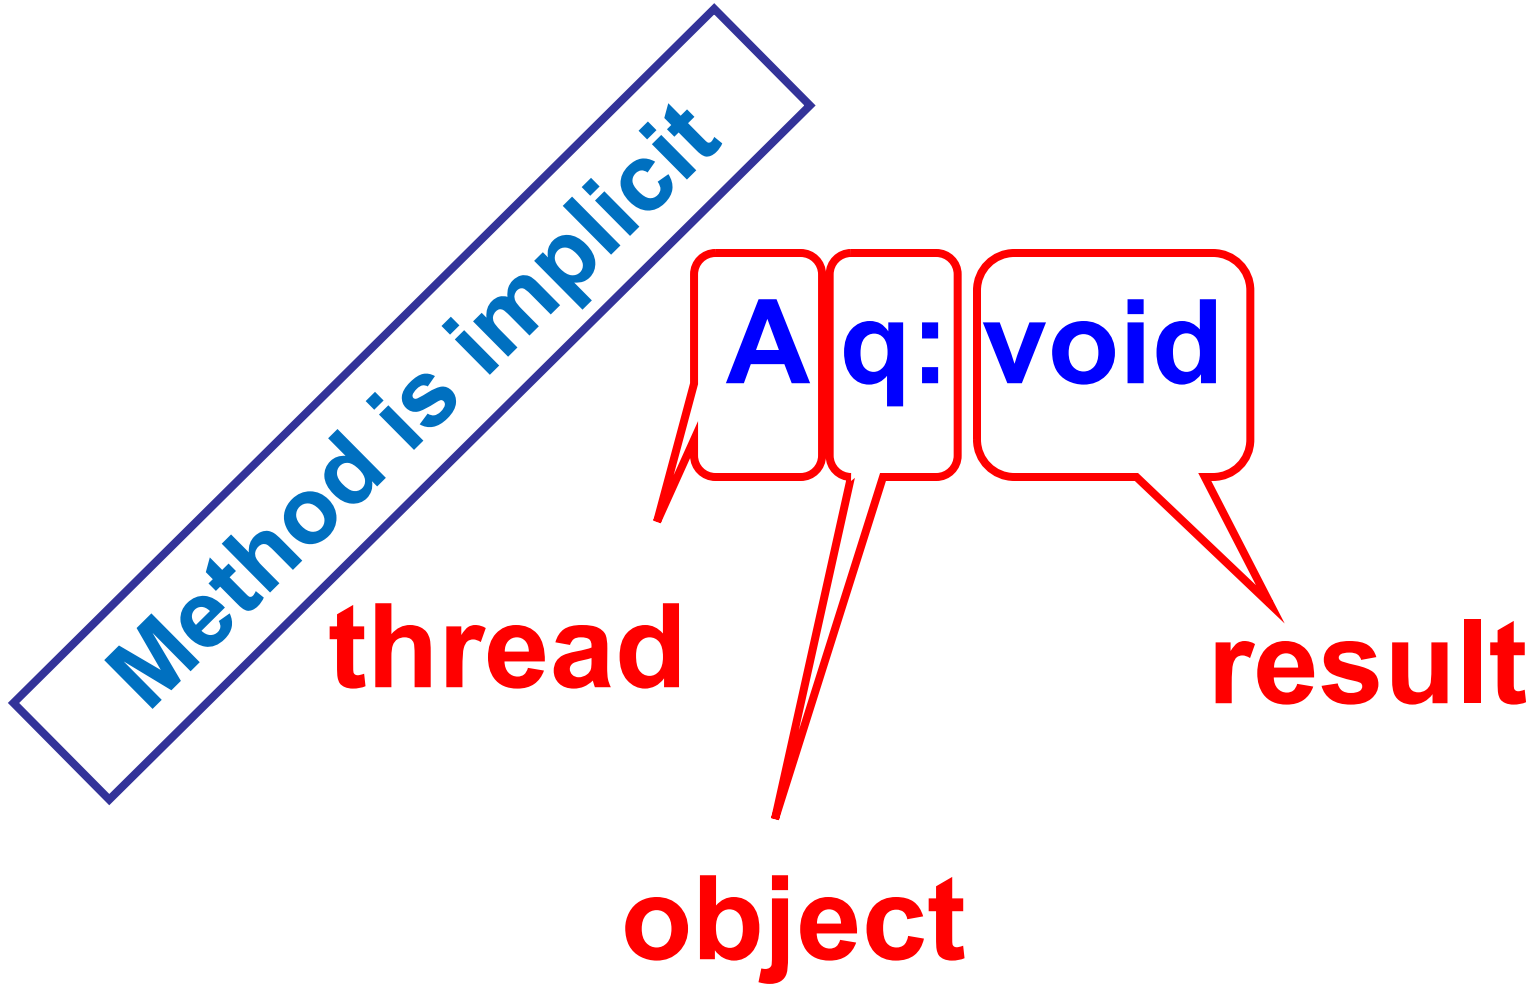
# Response Notation




# Response Notation



# Response Notation



# Response Notation



# History - Describing an Execution


**H =**

- A q.enq(3)
- A q:void
- A q.enq(5)
- B p.enq(4)
- B p:void
- B q.deq()
- B q:3

**Sequence of  
invocations and  
responses**

# Definition

- Invocation & response *match* if





# Object Projections

$H =$

- A `q.enqueue(3)`
- A `q: void`
- B `p.enqueue(4)`
- B `p: void`
- B `q.dequeue()`
- B `q: 3`

# Object Projections

A `q.enqueue(3)`

A `q:void`

$H|q =$

B `q.dequeue()`

B `q:3`

# Thread Projections

$H =$

- A `q.enqueue(3)`
- A `q: void`
- B `p.enqueue(4)`
- B `p: void`
- B `q.dequeue()`
- B `q: 3`

# Thread Projections

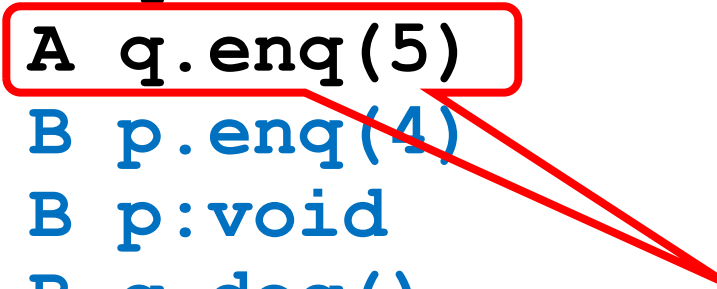
$H|B =$

- $B \text{ p.enq}(4)$
- $B \text{ p:void}$
- $B \text{ q.deq}()$
- $B \text{ q:3}$

# Complete Subhistory

**H =**

```
A q.enq(3)
A q:void
A q.enq(5)
B p.enq(4)
B p:void
B q.deq()
B q:3
```



**An invocation is  
*pending* if it has no  
matching response**

# Complete Subhistory

**H =**

A	q.enq(3)
A	q:void
A	q.enq(5)
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3

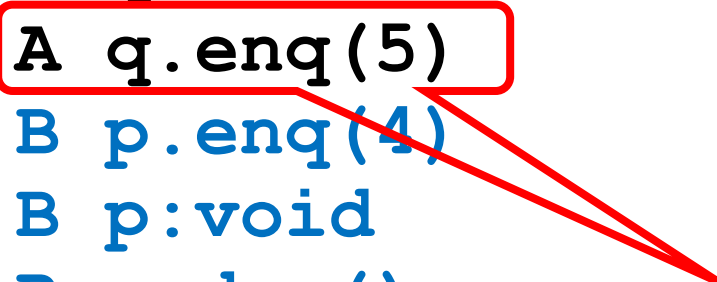
May or may not have taken effect

# Complete Subhistory

**H =**

```
A q.enq(3)
A q:void
A q.enq(5)
B p.enq(4)
B p:void
B q.deq()
B q:3
```

**discard pending  
invocations**



# Complete Subhistory

A q.enq(3)  
A q:void

**Complete(H) =** B p.enq(4)  
B p:void  
B q.deq()  
B q:3



# Sequential Histories

A q.enq(3)

A q:void

B p.enq(4)

B p:void

B q.deq()

B q:3

A q:enq(5)

# Sequential Histories

A q.enq(3)

A q:void

**match**

B p.enq(4)


B p:void

B q.deq()


B q:3

A q:enq(5)

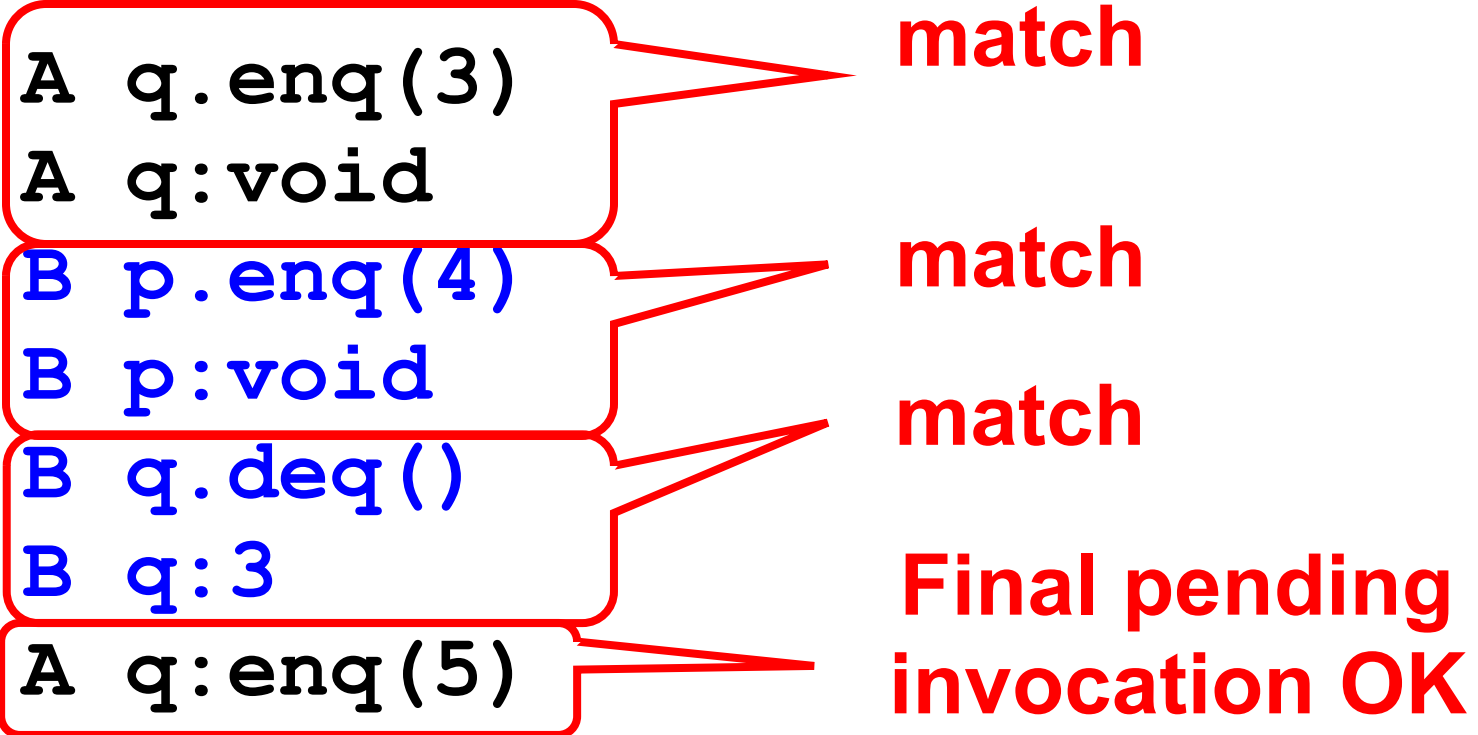
# Sequential Histories



# Sequential Histories



# Sequential Histories



# Sequential Histories

A q.enq(3)  
A q:void

**match**

B p.enq(4)  
B p:void

**match**

B q.deq()  
B q:3

**match**

A q:enq(5)

**Final pending  
invocation OK**

Method calls of different  
threads do not interleave

# Well-Formed Histories

H=

- A q.enq(3)
- B p.enq(4)
- B p:void
- B q.deq()
- A q:void
- B q:3

# Well-Formed Histories

Per-thread projections  
sequential

$H =$

- A q.enq(3)
- B p.enq(4)
- B p:void
- B q.deq()
- A q:void
- B q:3

$H \mid B =$

- B p.enq(4)
- B p:void
- B q.deq()
- B q:3



# Well-Formed Histories

Per-thread projections  
sequential

$H =$

- A q.enq(3)
- B p.enq(4)
- B p:void
- B q.deq()
- A q:void
- B q:3

$H \mid B =$

- B p.enq(4)
- B p:void
- B q.deq()
- B q:3

$H \mid A =$

- A q.enq(3)
- A q:void

# Equivalent Histories

Threads see the same  
thing in both

$$\left\{ \begin{array}{l} H|A = G|A \\ H|B = G|B \end{array} \right.$$

H=

```
A q.enq(3)
B p.enq(4)
B p:void
B q.deq()
A q:void
B q:3
```

G=

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
```

# Sequential Specifications


- A sequential specification is some way of telling whether a
  - Single-thread, single-object history
  - Is legal
- For example:
  - Pre and post-conditions
  - But plenty of other techniques exist ...

# Legal Histories


- A sequential (multi-object) history  $H$  is legal if
  - For every *object*  $x$
  - $H|x$  is in the sequential spec for  $x$
  - *Not talking about threads now!*

# Precedence

A q.enq(3)  
B p.enq(4)  
B p.void  
A q:void  
B q.deq()  
B q:3

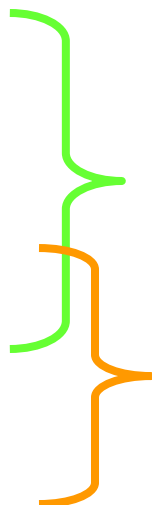


A method call **precedes**  
another if *response* event  
precedes *invocation* event




# Non-Precedence

```
A q.enq(3)
B p.enq(4)
B p.void
B q.deq()
A q:void
B q:3
```


A diagram showing two overlapping method calls. A green curly brace groups the first three lines: 'A q.enq(3)', 'B p.enq(4)', and 'B p.void'. An orange curly brace groups the last three lines: 'B q.deq()', 'A q:void', and 'B q:3'. The two braces overlap, indicating that the two method calls are concurrent.

Some method calls  
**overlap** one another



# Notation

- Given
  - History  $H$
  - method executions  $m_0$  and  $m_1$  in  $H$
- We say  $m_0 \rightarrow_H m_1$ , if
  - $m_0$  precedes  $m_1$
- Relation  $m_0 \rightarrow_H m_1$  is a
  - Partial order
  - Total order if  $H$  is sequential



# Linearizability

- History  $H$  is *linearizable* if it can be extended to  $G$  by
  - Appending zero or more responses to pending invocations
  - Discarding other pending invocations
- So that  $G$  is equivalent to
  - Legal sequential history  $S$
  - where  $\rightarrow_G \subset \rightarrow_S$




# Remarks

- Some pending invocations
  - Took effect, so keep them
  - Discard the rest
- Condition  $\rightarrow_{\mathbf{G}} \subset \rightarrow_{\mathbf{S}}$ 
  - Means that **S** respects “real-time order” of **G**

Ensuring  $\rightarrow_G \subset \rightarrow_S$

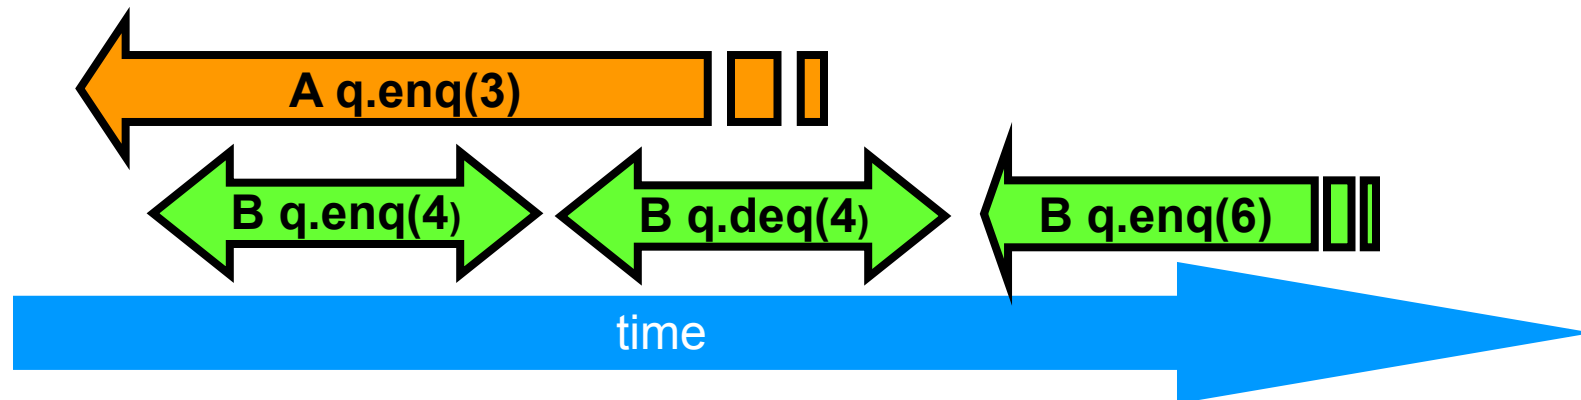
$$\rightarrow_G = \{a \rightarrow c, b \rightarrow c\}$$

$$\rightarrow_S = \{a \rightarrow b, a \rightarrow c, b \rightarrow c\}$$



# Example

A q.enq(3)  
B q.enq(4)  
B q:void  
B q.deq()  
B q:4  
B q:enq(6)



# Example

A q.enq(3)

B q.enq(4)

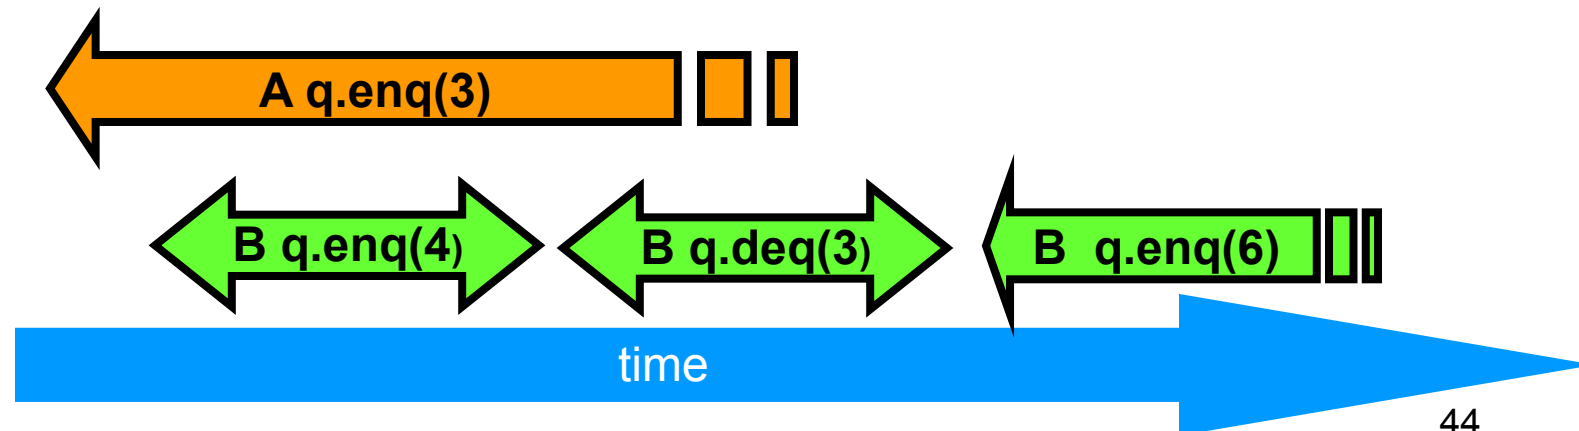
B q:void

B q.deq()

B q:4

B q:enq(6)

Complete this  
pending  
invocation



# Example

A q.enq(3)

B q.enq(4)

B q:void

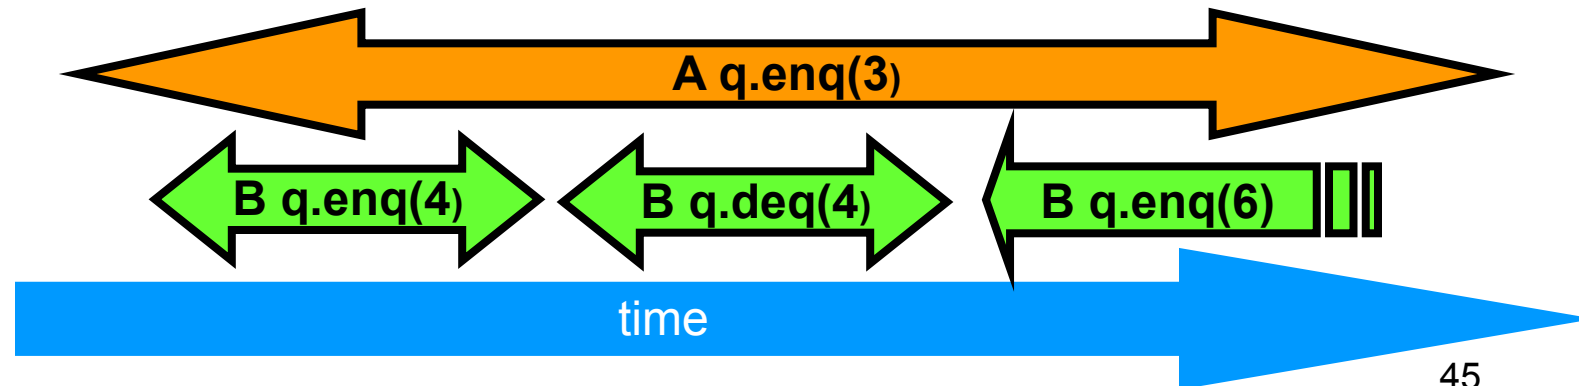
B q.deq()

B q:4

B q:enq(6)

A q:void

Complete this  
pending  
invocation



# Example

discard this one

A q.enq(3)

B q.enq(4)

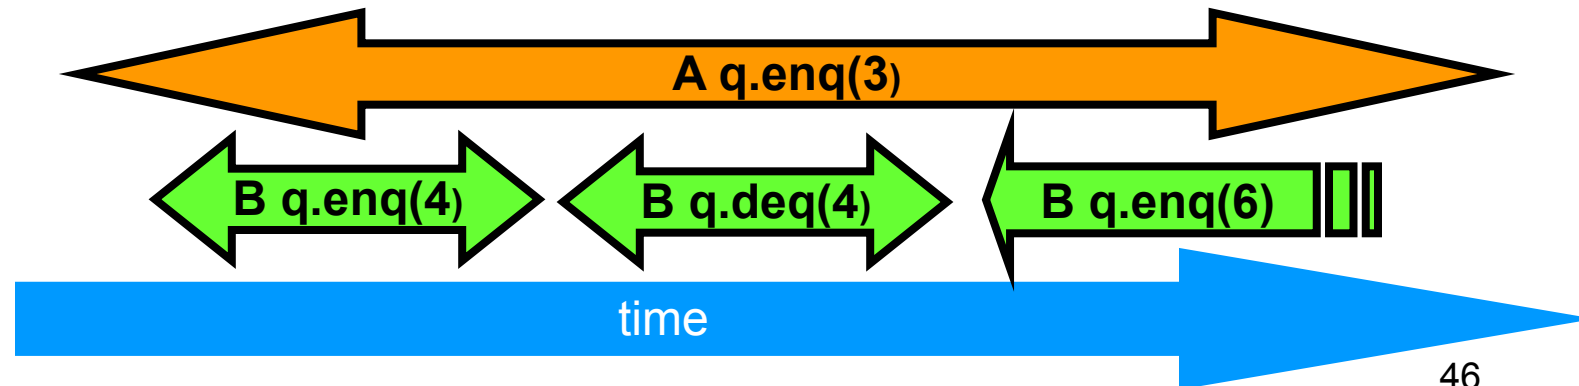
B q:void

B q.deq()

B q:4

B q:enq(6)

A q:void



# Example

A q.enq(3)

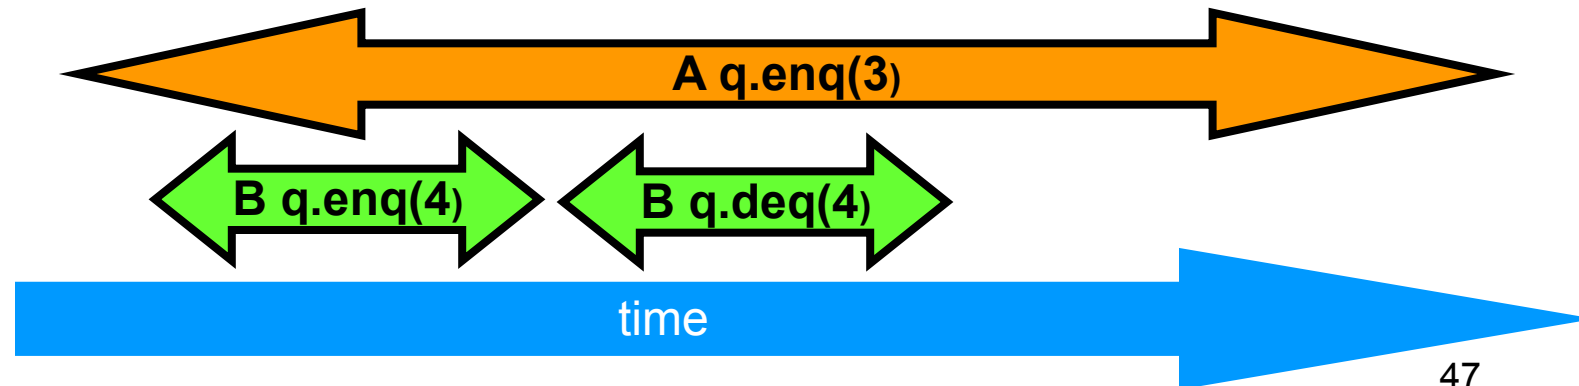
B q.enq(4)

B q:void

B q.deq()

B q:4

A q:void



# Example

A q.enq(3)

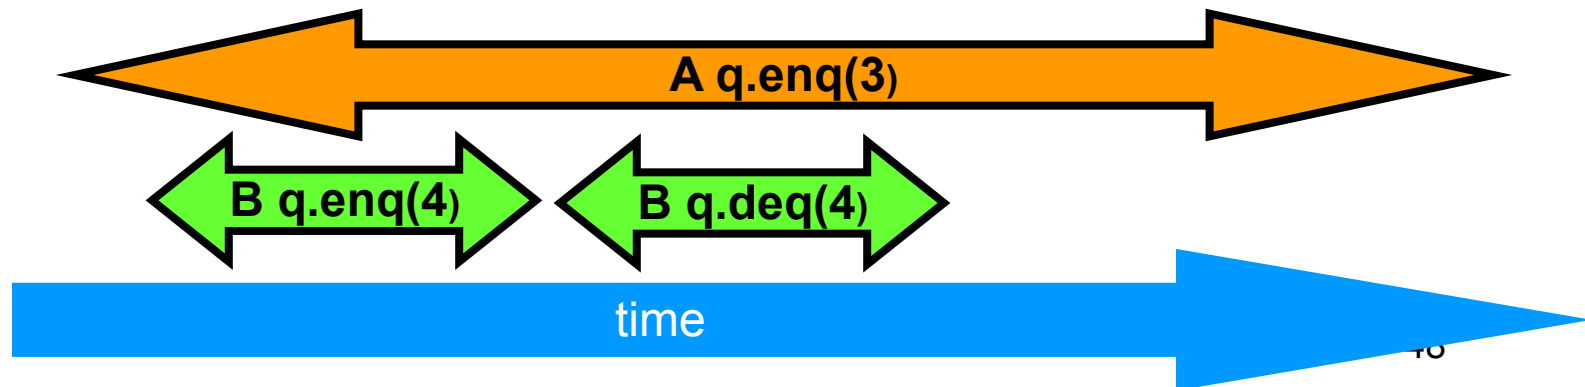
B q.enq(4)

B q:void

B q.deq()

B q:4

A q:void





# Example

A q.enq(3)

B q.enq(4)

B q:void

B q.deq()

B q:4

A q:void

B q.enq(4)

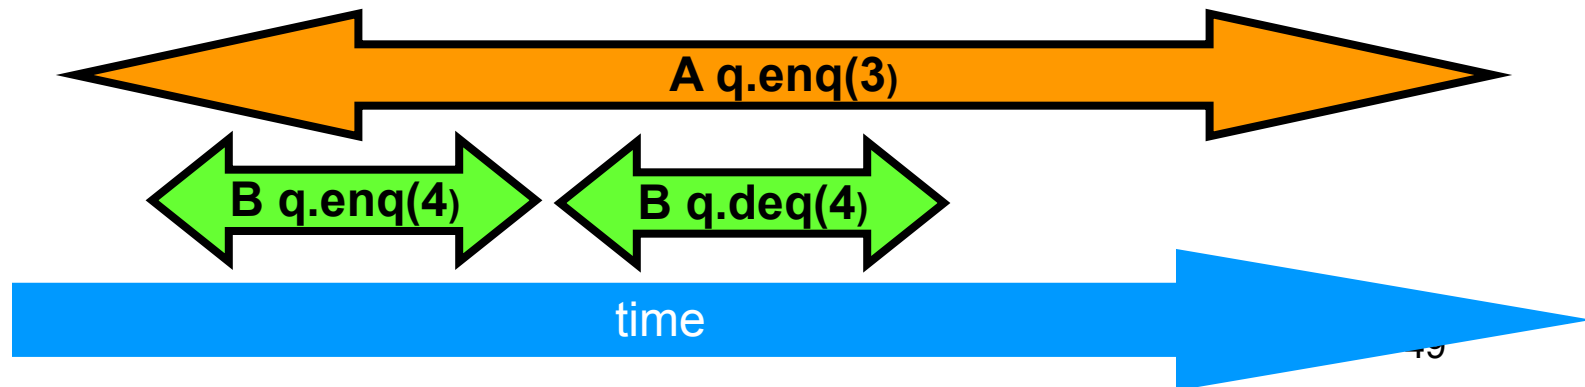
B q:void

A q.enq(3)

A q:void

B q.deq()

B q:4

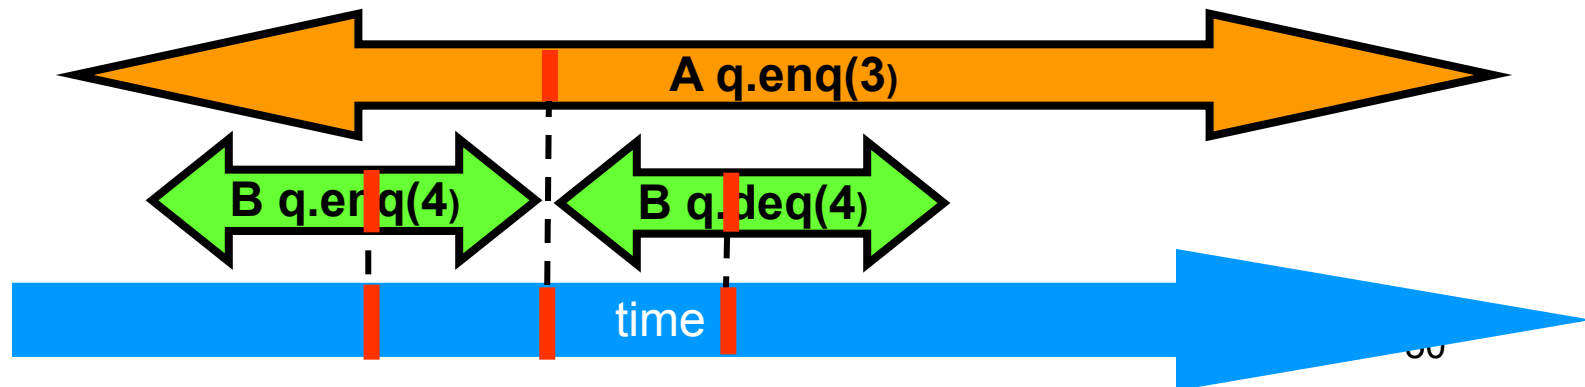


# Example

## Equivalent sequential history

A q.enq(3)  
B q.enq(4)  
B q:void  
B q.deq()  
B q:4  
A q:void

B q.enq(4)  
B q:void  
A q.enq(3)  
A q:void  
B q.deq()  
B q:4




# Why Does Composability Matter?

- Modularity
- Can prove linearizability of objects in isolation
- Can compose independently-implemented objects
  - A history of two linearizable objects is linearizable


# Reasoning About Linearizability: Locking

```
def deq() : T = {  
  myLock.lock()  
  try {  
    if (tail == head) {  
      throw EmptyException  
    }  
    val x = items(head % items.length)  
    head = head + 1  
    x  
  } finally {  
    myLock.unlock()  
  }  
}
```



# Reasoning About Linearizability: Locking


```
def deq() : T = {  
  myLock.lock()  
  try {  
    if (tail == head) {  
      throw EmptyException  
    }  
    val x = items(head % items.length)  
    head = head + 1  
    x  
  } finally {  
    myLock.unlock()  
  }  
}
```



Linearization points  
are when locks are  
released

# More Reasoning: Wait-free

```
class LockFreeQueue[T: ClassTag](val capacity: Int) {  
  
  @volatile  
  private var head, tail: Int = 0  
  private val items = new Array[T](capacity)  
  
  def enq(x: T): Unit = {  
    if (tail - head == items.length) throw FullException  
    items(tail % items.length) = x  
    tail = tail + 1  
  }  
  
  def deq(): T = {  
    if (tail == head) throw EmptyException  
    val x = items(head % items.length)  
    head = head + 1  
    x  
  }  
}
```



# More Reasoning: Wait-free

**Remember that there  
is only one enqueuer  
and only one dequeuer**

```
class LockFreeQueue[T: ClassTag](val capacity: Int) {  
    head, tail: Int = 0  
    items = new Array[T](capacity)  
  
    def enqueue(x: T): Unit = {  
        if (tail - head == capacity) throw EmptyException  
        items[tail % items.length] = x  
        tail = tail + 1  
    }  
  
    def dequeue(): T = {  
        if (tail == head) throw EmptyException  
        val x = items[head % items.length]  
        head = head + 1  
        x  
    }  
}
```

**Linearization order is order head  
and tail fields modified**

# Strategy

- Identify one atomic step where method “happens”
  - Critical section
  - Machine instruction
- Doesn't always work
  - Might need to define several different steps for a given method



# Linearizability: Summary

- Powerful specification tool for shared objects
- Allows us to capture the notion of objects being “atomic”
- Don't leave home without it

# Alternative: Sequential Consistency

- History  $H$  is **Sequentially Consistent** if it can be extended to  $G$  by
  - Appending zero or more responses to pending invocations
  - Discarding other pending invocations
- So that  $G$  is equivalent to a
  - Legal sequential history  $S$
  - ~~– Where  $\exists G \subseteq \exists S$~~


Differs from  
linearizability




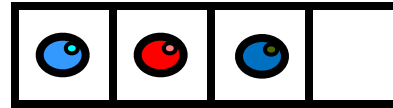
# Sequential Consistency

- No need to preserve real-time order
  - Cannot re-order operations done by the same thread
  - Can re-order non-overlapping operations done by different threads
- Often used to describe multiprocessor memory architectures


# Example



# Example




# Example






# Example





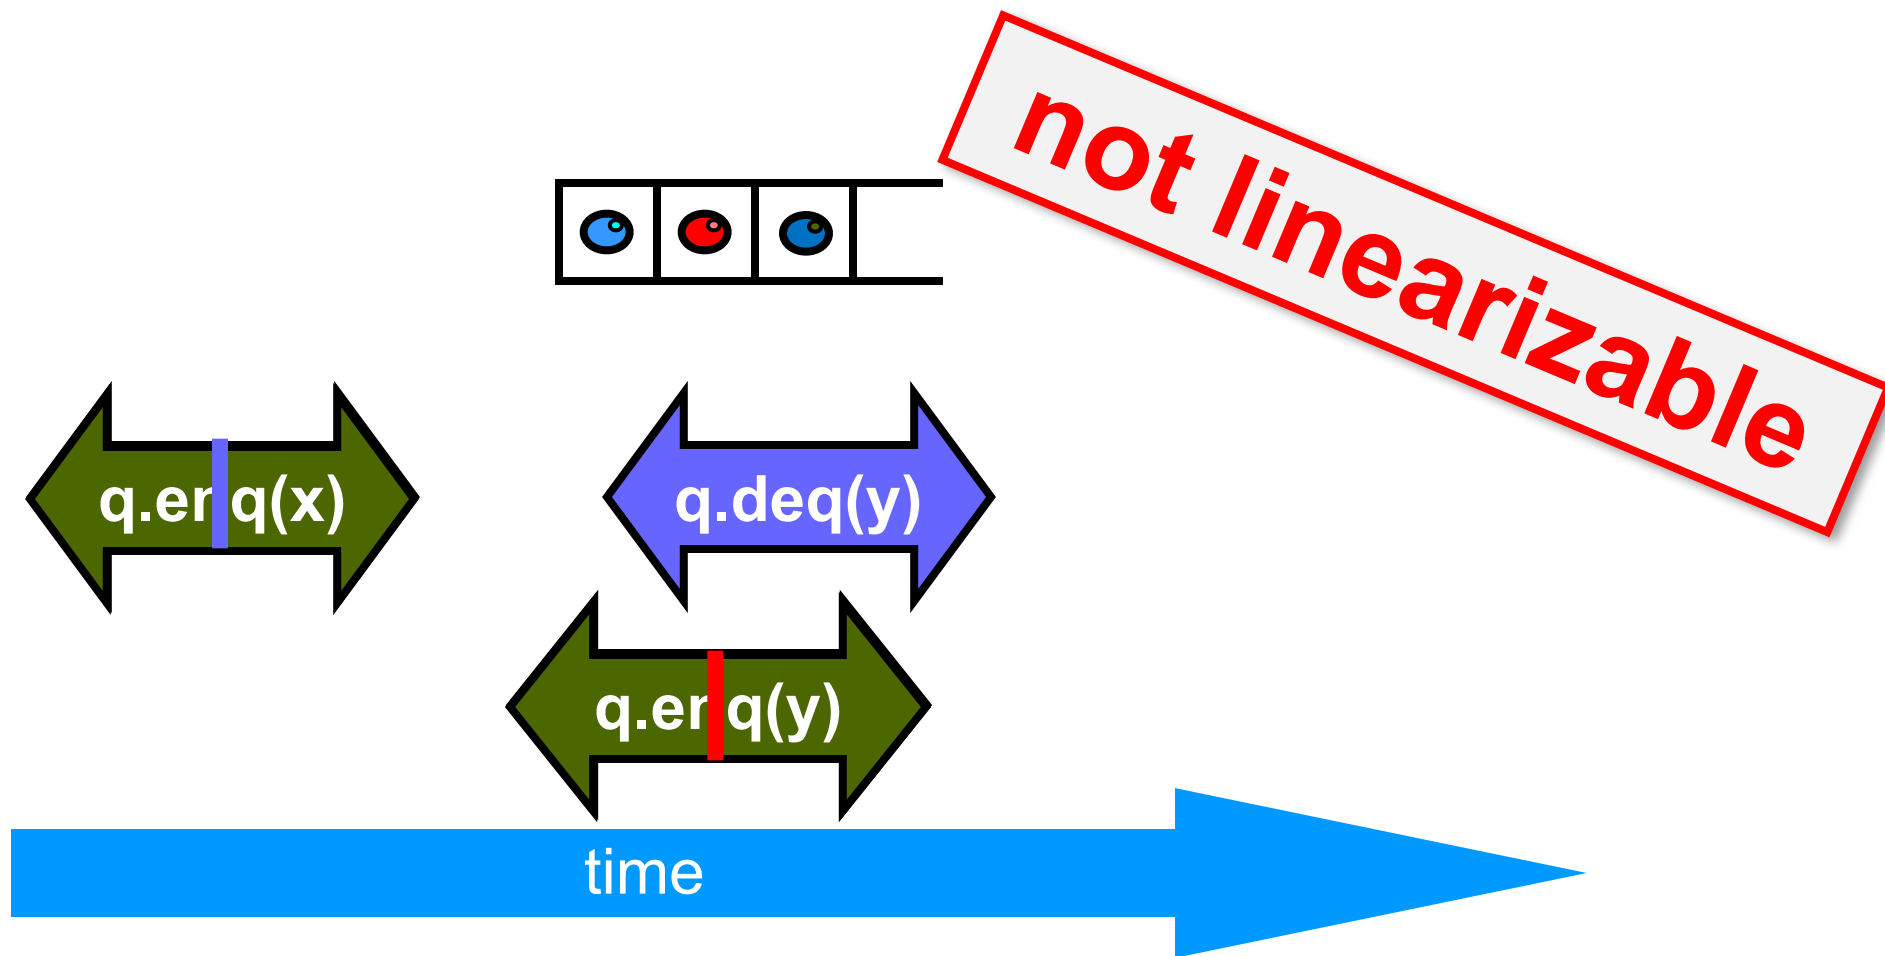
# Example







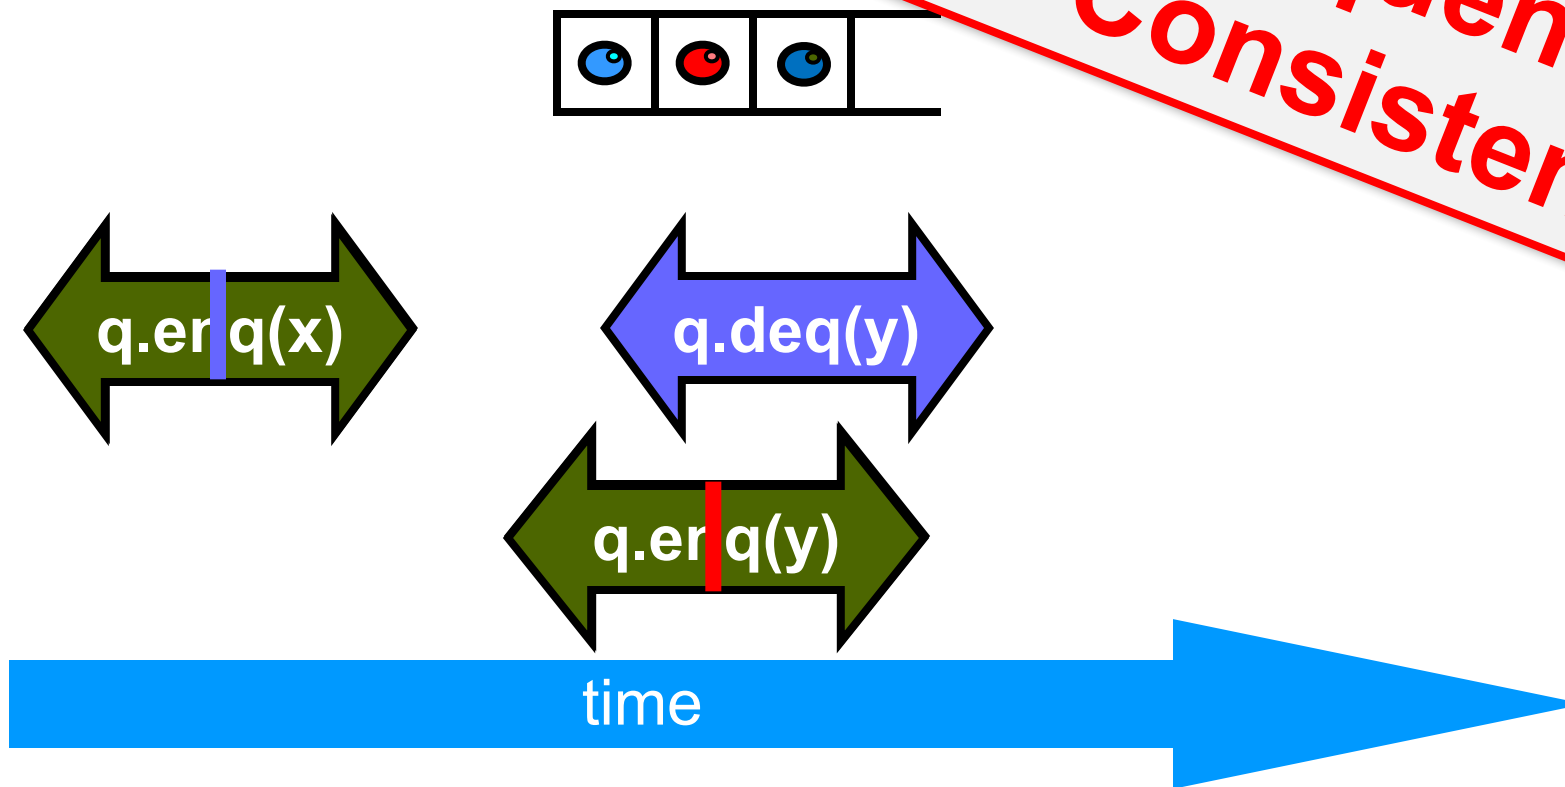
# Example





# Example


**Yet Sequentially Consistent**




# Theorem

Sequential Consistency is not composable


# FIFO Queue Example




# FIFO Queue Example




# FIFO Queue Example



# H/p Sequentially Consistent




# H/q Sequentially Consistent







# Ordering imposed by p




# Ordering imposed by q



# Ordering imposed by both




# Combining orders




# Fact

- Most hardware architectures don't support sequential consistency
- Because they think it's too strong
- Here's another story ...

# The Flag Example




# The Flag Example



- Each thread's view is sequentially consistent
  - It went first


# The Flag Example



- Entire history isn't sequentially consistent
  - Can't both go first



# The Flag Example



- Is this behavior really so wrong?
  - We can argue either way ...

# Opinion: It's Wrong

- This pattern
  - Write mine, read yours
- Is exactly the flag principle
  - Beloved of Alice and Bob
  - Heart of mutual exclusion
    - Peterson
    - Bakery, etc.
- It's non-negotiable!

# Peterson's Algorithm

```
def lock(): Unit = {  
    flag(i) = true  
    victim = i  
    while (flag(1 - i) && victim == i) {}  
}  
  
def unlock(): Unit = {  
    val i = ThreadID.get  
    flag(i) = false  
}
```

# Crux of Peterson Proof

(1)  $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$

(3)  $\text{write}_B(\text{victim}=B) \rightarrow$

(2)  $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$   
 $\rightarrow \text{read}_A(\text{victim})$

# Crux of Peterson Proof

- (1)  $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$
- (3)  $\text{write}_B(\text{victim}=B) \rightarrow$
- (2)  $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$   
 $\rightarrow \text{read}_A(\text{victim})$

Observation: proof relied on fact that if a location is stored, a later load by some thread will return this or a later stored value.

# Opinion: But It Feels So Right ...

- Many hardware architects think that sequential consistency is too strong
- Too expensive to implement in modern hardware
- OK if flag principle
  - violated by default
  - Honored by explicit request

# Hardware Consistency

**Initially,  $a = b = 0$ .**

**Processor 0**

```
mov 1, a    ;Store  
mov b, %ebx ;Load
```

**Processor 1**

```
mov 1, b    ;Store  
mov a, %eax ;Load
```

What are the final possible values of `%eax` and `%ebx` after both processors have executed?

Sequential consistency implies that no execution ends with `%eax = %ebx = 0`


# Hardware Consistency

- No modern-day processor implements sequential consistency.
- Hardware actively reorders instructions.
- Compilers may reorder instructions, too.
- Why?
- Because most of performance is derived from a single thread's unsynchronized execution of code!

This is known as **Weak (Relaxed) Memory Semantics**



# Weak-Memory Instruction Reordering




Program Order

Execution Order

- Q. Why might the hardware or compiler decide to reorder these instructions?
- A. To obtain higher performance by covering load latency — *instruction-level parallelism*.

# Weak-Memory Instruction Reordering




Q. When is it safe for the hardware or compiler to perform this reordering?

A. When  $a \neq b$ .

A'. And there's no concurrency.


# Hardware Reordering



- Processor can issue stores faster than the network can handle them  $\Rightarrow$  store buffer.
- Loads take priority, bypassing the store buffer.
- Except if a load address matches an address in the store buffer, the store buffer returns the result.

# X86 Relaxed Memory Model

## Thread's Code



A vertical list of memory operations on a light yellow background with a folded bottom-right corner. The operations are: Store1, Store2, Load1, Load2, Store3, Store4, Load3, Load4, and Load5. Red 'X' marks are drawn over Store1, Store2, Load2, Store3, Load4, and Load5.

~~Store1~~  
~~Store2~~  
Load1  
~~Load2~~  
~~Store3~~  
Store4  
Load3  
~~Load4~~  
~~Load5~~

1. Loads *are not* reordered with loads.
2. Stores *are not* reordered with stores.
3. Stores *are not* reordered with prior loads.
4. A load *may* be reordered with a prior store to a different location *but not* with a prior store to the same location.
5. Stores to the same location *respect a global total order*.

# X86 Relaxed Memory Model

## Thread's Code

Store1  
Store2  
Load1  
Load2  
Store3  
Store4  
Load3  
Load4  
Load5

↑  
LOADS

1. Loads *are not* reordered with loads.
2. Stores *are not* reordered with stores.
3. **Total Store Ordering (TSO)** ...weaker than sequential consistency
4. *...prior*
5. Stores to the same location *respect a global total order.*

**OK!**

# Memory Barriers (Fences)

- *A memory barrier (or memory fence) is a hardware action that enforces an ordering constraint between the instructions before and after the fence.*
- *A memory barrier can be issued explicitly as an instruction (x86: mfence)*
- *The typical cost of a memory fence is comparable to that of an L2-cache access.*

# X86 Relaxed Memory Model

## Thread's Code

Store1

Store2

Load1

Load2

Store3

Store4

**Barrier**

Load3

Load4

Load5

1. Loads
2. Store
3. Store loads
4. A load store with a prior store to the same location.
5. Stores to the same location respect a global total order.

Total Store Ordering +  
properly placed memory barriers  
= sequential consistency

# Memory Barriers

- Explicit Synchronization
- Memory barrier will
  - Flush write buffer
  - Bring caches up to date
- Compilers often do this for you
  - Entering and leaving critical sections



# Java/Scala Volatile Variables

- In Java, can ask compiler to keep a variable up-to-date by declaring it **volatile**
- In Scala, use **@volatile** annotation
- Adds a memory barrier after each store
- Inhibits reordering, removing from loops, & other “compiler optimizations”

Demo: Realistic Locks

# Summary: Real-World

- Hardware weaker than sequential consistency
- Can get sequential consistency at a price
- Linearizability better fit for high-level software

# Linearizability

- Linearizability
  - Operation takes effect instantaneously between invocation and response
  - Uses sequential specification, locality implies composability

# Summary: Correctness

- Sequential Consistency
  - Not composable
  - Harder to work with
  - Good way to think about hardware models
- We will use *linearizability* as our consistency condition in the remainder of this course unless stated otherwise



This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](http://creativecommons.org/licenses/by-sa/3.0/).

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.